

# CS 546 – Advanced Topics in NLP

Dilek Hakkani-Tür



UNIVERSITY OF  
**ILLINOIS**  
URBANA-CHAMPAIGN



Siebel School of  
Computing  
and Data Science



# Topic for Today: Reminders

- Machine Learning (ML) for NLP
- ML Examples: NLP Tasks
- ML Basics
- Calculus Reminders
- Deep Learning
- Linear Regression – Case Study

# What is Machine Learning?

- Machine learning (ML) is the field of artificial intelligence that develops **algorithms** and **statistical models** that **computers use to perform a specific task**
  - without using explicit instructions (i.e., coding the specific steps for performing the task),
  - relying on patterns and inference instead.

# Example Task: Sentiment Detection for Product Reviews



Approach for solving tasks:

Input: “I love this camera!”

↓  
**program.py**

+

if input contains “love”, “like”, etc.  
output = positive

“It makes too much noise.”

↓  
**program.py**

-

if input contains “too much”, “bad”, etc.  
output = negative

“It’s a very bulky.”

↓  
**program.py**

?

Some tasks are complex, and we don't know how to write a program to solve them

Example: speech recognition, summarization, ...

# Example Task: Sentiment Detection for Product Reviews



- Learning  $\approx$  Looking for a Function

Input: “I love this camera!”

$$\begin{array}{c} \downarrow \\ f \\ + \end{array}$$

“It makes too much noise.”

$$\begin{array}{c} \downarrow \\ f \\ - \end{array}$$

“It’s a very bulky.”

$$\begin{array}{c} \downarrow \\ f \\ ? \end{array}$$

Given a large amount of data, the machine learns what the function  $f$  should be.

# Other Example Tasks

- Speech Recognition

$$f( \quad \text{[sound波形图]} \quad ) = \text{"the book"}$$

- Handwriting Recognition

$$f( \quad \text{[手写数字2]} \quad ) = \text{"2"}$$

- Weather forecast

$$f( \quad \text{[晴天图标]} \quad \text{Thursday} \quad ) = \text{“} \quad \text{[下雨图标]} \quad \text{Saturday”}$$

- Playing video games

$$f( \quad \text{[游戏截图]} \quad ) = \text{“move left”}$$



# Topic for Today: Reminders

- Machine Learning (ML) for NLP

- ML Examples: NLP Tasks

- ML Basics

- Calculus Reminders

- Deep Learning

- Linear Regression – Case Study

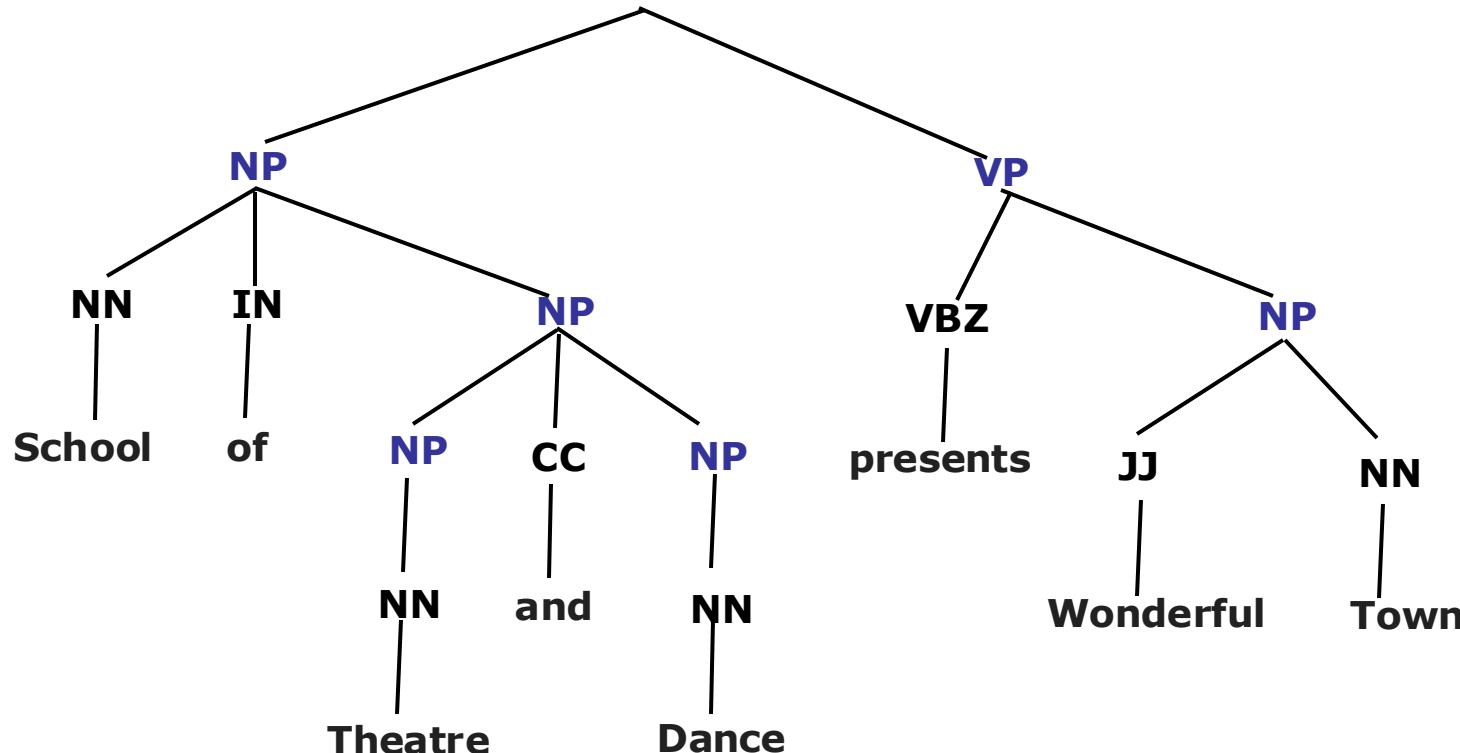
# Other Examples: Classical NLP Pipeline



- Sentence Segmentation
- Tokenization

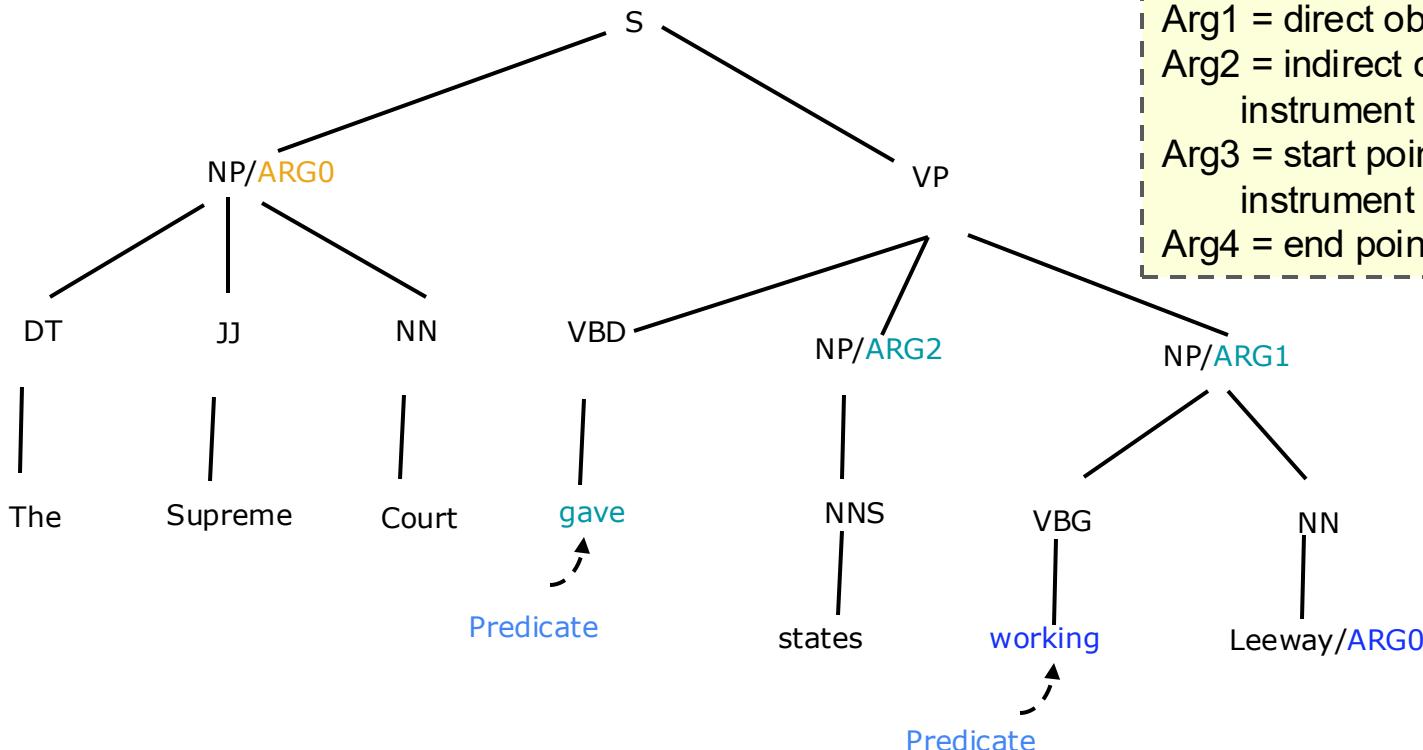
# Other Examples: Classical NLP Pipeline

- POS tagging and Syntactic Parsing



# Other Examples: Classical NLP Pipeline

- Semantic Role Labeling



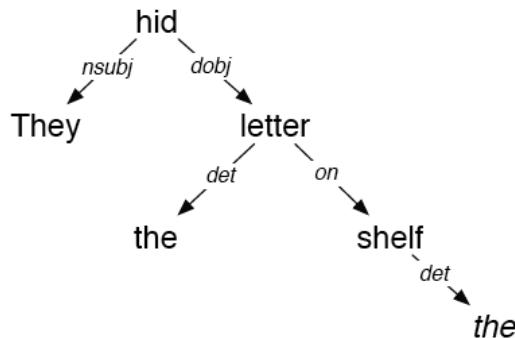
## CORE ARGUMENTS

Arg0 = agent  
 Arg1 = direct object / theme / patient  
 Arg2 = indirect object / benefactive / instrument / attribute / end state  
 Arg3 = start point / benefactive / instrument / attribute  
 Arg4 = end point

# Other Examples: Classical NLP Pipeline



## - Dependency Parsing



They hid the letter on the shelf.

Argument Dependencies	Description
<b>nsubj</b>	nominal subject
<b>csubj</b>	clausal subject
<b>dobj</b>	direct object
<b>iobj</b>	indirect object
<b>pobj</b>	object of preposition

Modifier Dependencies	Description
<b>tmmod</b>	temporal modifier
<b>appos</b>	appositional modifier
<b>det</b>	determiner
<b>prep</b>	prepositional modifier

# Other Examples: Classical NLP Pipeline



- Named Entity Extraction

Since its inception in 2001, `<name ID="1" type="organization">Red</name>` has caused a stir in `<name ID="2" type="location">Northeast Ohio</name>` by stretching the boundaries of classical by adding multi-media elements to performances and looking beyond the expected canon of composers.

Under the baton of `<name ID="3" type="organization"> Red</name>` Artistic Director `<name ID="4" type="person"> Jonathan Sheffer</name>`, `<name ID="5" type="organization">Red</name>` makes its debut appearance at `<name ID="6" type="organization"> Kent State University</name>` on March 7 at 7:30 p.m.

# Other Examples: Classical NLP Pipeline



- Coreference Resolution

But **the little prince** could not restrain admiration:

"Oh! How beautiful **you** are!"

"Am **I** not?" **the flower** responded, sweetly. "And **I** was born at the same moment as **the sun** . . ."

**The little prince** could guess easily enough that **she** was not any too modest--but how moving--and exciting--**she** was!

"**I** think it is time for breakfast," **she** added an instant later. "If **you** would have the kindness to think of my needs--"

And **the little prince**, completely abashed, went to look for a sprinkling-can of fresh water. So, **he** tended **the flower**.



# Other Examples: NLP Tasks

- Machine Translation

$$f(\text{ "Kitap masada." }) = \text{ "The book is on the table."}$$

- Question Answering

$$f(\text{ "Who wrote Leviathan?" }) = \text{ "Paul Auster"}$$

- Summarization

$$f(\text{ [document icon] }) = \text{ [summary icon]}$$



# Topic for Today: Reminders

- Machine Learning (ML) for NLP
- ML Examples: NLP Tasks
- ML Basics
- Calculus Reminders
- Deep Learning
- Linear Regression – Case Study

# Kinds of Machine Learning



- Supervised learning
- Unsupervised learning
- Self-supervised learning
- Reinforcement learning

# Supervised Learning

- Learning functions for predicting **targets** from given inputs, using pairs of **inputs and targets**.
- The targets, which we often call ***labels***, are generally denoted by  $y$ .
- The input data, also called the ***features*** or covariates, are typically denoted  $x$ .

# Supervised Learning (cont.)

$x$ : "It makes too much noise."

function input

$y$ : - (negative)

function output

Training  
:

**x**

Training inputs

Supervised learning

**y**

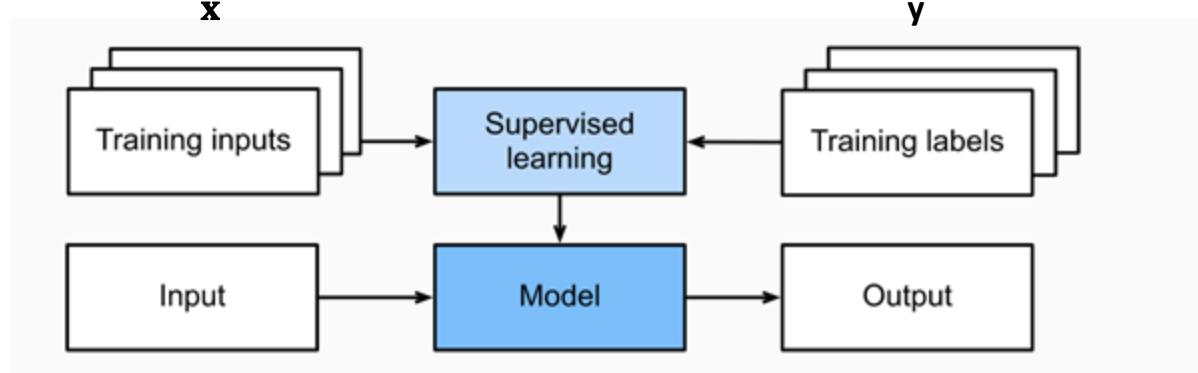
Training labels

Testing  
:

Input

Model

Output



# Supervised Learning (cont.)

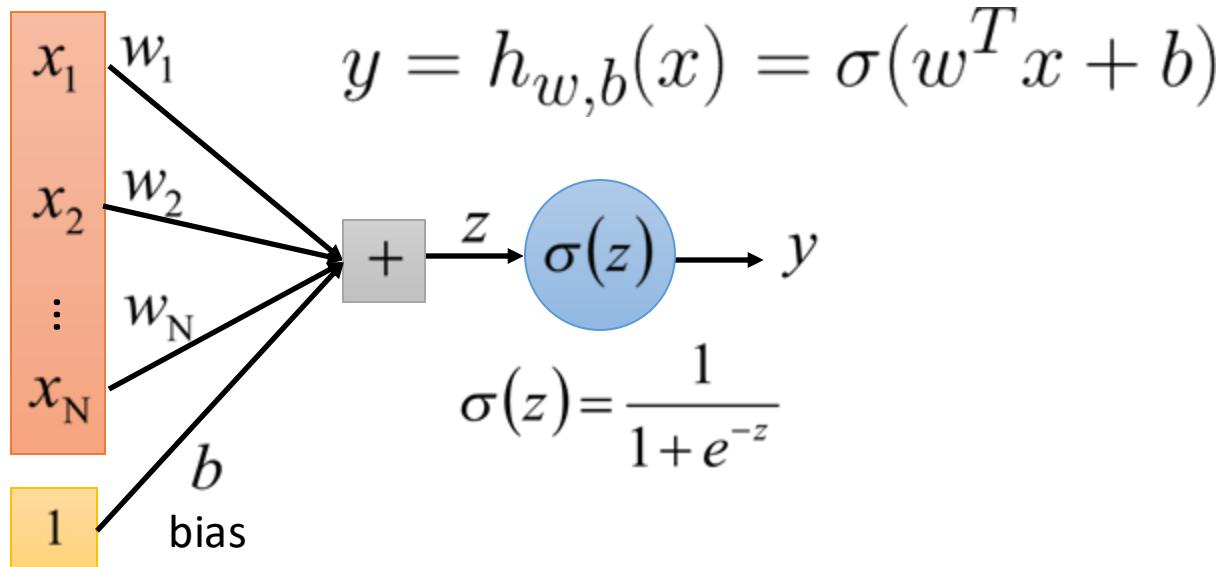
- Each (input, target) pair is called an *example* or an *instance*.
- Supervised learning learns the model/function from examples.
- For example, machine translation models are trained from pairs of source and target language sentences.

# Classification Example: Apartment search



- Whether an apartment will be preferred by a user or not?
- Features:
  - Number of rooms, e.g., 1, 2, 3
  - Whether the apartment has a dish washer or not, e.g. yes, no
  - Style of the apartment building, e.g., high-rise, low-rise, duplex,
  - ...

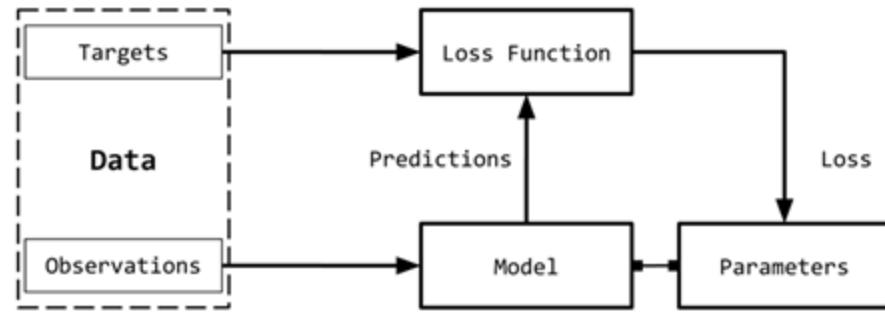
# Example Function: a Single Neuron



$w, b$  are the parameters of this neuron

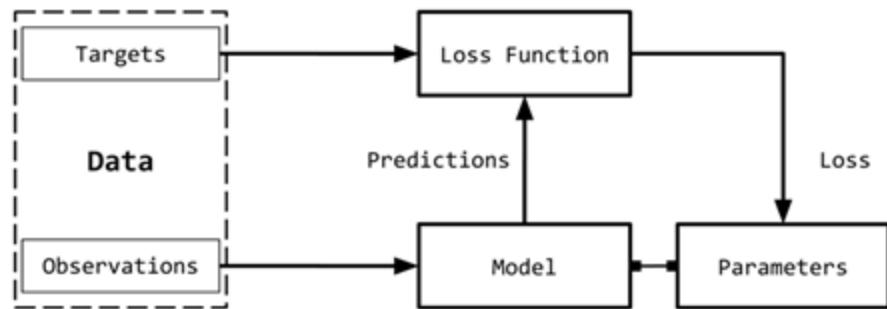
# Terminology

- **Observations:** inputs,  $x$ .
- **Targets:** labels,  $y$ , corresponding to the observations (also called ground truth sometimes).
- **Model:** Function that takes an observation and predicts the value of its target label.



# More Terminology

- **Parameters:** weights (for deep learning),  $w$ , that parameterize the model.
- **Predictions:** Also called estimates, target values estimated by the model,  $\hat{y}$ .
- **Loss Function:** a function (denoted  $\mathcal{L}$ ) that compares how far off a prediction is from its target. The lower the loss, the better is the model at predicting the target.



# Supervised Learning - Regression



- When the targets take arbitrary values in a range (instead of specific categories such as +/- for sentiment classification), we call it a **regression** problem.
- Examples:
  - Estimating the price of a house given attributes such as location, square footage, number of rooms, etc.
  - Predicting how many miles can a dog run given attributes such as its breed, age, size, etc.

# Supervised Learning – Classification



- When the targets take a limited number of pre-defined categories, we call it a classification problem.
- **Binary classification:** only two values
  - Positive/negative sentiment classification
  - Sentence segmentation from speech: each word boundary is classified into sentence boundary versus not a sentence boundary.
  - Sentence segmentation for text: each “.” is classified into two classes similarly.
- **Multi-class classification:** more than two values
  - Digit recognition: each digit is classified into one of 10 values for 0,...,9.

# Supervised Learning - Tagging



- Some tasks are not simply binary or multi-class classification tasks, but we need to
  - Tagging images with multiple objects
  - Tagging tokens/words in natural language utterances



Locations in utterances:

I want to fly from Boston to Seattle via Chicago. =>  
I want to fly from **Boston** to **Seattle** via **Chicago**.

Sequence Tagging

# Kinds of Machine Learning

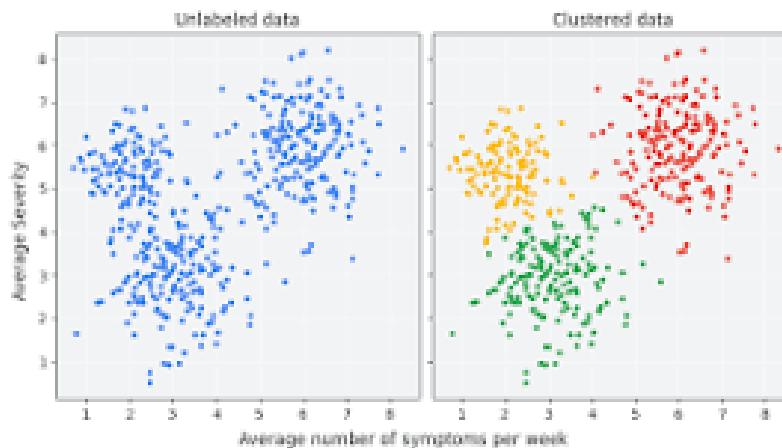


- Supervised learning
- Unsupervised learning
- Self-supervised learning
- Reinforcement learning

# Unsupervised Learning



- Finding previously unknown patterns in data set without pre-existing labels.
  - Clustering
  - Zero-Shot Learning
  - Few-Shot Learning
  - Semi-supervised Learning
  - Weakly-supervised Learning
  - Transfer Learning



# Kinds of Machine Learning

- Supervised learning
- Unsupervised learning
- Self-supervised learning
- Reinforcement learning

# Self-supervised Learning

- Benefiting from enormous amount of data available on the web and other resources.
- Examples:
  - For text, masking tokens and predicting them or predicting the next token,
  - For image, image reconstruction, and
  - For videos, predicting video frames.

# Kinds of Machine Learning

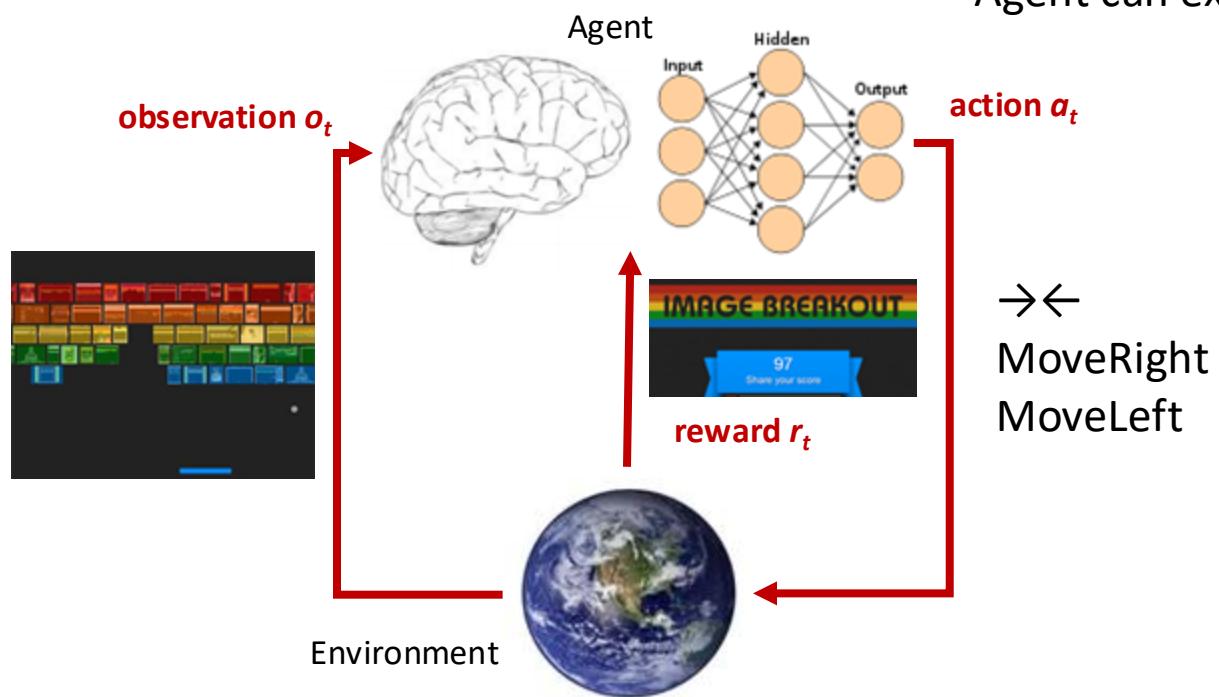
- Supervised learning
- Unsupervised learning
- Self-supervised learning
- Reinforcement learning

# Reinforcement Learning



- Learning action prediction policies through interactions with an environment over a series of *timesteps*.
  - At each timestep  $t$ , the learning agent receives some observation  $o_t$  from the environment and must choose an action  $a_t$  that is subsequently transmitted back to the environment via some mechanism.
  - Finally, the agent receives a reward  $r_t$  from the environment.
  - The agent then receives a subsequent observation, and chooses a subsequent action, and so on.
  - A *policy* is a function that maps from observations to actions.

# Agent and Environment





# Topic for Today: Reminders

- Machine Learning (ML) for NLP
- ML Examples: NLP Tasks
- ML Basics
- Calculus Reminders
- Deep Learning
- Linear Regression – Case Study

# Derivatives and Differentiation



- Training deep learning models: updating them successively so that they get better as they see more data
- Getting better: minimizing a loss function
- A crucial step in nearly all deep learning optimization algorithms.
- Loss functions that are differentiable with respect to our model's parameters.

# Derivatives and Differentiation (cont.)



- Suppose we have function  $f: \mathbb{R} \rightarrow \mathbb{R}$  with scalar input and outputs.
- The *derivative* of  $f$  is defined as:

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}$$

- If  $f'(a)$  exists,  $f$  is said to be *differentiable* at  $a$ .
- If  $f$  is differentiable at every number of an interval, then this function is differentiable on this interval.
- Alternative notations for derivatives:

$$f'(x) = y' = \frac{dy}{dx} = \frac{df}{dx} = \frac{d}{dx} f(x) = Df(x) = D_x f(x).$$

# Rules for Differentiating Common Functions



- $D\mathbf{C} = 0$  ( $\mathbf{C}$  is a constant),
- $Dx^n = nx^{n-1}$  (the *power rule*,  $n$  is any real number),
- $De^x = e^x$ ,
- $D \ln(x) = 1/x$ .

# Rules for Differentiating Common Functions



$$\frac{d}{dx}[Cf(x)] = C \frac{d}{dx}f(x),$$

Constant multiple rule

$$\frac{d}{dx}[f(x) + g(x)] = \frac{d}{dx}f(x) + \frac{d}{dx}g(x),$$

Sum rule

$$\frac{d}{dx}[f(x)g(x)] = f(x)\frac{d}{dx}[g(x)] + g(x)\frac{d}{dx}[f(x)],$$

Product rule

$$\frac{d}{dx}\left[\frac{f(x)}{g(x)}\right] = \frac{g(x)\frac{d}{dx}[f(x)] - f(x)\frac{d}{dx}[g(x)]}{[g(x)]^2}.$$

Quotient rule

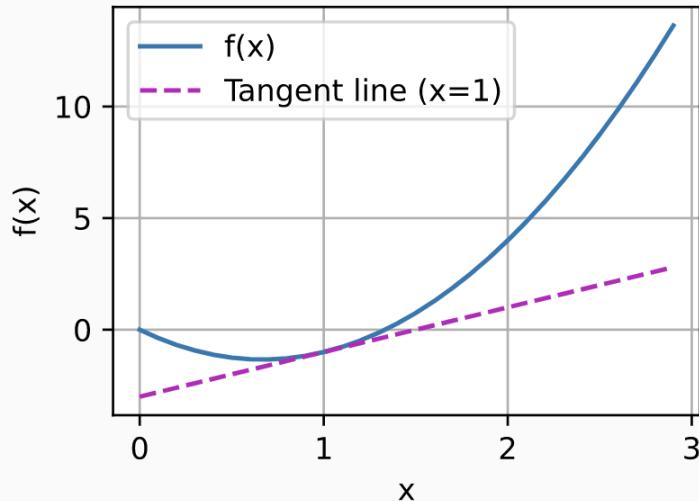
# Example



- Given function,  $u = f(x) = 3x^2 - 4x$
- We can apply the previous set of rules to get its derivative:

$$u' = f'(x) = 3 \frac{d}{dx} x^2 - 4 \frac{d}{dx} x = 6x - 4$$

- For  $x=1$ , we have  $u'=2$ .
- This derivative is also the slope of the tangent line to the curve  $u=f(x)$  when  $x=1$ .



# Partial Derivatives



- If  $y$  is a function with  $n$  variables, i.e.,

$$y=f(x_1, x_2, \dots, x_n)$$

- The **partial derivative** of  $y$  with respect to its  $i^{\text{th}}$  parameter  $x_i$  is:

$$\frac{\partial y}{\partial x_i} = \lim_{h \rightarrow 0} \frac{f(x_1, \dots, x_{i-1}, x_i + h, x_{i+1}, \dots, x_n) - f(x_1, \dots, x_i, \dots, x_n)}{h}.$$

# Gradients



- For functions with multiple variables, we obtain the **gradient vector** of the function by concatenating partial derivatives of that function w.r.t. all its variables.
- For example, for  $f: \mathbb{R}^n \rightarrow \mathbb{R}$  is a function with an  $n$ -dimensional vector  $\mathbf{x} = [x_1, x_2, \dots, x_n]^T$  input and a scalar output. The gradient of the function  $f(\mathbf{x})$  w.r.t.  $\mathbf{x}$  is a vector of  $n$  partial derivatives:

$$\nabla_{\mathbf{x}} f(\mathbf{x}) = \left[ \frac{\partial f(\mathbf{x})}{\partial x_1}, \frac{\partial f(\mathbf{x})}{\partial x_2}, \dots, \frac{\partial f(\mathbf{x})}{\partial x_n} \right]^T$$

# The Chain Rule

- Suppose that functions  $y=f(u)$  and  $u=g(x)$  are both differentiable, then,

$$\frac{dy}{dx} = \frac{dy}{du} \frac{du}{dx}$$

- More general scenario, where a differentiable function  $y$  has variables  $u_1, u_2, \dots, u_m$ , where each differentiable function  $u_i$  has variables  $x_1, x_2, \dots, x_n$ . Then, for any  $i=1, 2, \dots, n$ ,

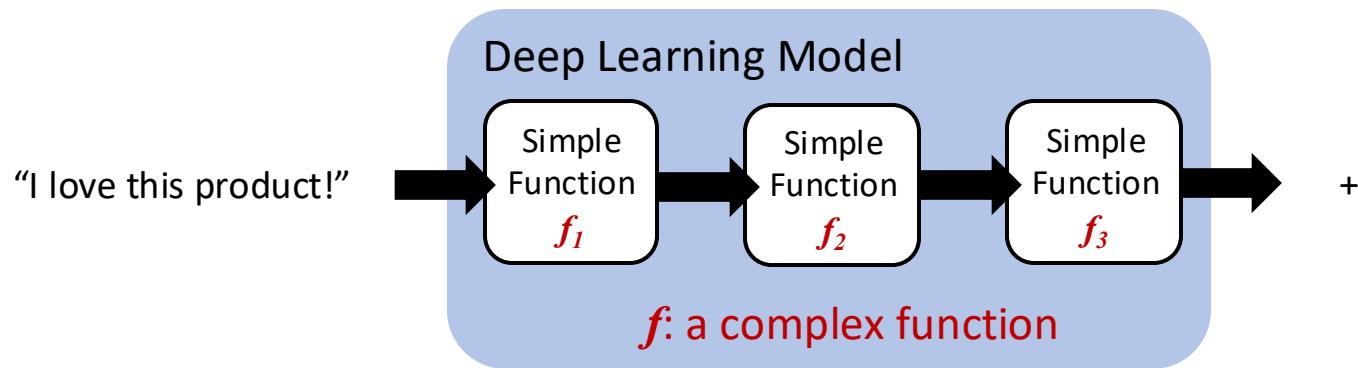
$$\frac{dy}{dx_i} = \frac{dy}{du_1} \frac{du_1}{dx_i} + \frac{dy}{du_2} \frac{du_2}{dx_i} + \dots + \frac{dy}{du_m} \frac{du_m}{dx_i}$$



# Topic for Today: Reminders

- Machine Learning (ML) for NLP
- ML Examples: NLP Tasks
- ML Basics
- Calculus Reminders
- Deep Learning
- Linear Regression – Case Study

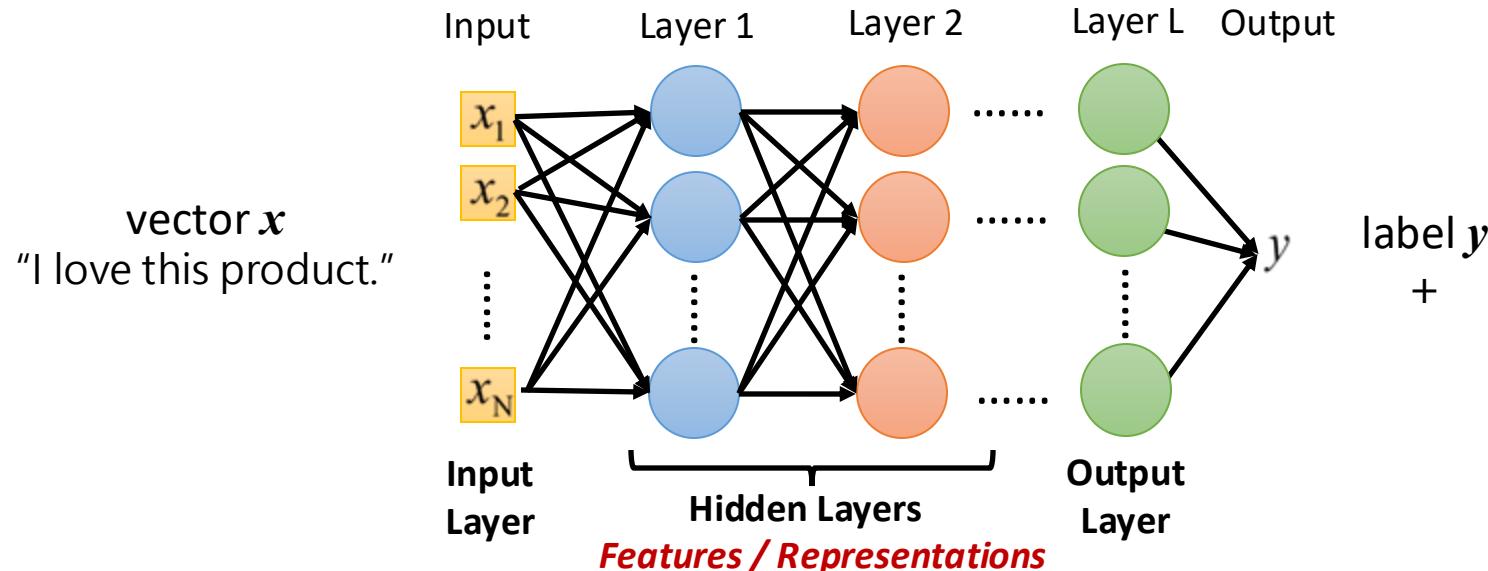
# Learning Stacked Functions



End-to-end training: what each function should do is learned automatically

Deep learning usually refers to *neural network*-based model

# Learning Stacked Functions (cont.)



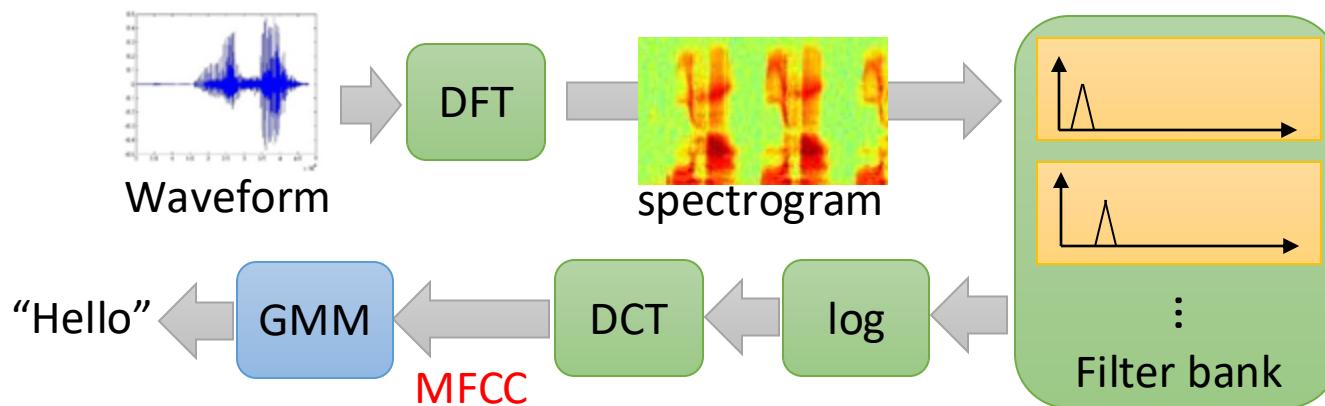
Representation Learning attempts to learn good features/representations

Deep Learning attempts to learn (multiple levels of) representations and an output

# Example: Speech Recognition (Before DL)



## - Shallow Model



Each box is a simple function in the production line:



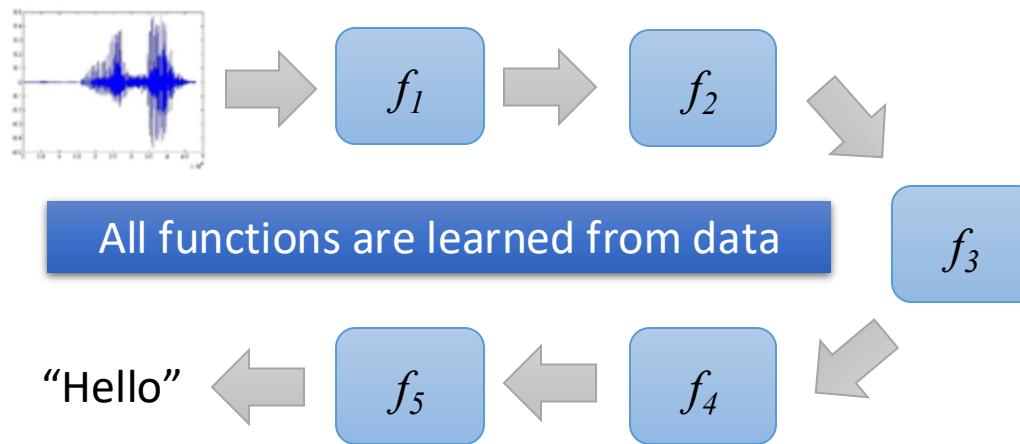
:hand-crafted



:learned from data

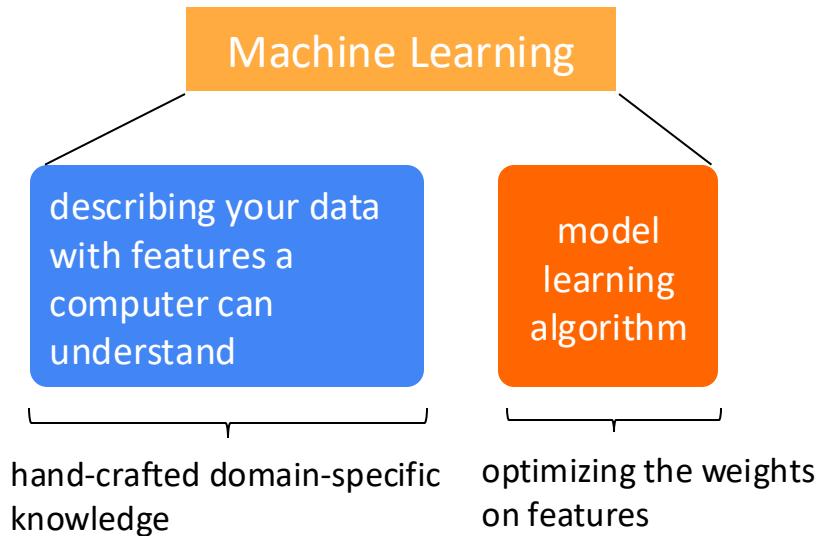
# Example: Speech Recognition (With DL)

- Deep Model

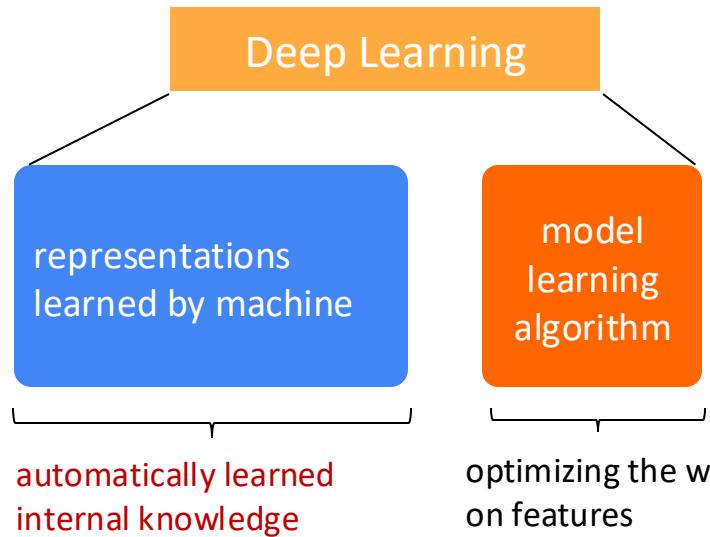


Less engineering labor, but machine learns more

# Machine Learning

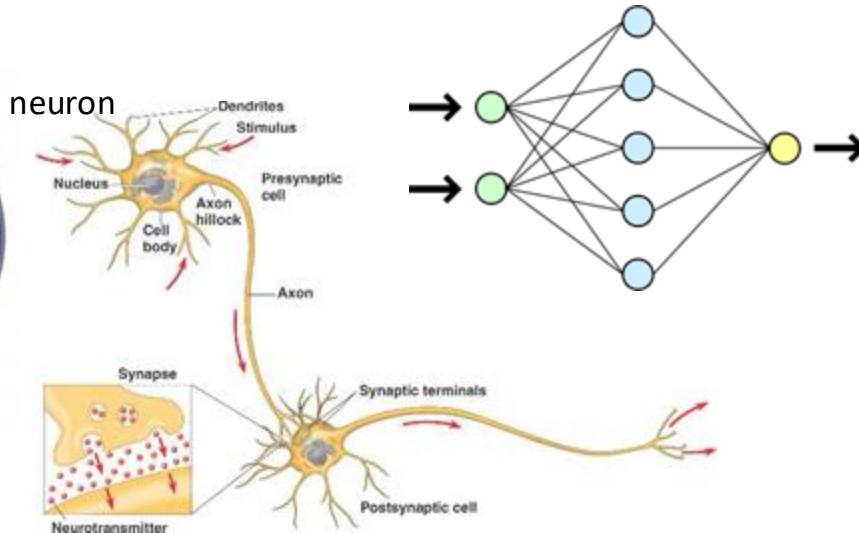


# Deep Learning



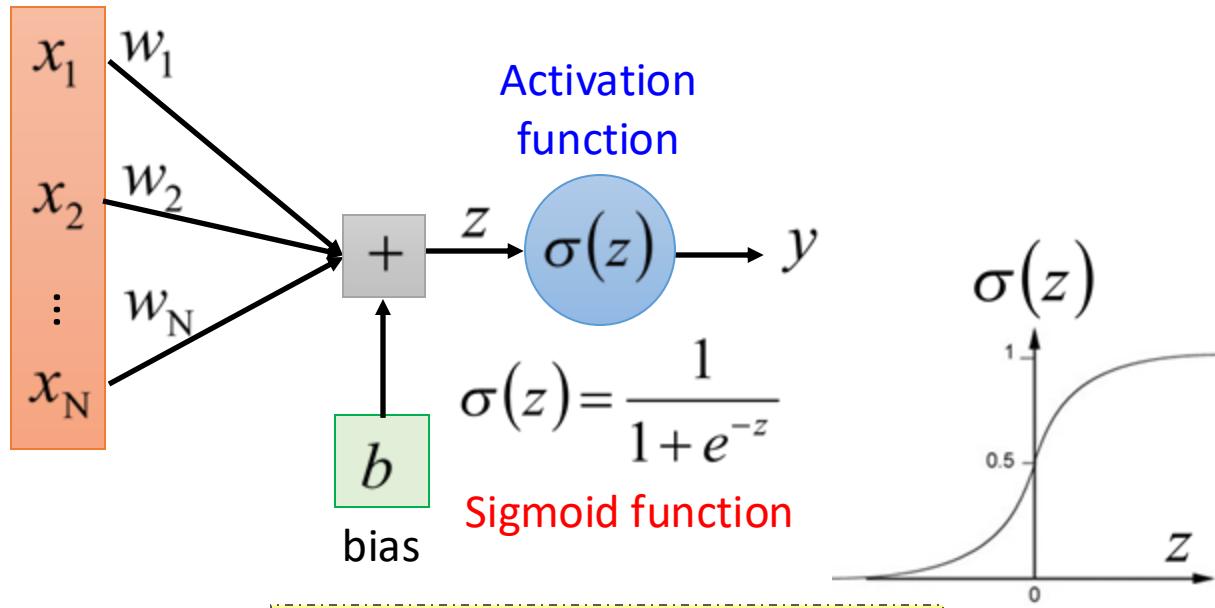
Deep learning usually refers to *neural network* based model

# Inspired by the Human Brain



Neurotransmission section of: [https://en.wikipedia.org/wiki/Human\\_brain](https://en.wikipedia.org/wiki/Human_brain)

# A Single Neuron



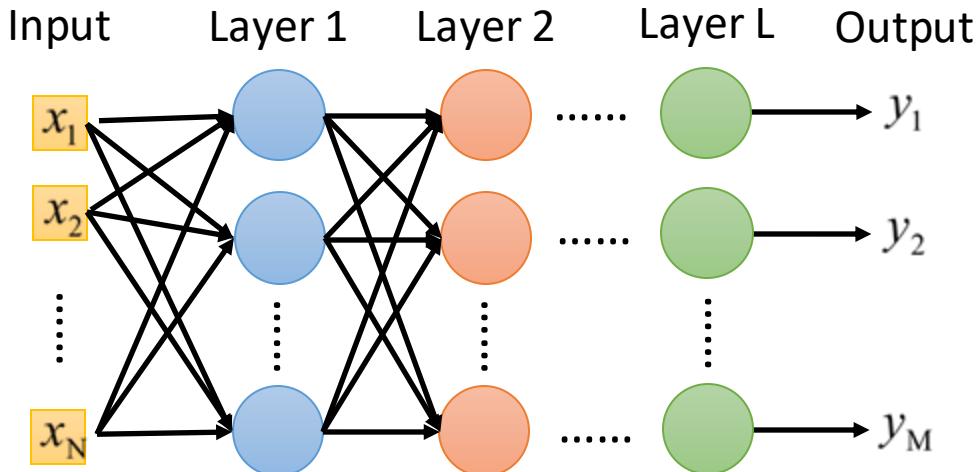
# Deep Neural Network



A neural network is a complex function:

$$f : R^N \rightarrow R^M$$

- Cascading the neurons to form a neural network



Each layer is a simple function in the production line

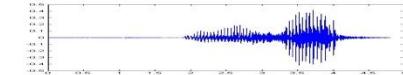


# Topic for Today: Reminders

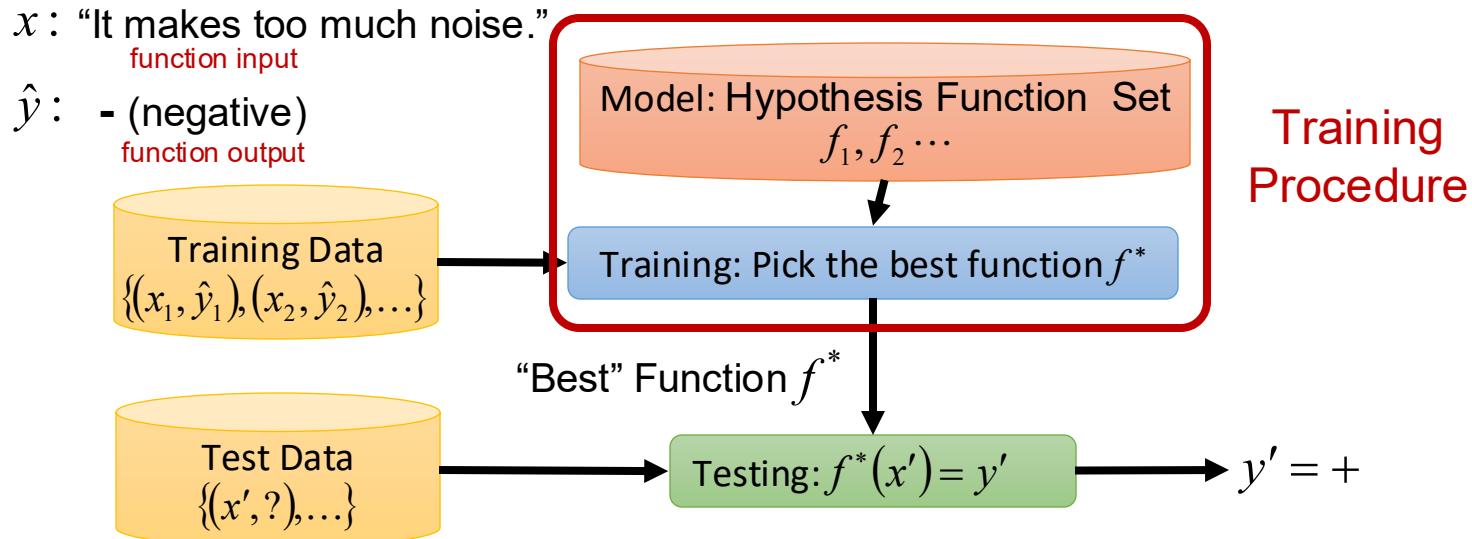
- Machine Learning (ML) for NLP
- ML Examples: NLP Tasks
- ML Basics
- Calculus Reminders
- Deep Learning
- Linear Regression – Case Study

# Learning ≈ Looking for a Function



- Speech Recognition  $f($   ) = “the book”
- Handwriting Recognition  $f($   ) = “2”
- Weather forecast  $f($   Thursday ) = “ Saturday”
- Play video games  $f($   ) = “move left”

# Machine Learning Framework

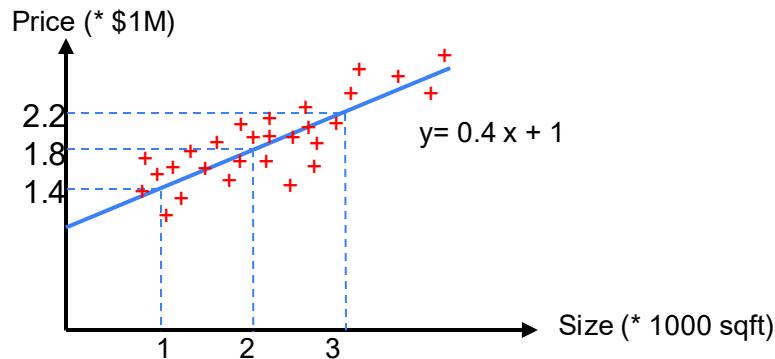


Training aims to pick the best function given the observed examples  
Testing predicts the label using the learned function

# Linear Regression

- **Regression:** modeling the relationship between data points  $x$  and corresponding real-valued targets  $y$ 
  - Example: Predicting prices of homes
  - Label: Price,  $y^{(i)}$
  - Features: Size, Age
    - $x^{(i)} = [x_{1}^{(i)}, x_{2}^{(i)}]$

	Size (sqft)	Price
2019 Sales:	2100	\$1.90M
	1600	\$1.56M
	3200	\$2.40M
	...	...



# Linear Regression (cont.)

- Assumption:
  - the relationship between the *features*  $\mathbf{x}$  and targets  $y$  is linear,
  - i.e.,  $y$  can be expressed as a weighted sum of the inputs  $\mathbf{x}$ , give or take some noise on the observations



$$\text{price} = w_1 \cdot x_1 + w_2 \cdot x_2 + b$$

weights                      Bias  
(offset or intercept)

# Linear Regression: Modeling Bivariate Data



- Bivariate data:  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Model:  $y_i = f(x_i) + e_i$  Random error
- Supervised approach!
- Model allows us to predict the value of y for any given value of x.
  - x is called the independent or predictor variable.
  - y is the dependent or response variable.

# Other examples of $f(x)$

- lines:  $f(x) = mx + b$
- polynomials:  $f(x) = ax^2 + bx + c$
- others:  $f(x) = a/x + b$

$$f(x) = a \sin(x) + b$$

$$f(x) = a \sqrt{x}$$

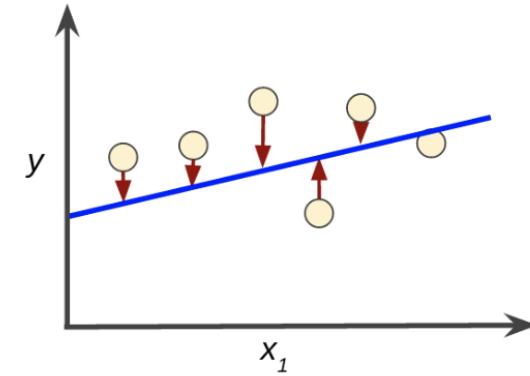
# Assumptions of Linear Regression

- The relationship between  $x$  and  $y$  is linear
- $y$  is distributed normally at each value of  $x$ , and the variance of  $y$  at every value of  $x$  is the same
- The observations are independent

# Mean Squared Error

- $y_i$  is the actual value,  $f(x_i)$  is the prediction, and their difference is the error.
- Commonly used function, (mean) squared error:

$$MS = \sum_{i=1}^n (y_i - f(x_i))^2$$



# Case Study: $f(x) = mx + b$

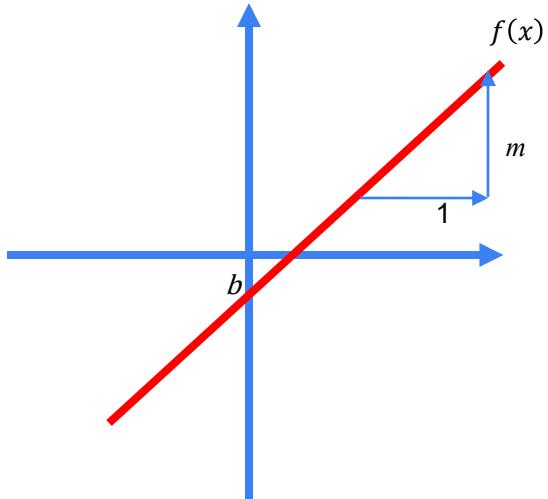
- (Mean) squared error:

$$MS(m, b) = \sum_{i=1}^n (y_i - mx_i - b)^2$$

- To estimate the values that minimize the mean squared error, we need to take the derivative with respect to the two variables  $m$  and  $b$ :

$$\frac{\partial MS(m, b)}{\partial b} = 0 \quad \text{and} \quad \frac{\partial MS(m, b)}{\partial m} = 0$$

# Case Study: $f(x) = mx + b$



- $b$  is called bias.
- $m$  is the slope of  $f(x)$
- Every time  $x$  moves by 1,  $y$  moves by  $m$

# Case Study: $f(x) = mx + b$

- We need to solve these two equations to find  $m$  and  $b$ , to be able to find the function  $f(x_i)$ .
- Beginning with the first one, for  $b$ :

$$\frac{\partial SS(m, b)}{\partial b} = \frac{\partial \sum_{i=1}^n (y_i - mx_i - b)^2}{\partial b} = \sum_{i=1}^n 2(y_i - mx_i - b)(-1) = 0$$

$$\sum_{i=1}^n y_i - m \sum_{i=1}^n x_i - nb = 0$$

$$b = \frac{1}{n} \sum_{i=1}^n y_i - m \frac{1}{n} \sum_{i=1}^n x_i$$

# Case Study: $f(x) = mx + b$

- We can use  $\bar{y}$  and  $\bar{x}$  to represent the mean values of the  $x$  and  $y$  coordinates, that is all the  $x$  and  $y$  values in our data:

$$\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i \quad \text{and} \quad \bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

- Then we can simply re-write the equation for  $b$ :

$$b = \frac{1}{n} \sum_{i=1}^n y_i - m \frac{1}{n} \sum_{i=1}^n x_i$$

$$b = \bar{y} - m \bar{x}$$

# Case Study: $f(x) = mx + b$

- Now, let's find the value of  $m$ :

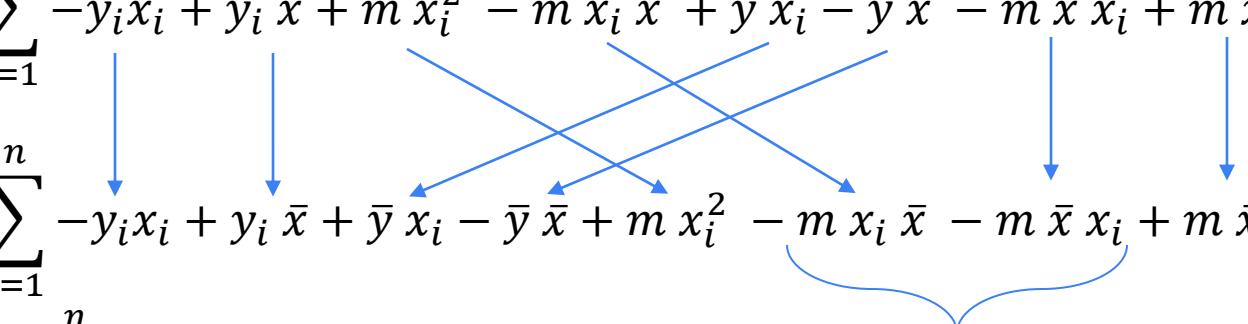
$$\frac{\partial SS(m,b)}{\partial m} = \frac{\partial \sum_{i=1}^n (y_i - mx_i - b)^2}{\partial m} = \frac{\partial \sum_{i=1}^n (y_i - mx_i - \bar{y} + m \bar{x})^2}{\partial m} = 0$$

$$\sum_{i=1}^n 2 (y_i - mx_i - \bar{y} + m \bar{x}) (-x_i + \bar{x}) = 0$$

$$\sum_{i=1}^n -y_i x_i + y_i \bar{x} + m x_i^2 - m x_i \bar{x} + \bar{y} x_i - \bar{y} \bar{x} - m \bar{x} x_i + m \bar{x}^2 = 0$$

# Case Study: $f(x) = mx + b$

$$\sum_{i=1}^n -y_i x_i + y_i \bar{x} + m x_i^2 - m x_i \bar{x} + \bar{y} x_i - \bar{y} \bar{x} - m \bar{x} x_i + m \bar{x}^2 = 0$$



$$\sum_{i=1}^n -y_i x_i + y_i \bar{x} + \bar{y} x_i - \bar{y} \bar{x} + m x_i^2 - m x_i \bar{x} - m \bar{x} x_i + m \bar{x}^2 = 0$$

$$\sum_{i=1}^n -y_i x_i + y_i \bar{x} + \bar{y} x_i - \bar{y} \bar{x} + m x_i^2 - 2m x_i \bar{x} + m \bar{x}^2 = 0$$

# Case Study: $f(x) = mx + b$

$$\sum_{i=1}^n -y_i x_i + y_i \bar{x} + \bar{y} x_i - \bar{y} \bar{x} + m x_i^2 - 2m x_i \bar{x} + m \bar{x}^2 = 0$$

$$\sum_{i=1}^n (\bar{y} - y_i)(x_i - \bar{x}) + m (x_i - \bar{x})^2 = 0$$

Therefore,

$$m = \frac{\sum_{i=1}^n (\bar{y} - y_i)(\bar{x} - x_i)}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

# Case Study: $f(x) = mx + b$

Summary:

$$b = \bar{y} - m \bar{x}$$

$$m = \frac{\sum_{i=1}^n (\bar{y} - y_i)(\bar{x} - x_i)}{\sum_{i=1}^n (x_i - \bar{x})^2}$$

# Linear Regression (cont.)



- Goal: choose the weights  $w$  and bias  $b$  to best fit the true values of  $y$  observed in the data
- When our inputs consist of  $d$  features, we express our prediction  $\hat{y}$  as:

$$\hat{y} = w_1 \cdot x_1 + \dots + w_d \cdot x_d + b$$

- Collecting all features into a vector  $\mathbf{x}$  and all weights into a vector  $\mathbf{w}$ , our model can be expressed compactly using a dot product:

$$\hat{y} = \mathbf{w}^T \mathbf{x} + b$$

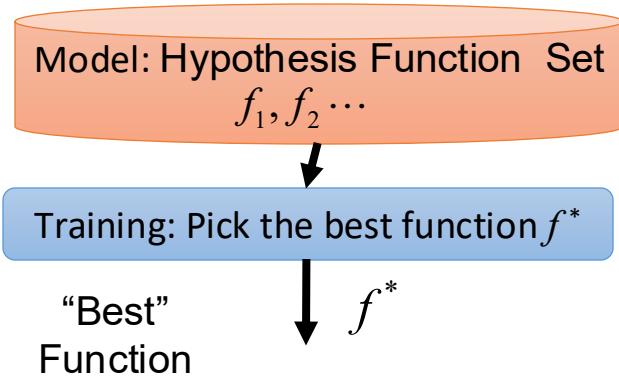
- $\mathbf{x}$ : feature vector for a single data point
- $\mathbf{X}$ : a collection of data points
- $\hat{\mathbf{y}}$ : vector of estimated predictions

$$\hat{\mathbf{y}} = \mathbf{X}\mathbf{w} + b$$

# Training Procedure: Searching for the best parameters

$$\hat{y} = \mathbf{X}\mathbf{w} + b$$

Parameters: elements of  $\mathbf{w}$  and  $b$

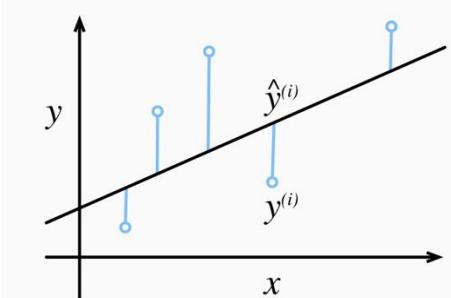


- Q1. What is the model? (function hypothesis set)
- Q2. What does a “good” function mean?
- Q3. How do we find the “best” function?

# Loss Function

- Measure of *fitness*, quantifies the distance between  $y$  and  $\hat{y}$ .
  - Sum of squared errors:
- $$l^{(i)}(\mathbf{w}, b) = \frac{1}{2} (y^{(i)} - \hat{y}^{(i)})^2$$
- To measure quality over the entire training set, we average over all  $K$  training examples:

$$\begin{aligned} L(\mathbf{w}, b) &= \frac{1}{K} \sum_{i=1}^K l^{(i)}(\mathbf{w}, b) \\ &= \frac{1}{K} \sum_{i=1}^K \frac{1}{2} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)})^2 \end{aligned}$$



Residual = truth - predicted

# Training Procedure (cont.)

- Training the model  $\approx$  search for parameters  $(\mathbf{w}^*, b^*)$  that minimize the total loss across all training samples:

$$\mathbf{w}^*, b^* = \operatorname{argmin}_{\mathbf{w}, b} L(\mathbf{w}, b)$$

# Simplification in Notation

- $\hat{y} = \mathbf{X}\mathbf{w} + b$
- Append 1 to features and bias to the weights:  $\mathbf{X} \rightarrow [\mathbf{X} \ 1]$  and  $\mathbf{w} \rightarrow \begin{pmatrix} \mathbf{w} \\ b \end{pmatrix}$

$$\left( \begin{array}{c}
 \left. \begin{array}{cccc}
 x_1^{(1)} & x_2^{(1)} & \dots & x_n^{(1)} \\
 x_1^{(2)} & x_2^{(2)} & \dots & x_n^{(2)} \\
 \dots & & & \\
 x_1^{(K)} & x_2^{(K)} & \dots & x_n^{(K)}
 \end{array} \right\} \\
 \end{array} \right) \left( \begin{array}{c}
 w_1 \\
 w_2 \\
 \dots \\
 w_n
 \end{array} \right) + b = \left( \begin{array}{c}
 \left. \begin{array}{cccc}
 x_1^{(1)} & x_2^{(1)} & \dots & x_n^{(1)} & 1 \\
 x_1^{(2)} & x_2^{(2)} & \dots & x_n^{(2)} & 1 \\
 \dots & & & \\
 x_1^{(K)} & x_2^{(K)} & \dots & x_n^{(K)} & 1
 \end{array} \right\} \\
 \end{array} \right) \left( \begin{array}{c}
 w_1 \\
 w_2 \\
 \dots \\
 w_n \\
 b
 \end{array} \right) = \mathbf{X} \mathbf{w}$$

# Solution for Linear Regression

$$L(\mathbf{X}, \mathbf{y}, \mathbf{w}) = \frac{1}{K} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2$$

$$\frac{d}{d\mathbf{w}} L(\mathbf{X}, \mathbf{y}, \mathbf{w}) = \frac{2}{K} (\mathbf{y} - \mathbf{X}\mathbf{w})^\top \mathbf{X}$$

Loss is convex, hence the optimum solution is at:  $\frac{d}{d\mathbf{w}} L(\mathbf{X}, \mathbf{y}, \mathbf{w}) = 0$

Therefore,

$$\frac{2}{K} (\mathbf{y} - \mathbf{X}\mathbf{w}^*)^\top \mathbf{X} = 0 \quad \text{and} \quad \mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

# Final Project Signup Sheet



[https://docs.google.com/spreadsheets/d/1EJ\\_5Xby0mRhHFmSRSmxlv6Gws4Qs5T8P5JUiKYKAZcA/edit?usp=sharing](https://docs.google.com/spreadsheets/d/1EJ_5Xby0mRhHFmSRSmxlv6Gws4Qs5T8P5JUiKYKAZcA/edit?usp=sharing)

(Everyone should sign up by: Sept 23rd)

You can start a new team or join an existing one. Teams should be 4-6 people!

CS 546 - Fall 2025 - Project Team SignUp Sheet							
	A	B	C	D	E	F	G
1	Team ID	Project Title (can change later)	Brief Description	Team Lead (Name and email)	Team Members (Names and Emails)	TA (or ConvAI PhD student) Mentor	Notes
2							
3							
4							

# Final Project Teaming Tip



- Make sure that you can contribute with meaningful work before joining a team!
- Final reports will include a section on who did what...

# Topics for Next Week

## Tuesday:

- Gradient Descent
- Softmax Regression
- Multi-layer Perceptron

## Thursday:

- Distributional Similarity
- Sparse Word Representations
- Word Embeddings and Word2Vec
- Language Modeling
- Unexpected things we learn with word embeddings