

CS 546 – Advanced Topics in NLP

Dilek Hakkani-Tür



UNIVERSITY OF
ILLINOIS
URBANA-CHAMPAIGN



Siebel School of
Computing
and Data Science



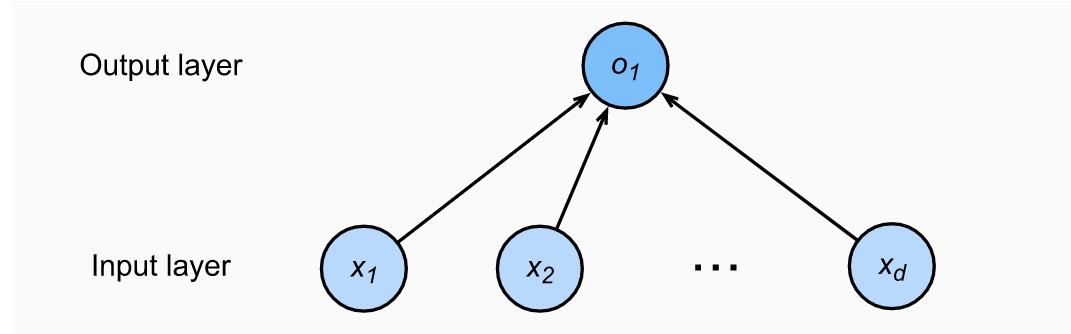
Topics for Today

- Gradient Descent
- Softmax Regression
- Multi-layer Perceptron

From Linear Regression to Neural Networks



- Linear regression is a single-layer neural network, consisting of just a single neuron!
- Fully connected layer (also called *dense* layer): every input is connected to every output.



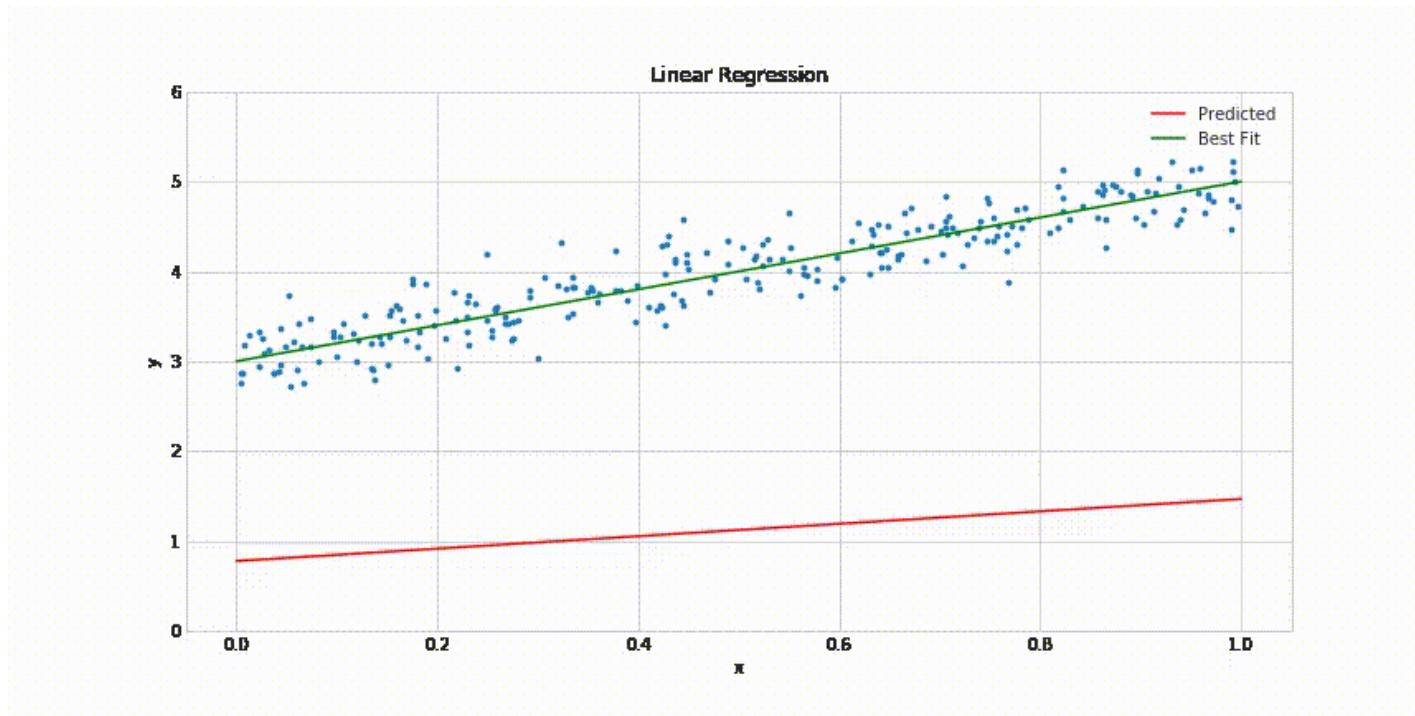
Gradient Descent



- Seen the solution to linear regression last Thursday!
- Even when we cannot solve the models analytically, we can still train models effectively by **iteratively reducing the error by updating the parameters in the direction that incrementally lowers the loss function:**
 - Derivative of the true loss (i.e., average of the losses computed over all examples in the training data).
 - Expressing the updates mathematically:

$$(\mathbf{w}, b) \leftarrow (\mathbf{w}, b) - \frac{\eta}{|K|} \sum_{i \in K} \partial_{(\mathbf{w}, b)} L^{(i)}(\mathbf{w}, b)$$

Case Study: $f(x) = mx + b$



Example from the web

Gradient Descent (cont.)

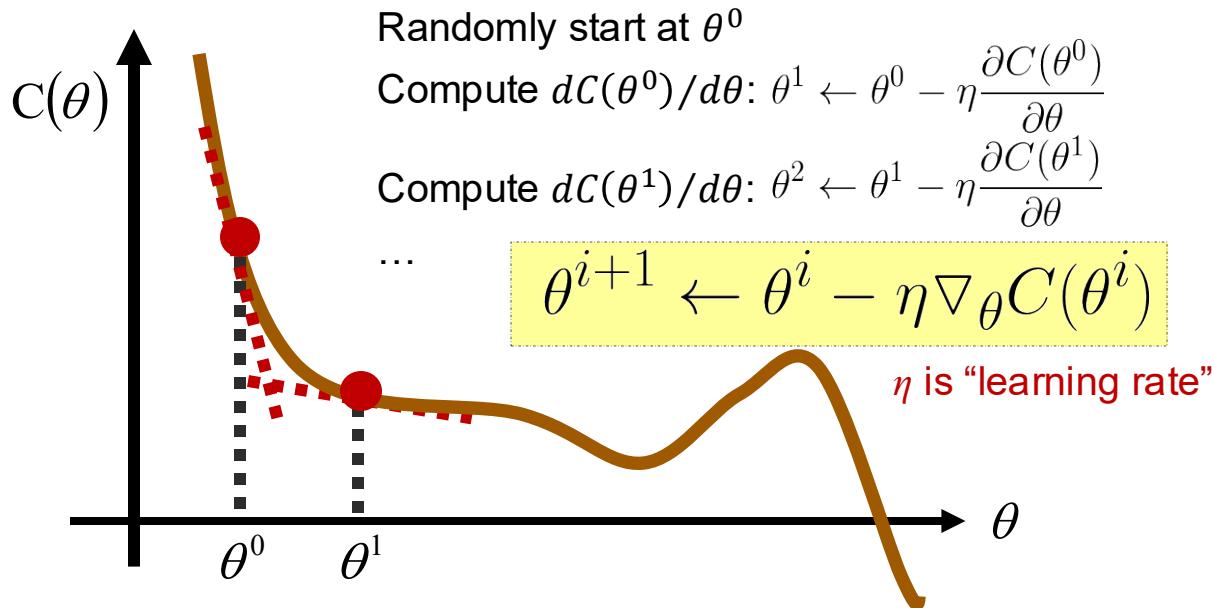
- Steps of the algorithm:
 1. initialize the values of the model parameters, typically at random
 2. iteratively update the parameters in the direction of the negative gradient.

Notation:

θ : model parameters (\mathbf{w}, b)
 $C(\theta)$: $L(\mathbf{w}, b)$

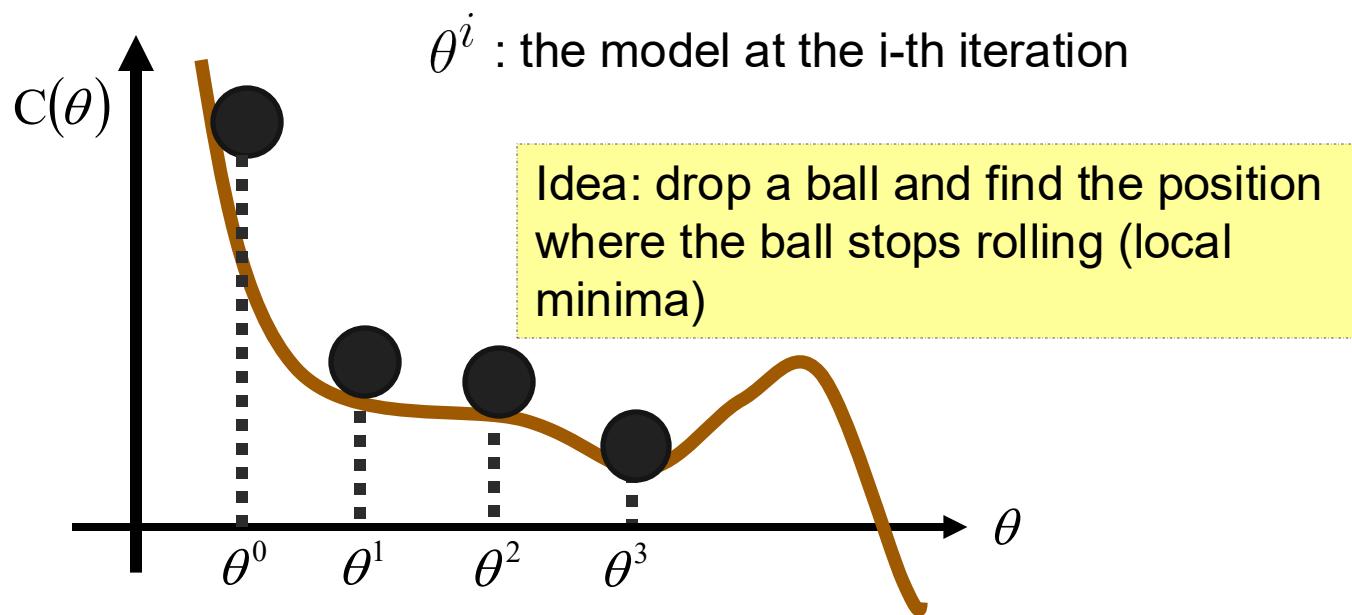
Gradient Descent (cont.)

- Assume that θ has only one variable



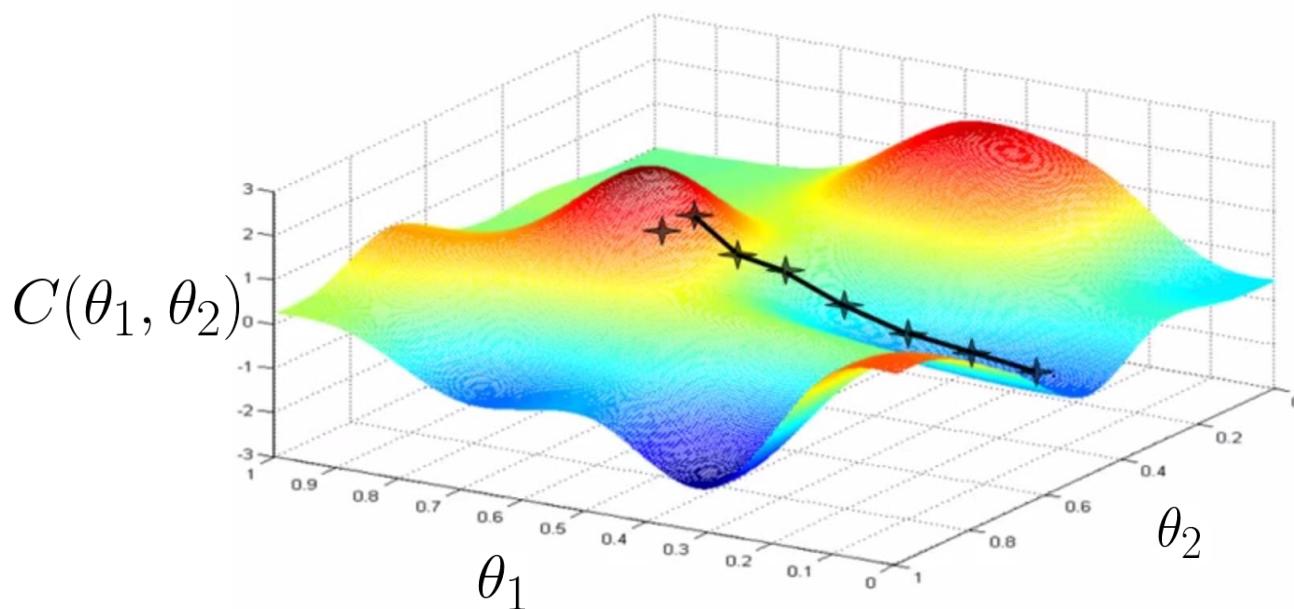
Gradient Descent (cont.)

- Assume that θ has only one variable



Gradient Descent (cont.)

- Assume that θ has two variables $\{\theta_1, \theta_2\}$



Gradient Descent (cont.)

- Assume that θ has two variables $\{\theta_1, \theta_2\}$

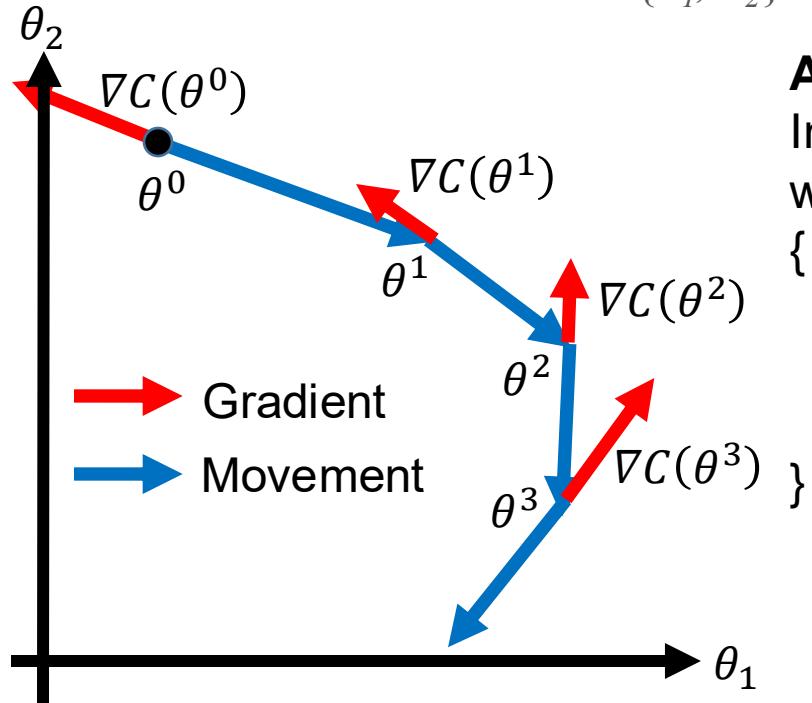
- Randomly start at θ^0 : $\theta^0 = \begin{bmatrix} \theta_1^0 \\ \theta_2^0 \end{bmatrix}$
- Compute the gradients of $C(\theta)$ at θ^0 : $\nabla_{\theta} C(\theta^0) = \begin{bmatrix} \frac{\partial C(\theta_1^0)}{\partial \theta_1} \\ \frac{\partial C(\theta_2^0)}{\partial \theta_2} \end{bmatrix}$
- Update parameters:

$$\begin{bmatrix} \theta_1^1 \\ \theta_2^1 \end{bmatrix} = \begin{bmatrix} \theta_1^0 \\ \theta_2^0 \end{bmatrix} - \eta \begin{bmatrix} \frac{\partial C(\theta_1^0)}{\partial \theta_1} \\ \frac{\partial C(\theta_2^0)}{\partial \theta_2} \end{bmatrix}$$

$$\theta^{i+1} \leftarrow \theta^i - \eta \nabla_{\theta} C(\theta^i)$$
- Compute the gradients of $C(\theta)$ at θ^1 : $\nabla_{\theta} C(\theta^1) = \begin{bmatrix} \frac{\partial C(\theta_1^1)}{\partial \theta_1} \\ \frac{\partial C(\theta_2^1)}{\partial \theta_2} \end{bmatrix}$
- ...

Gradient Descent (cont.)

- Assume that θ has two variables $\{\theta_1, \theta_2\}$



Algorithm

```

Initialization: start at  $\theta^0$ 
while( $\theta^{(i+1)} \neq \theta^i$ )
{
    compute gradient at  $\theta^i$ 
    update parameters
     $\theta^{i+1} \leftarrow \theta^i - \eta \nabla_{\theta} C(\theta^i)$ 
}
  
```

An Issue with Gradient Descent

$$\theta^{i+1} = \theta^i - \eta \nabla C(\theta^i)$$

Training Data
 $\{(x_1, \hat{y}_1), (x_2, \hat{y}_2), \dots\}$

$$C(\theta) = \frac{1}{K} \sum_k \|f(x_k; \theta) - \hat{y}_k\| = \frac{1}{K} \sum_k C_k(\theta)$$

$$\nabla C(\theta^i) = \frac{1}{K} \sum_k \nabla C_k(\theta^i)$$

After seeing all training samples, the model can be updated → slow

Stochastic Gradient Descent (SGD)

- Gradient Descent

$$\theta^{i+1} = \theta^i - \eta \nabla C(\theta^i) \quad \nabla C(\theta^i) = \frac{1}{K} \sum_k \nabla C_k(\theta^i)$$

- Stochastic Gradient Descent (SGD)

- Pick a training sample x_k

$$\theta^{i+1} = \theta^i - \eta \nabla C_k(\theta^i)$$

- If all training samples have the same probability to be picked, then

$$E[\nabla C_k(\theta^i)] = \frac{1}{K} \sum_k \nabla C_k(\theta^i)$$

Training Data
 $\{(x_1, \hat{y}_1), (x_2, \hat{y}_2), \dots\}$

The model can be updated after seeing one training sample → faster

Stochastic Gradient Descent (SGD) (cont.)



- When running SGD, the model starts θ^0

$$\text{pick } x_1 \quad \theta^1 = \theta^0 - \eta \nabla C_1(\theta^0)$$

$$\text{pick } x_2 \quad \theta^2 = \theta^1 - \eta \nabla C_2(\theta^1)$$

⋮

$$\text{pick } x_k \quad \theta^k = \theta^{k-1} - \eta \nabla C_k(\theta^{k-1})$$

⋮

$$\text{pick } x_K \quad \theta^K = \theta^{K-1} - \eta \nabla C_K(\theta^{K-1}) \quad \rightarrow \text{one epoch}$$

Training Data
 $\{(x_1, \hat{y}_1), (x_2, \hat{y}_2), \dots\}$

see all
training
samples once

$$\text{pick } x_1 \quad \theta^{K+1} = \theta^K - \eta \nabla C_1(\theta^K)$$

Mini-Batch SGD

- Batch Gradient Descent

Use all K samples in each iteration

$$\theta^{i+1} = \theta^i - \eta \frac{1}{K} \sum_k \nabla C_k(\theta^i)$$

- Stochastic Gradient Descent (SGD)

- Pick a training sample x_k

Use 1 sample in each iteration

$$\theta^{i+1} = \theta^i - \eta \nabla C_k(\theta^i)$$

- Mini-Batch SGD

- Pick a set of B training samples as a batch b

Use all B samples in each iteration

B is the “batch size”

$$\theta^{i+1} = \theta^i - \eta \frac{1}{B} \sum_{x_k \in b} \nabla C_k(\theta^i)$$



Topics for Today

- Gradient Descent
- Softmax Regression
- Multi-layer Perceptron

Softmax Regression

- Regression is useful for “how much?” and “how many?” questions, where the target can take a real value.
- Many NLP problems are about “which one?”
 - Sentiment classification: positive, negative(, neutral)
 - Sentence boundary detection from text: is the punctuation mark defining a sentence boundary or not
 - What is the part-of-speech tag of the word “word” in this sentence? NOUN, VERB, etc.?
- We can represent target labels with one-hot encodings as well. For example, for sentiment classification:

$$y \in \{(1,0,0), (0,1,0), (0,0,1)\}.$$



A Network Architecture for Multi-Class Problems



- To estimate the conditional probabilities associated with each class, we need a model with multiple outputs, one per class.
- As many linear functions as we have outputs.
- Example: Assume our input is represented with 4 features, then to compute the logits o_1, o_2, o_3 , for each output, we need:

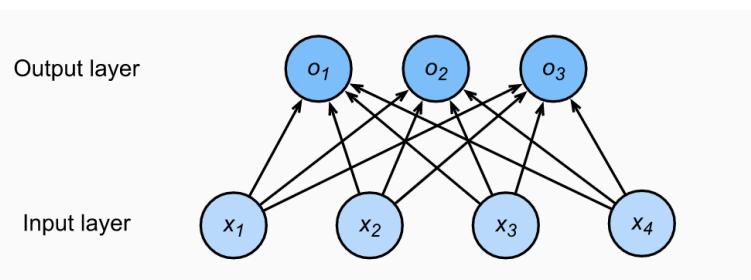
$$o_1 = x_1 w_{11} + x_2 w_{12} + x_3 w_{13} + x_4 w_{14} + b_1$$

$$o_2 = x_1 w_{21} + x_2 w_{22} + x_3 w_{23} + x_4 w_{24} + b_2$$

$$o_3 = x_1 w_{31} + x_2 w_{32} + x_3 w_{33} + x_4 w_{34} + b_3$$

$$\mathbf{o} = \mathbf{Wx} + \mathbf{b}$$

3 x 4 matrix



Multi-Class Problems



- Need to interpret model outputs as probabilities and optimize our parameters to produce probabilities that maximize the likelihood of the observed data
- To generate predictions, we can set a threshold or choose the *argmax* of the predicted probabilities
- Interpret the logits o directly as our outputs of interest, however:
 - Nothing constrains these numbers to sum to 1.
 - Depending on the inputs, they can take negative values.

Softmax



- Transforms logits such that they become nonnegative and sum to 1, while requiring that the model remains differentiable.

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{o}) \quad \text{where} \quad \hat{y}_i = \frac{\exp(o_i)}{\sum_j \exp(o_j)}.$$

- $\hat{y_1} + \hat{y_2} + \hat{y_3} = 1$,
- $0 \leq \hat{y}_i \leq 1$ for all i , and
- The ordering of the logits has not changed; hence we can still pick the output using argmax!

Maximum Likelihood Estimation

- The softmax function gives us a vector \hat{y} , interpreted as estimated conditional probabilities of each class given the input x , e.g., $\hat{y}_1 = P(y=\text{cat}|\mathbf{x})$.

$$P(Y | X) = \prod_{i=1}^n P(y^{(i)} | x^{(i)})$$

$$-\log P(Y | X) = \sum_{i=1}^n -\log P(y^{(i)} | x^{(i)})$$

- Maximizing $P(Y|X)$ (and thus equivalently minimizing $-\log P(Y|X)$) corresponds to predicting the label well!

Cross-Entropy Loss and Softmax

$$l(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{j=1}^q y_j \log \hat{y}_j.$$

q : the number of classes

- If we add o into the definition of the loss l and use the definition of the softmax we obtain:

$$\begin{aligned} l(\mathbf{y}, \hat{\mathbf{y}}) &= - \sum_{j=1}^q y_j \log \frac{\exp(o_j)}{\sum_{k=1}^q \exp(o_k)} \\ &= \sum_{j=1}^q y_j \log \sum_{k=1}^q \exp(o_k) - \sum_{j=1}^q y_j o_j \\ &= \log \sum_{k=1}^q \exp(o_k) - \sum_{j=1}^q y_j o_j. \end{aligned}$$

- If we compute its derivative with respect to o_i , we get:

$$\partial_{o_j} l(\mathbf{y}, \hat{\mathbf{y}}) = \frac{\exp(o_j)}{\sum_{k=1}^q \exp(o_k)} - y_j = \text{softmax}(\mathbf{o})_j - y_j.$$

The gradient is the difference between the probability assigned to the true class by our model. This makes computing gradients very easy in practice!

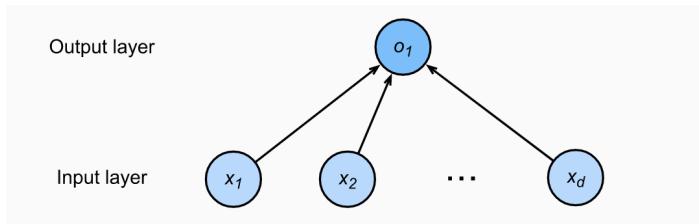
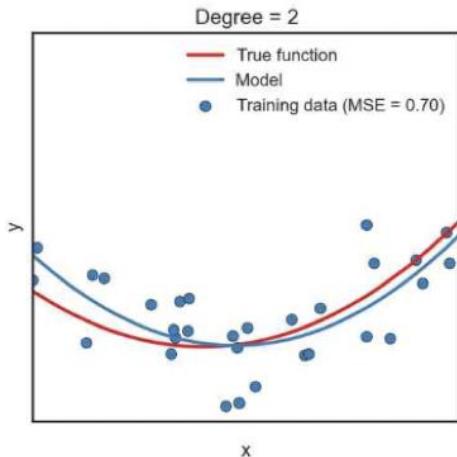
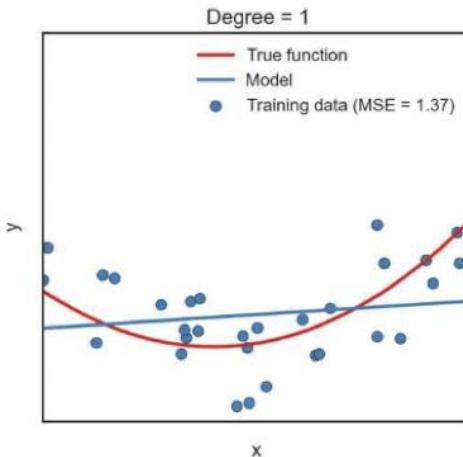


Topics for Today

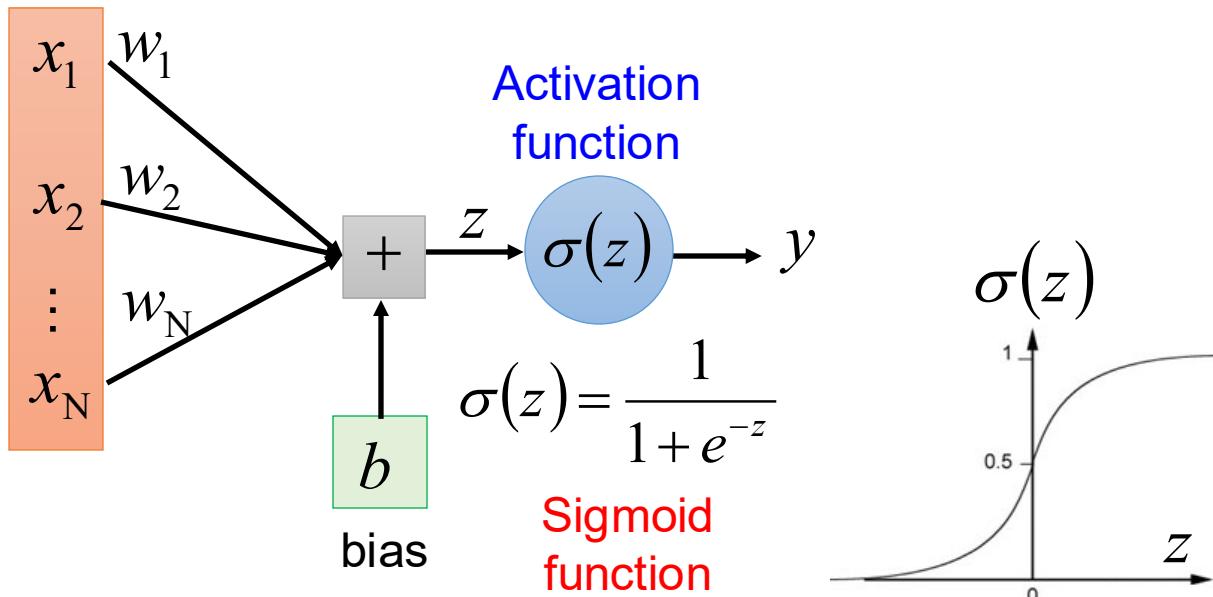
- Gradient Descent
- Softmax Regression
- Multi-layer Perceptron

Non-Linearity

- Linear regression inputs directly to our outputs via a single linear transformation.
- But what if our data is not linear?



Perceptron – A Single Neuron



Each neuron is a very simple function

Why non-linearity?

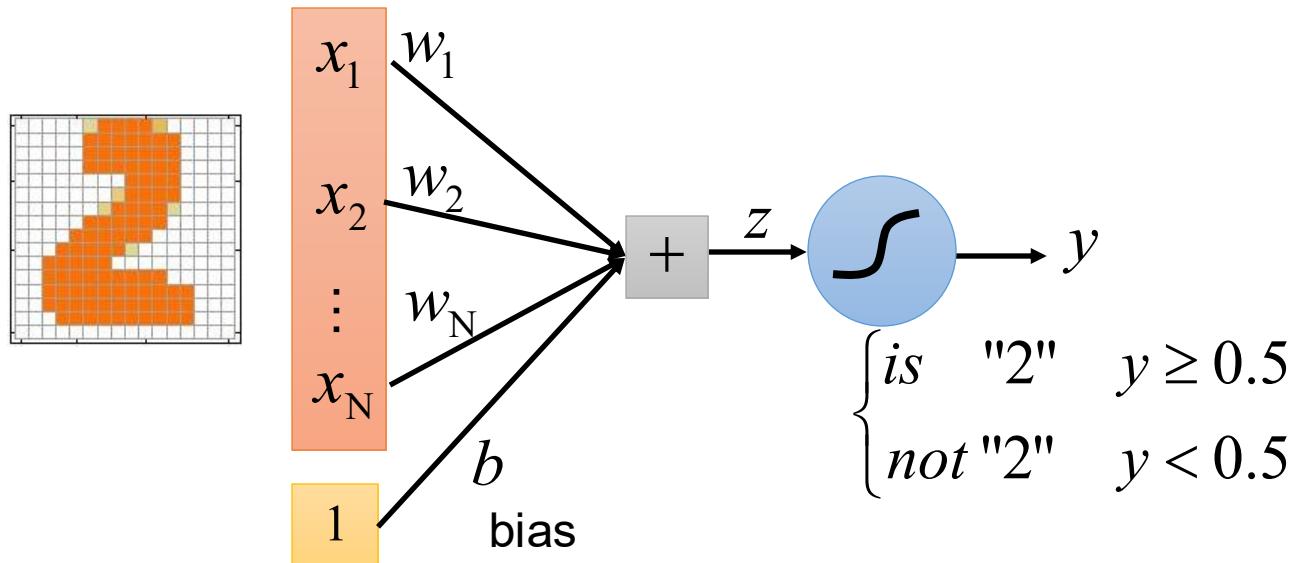
- Function approximation
 - ***Without non-linearity***, deep neural networks work the same as linear transform
$$W_1(W_2 \cdot x) = (W_1 W_2)x = Wx$$
 - ***With non-linearity***, networks with more layers can approximate more complex functions



Figure from: <http://cs224d.stanford.edu/lectures/CS224d-Lecture4.pdf>

Perceptron – A Single Neuron

$$f : R^N \rightarrow R^M$$

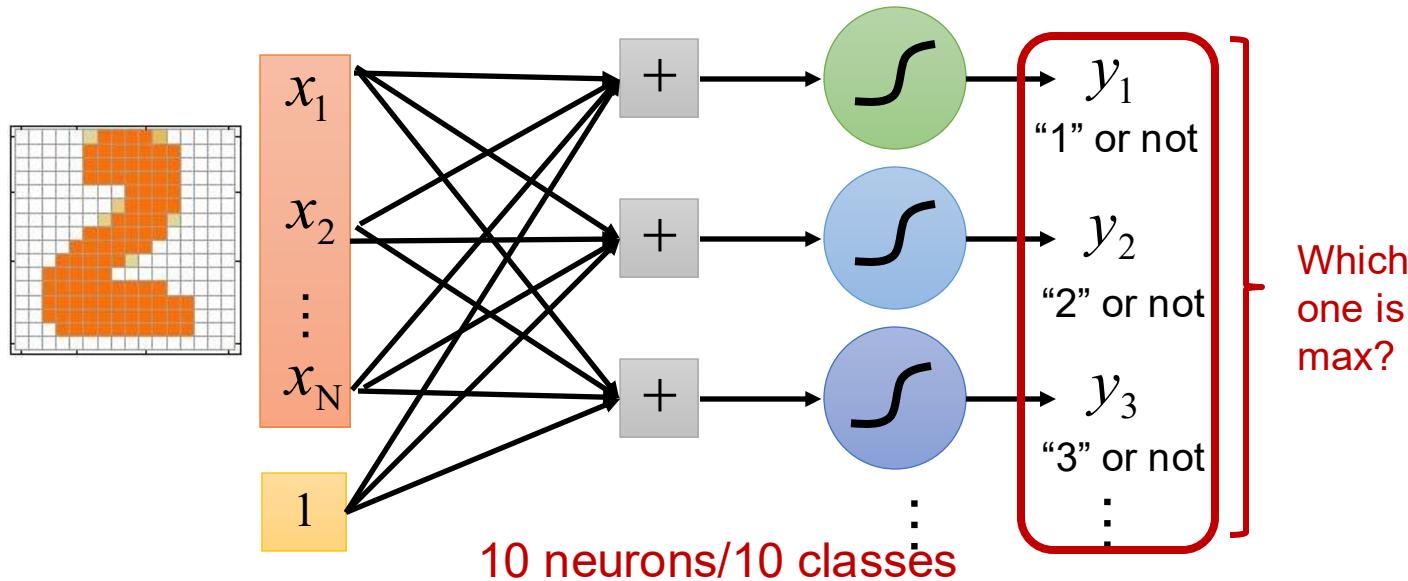


A single neuron can only handle binary classification

A Layer of Neurons

$$f : R^N \rightarrow R^M$$

- Handwriting digit classification

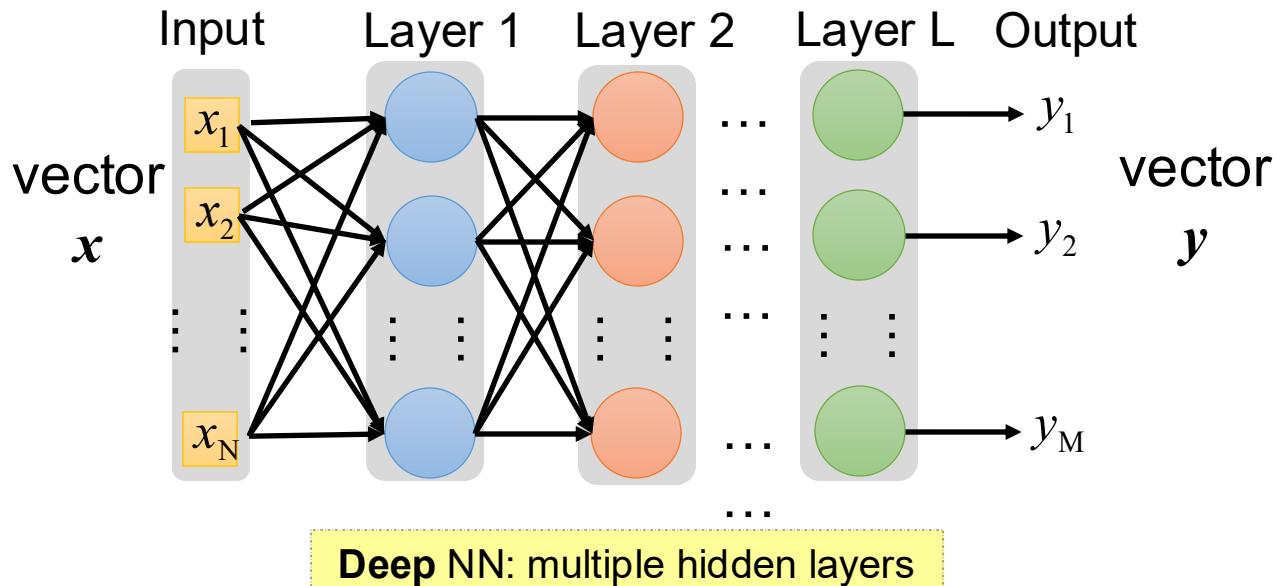


A layer of neurons can handle multiple possible output, and the result depends on the max one

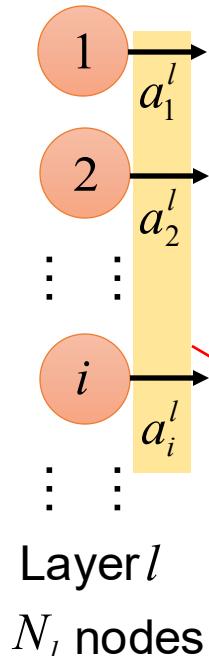
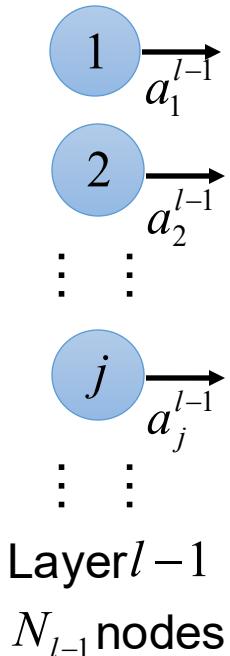
Deep Neural Networks (DNNs)



$$f : R^N \rightarrow R^M$$



DNNs – Notation



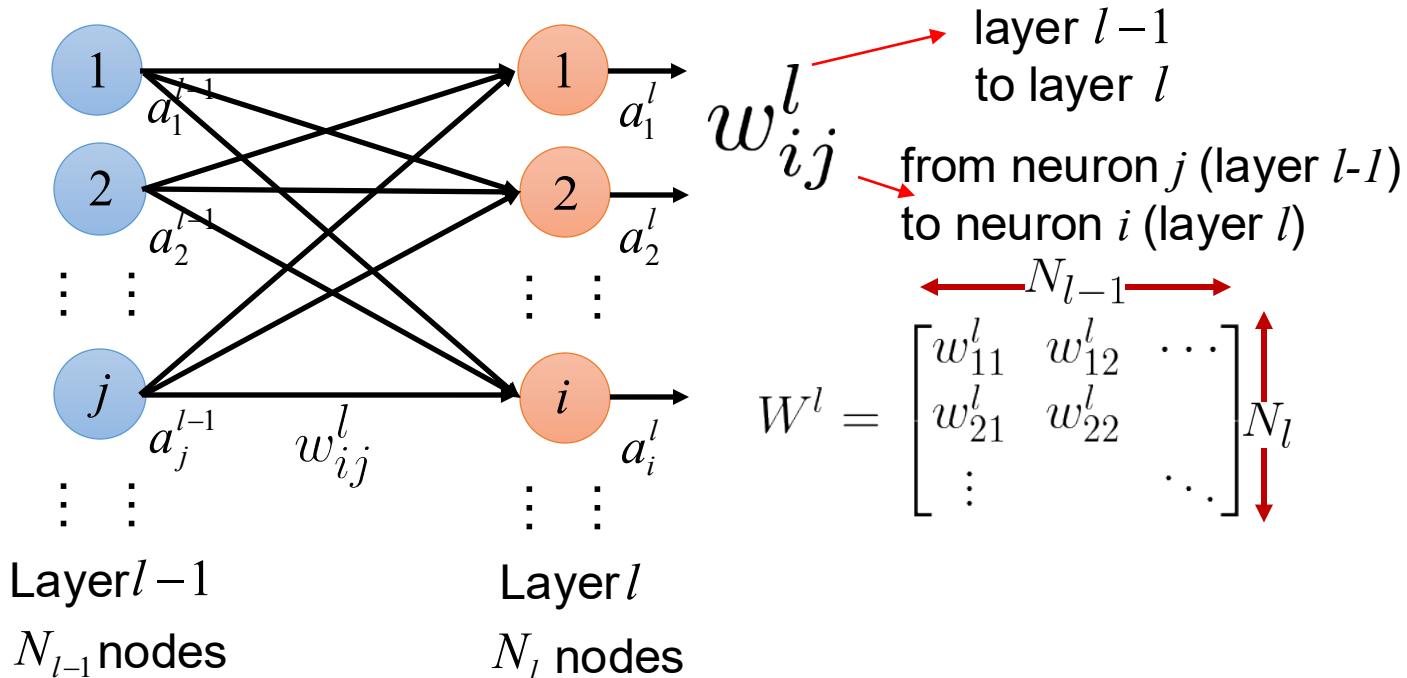
Output of a neuron:

a_i^l → neuron i

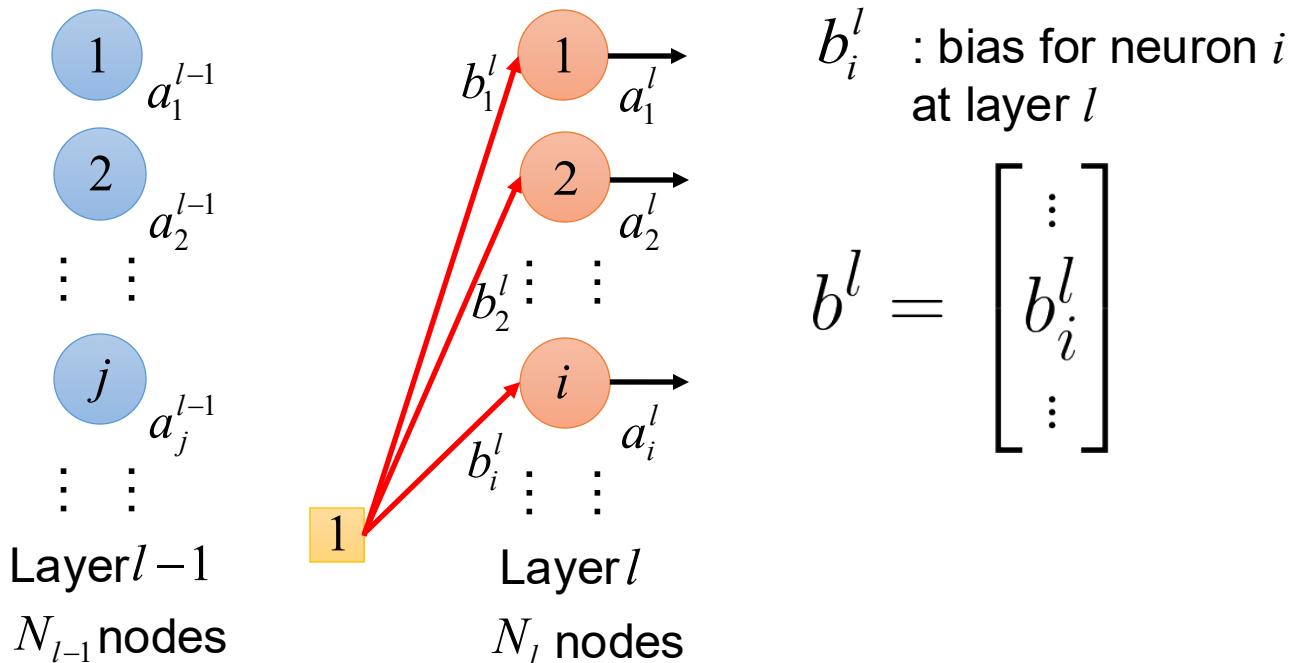
$$a^l = \begin{bmatrix} \vdots \\ a_i^l \\ \vdots \end{bmatrix}$$

output of one layer → a vector

DNNs – Notation (cont.)

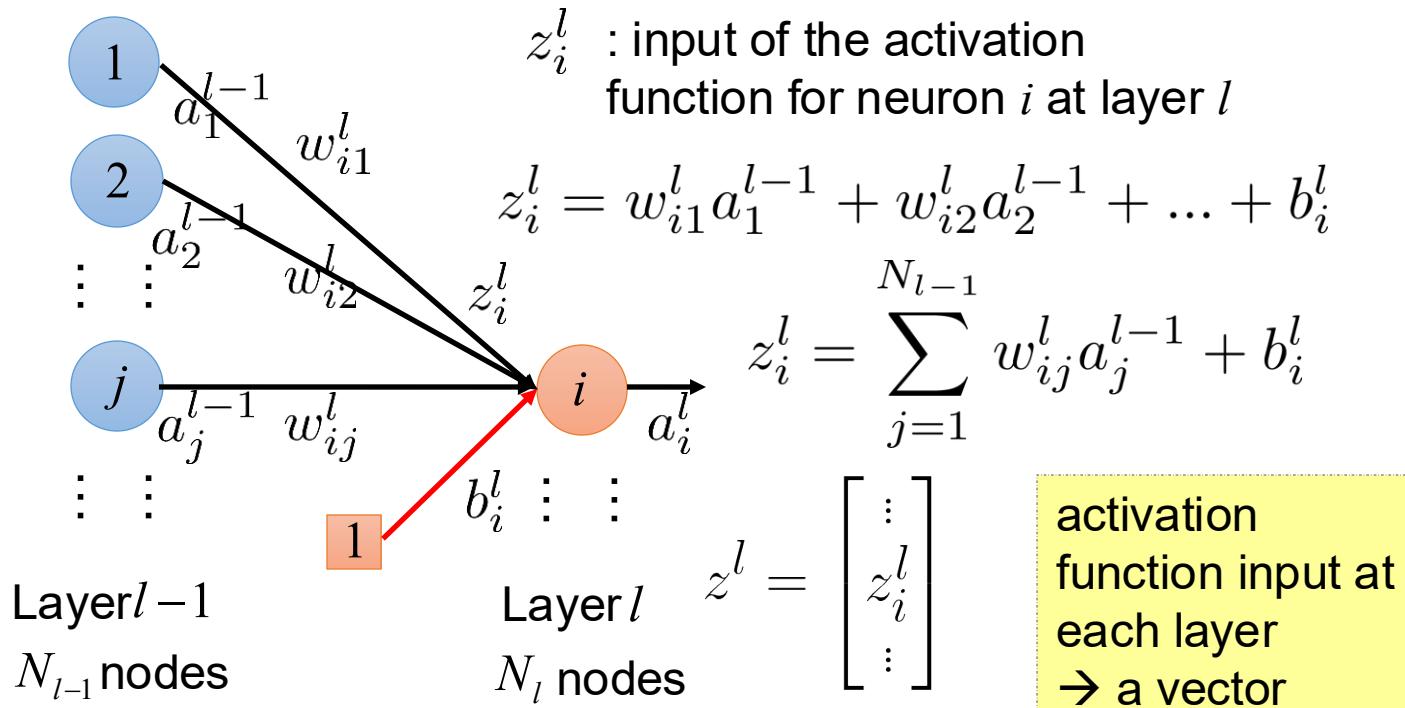


DNNs – Notation (cont.)



bias of all neurons at each layer → a vector

DNNs – Notation (cont.)



DNNs – Notation Summary



a_i^l : output of a neuron

a^l : output vector of a layer

z_i^l : input of activation function

z^l : input vector of activation function for a layer

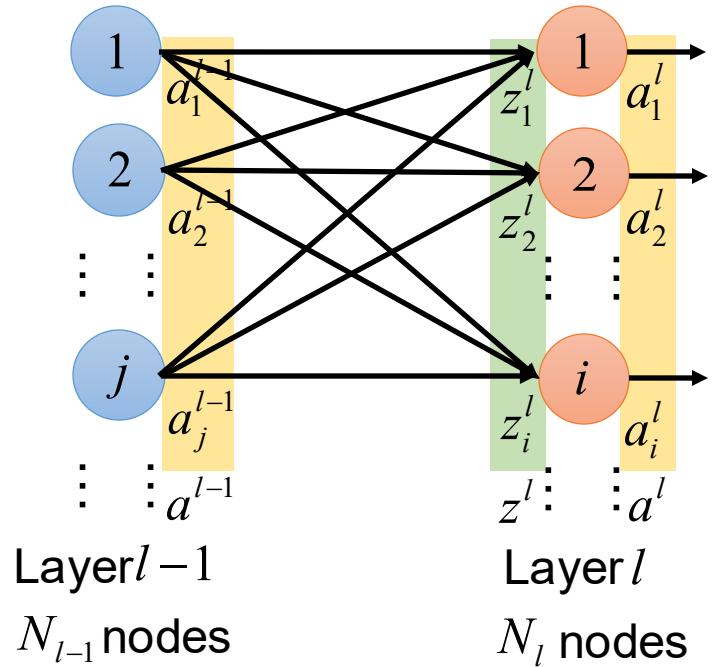
w_{ij}^l : a weight

W^l : a weight matrix

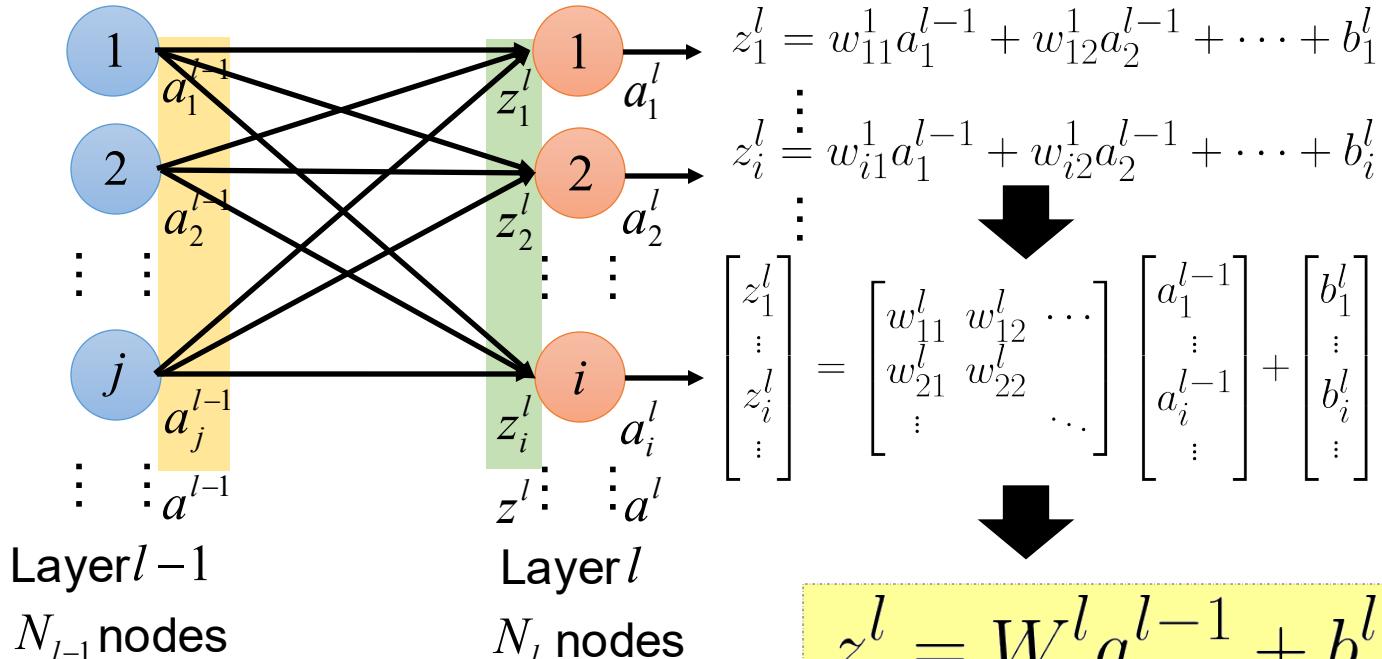
b_i^l : a bias

b^l : a bias vector

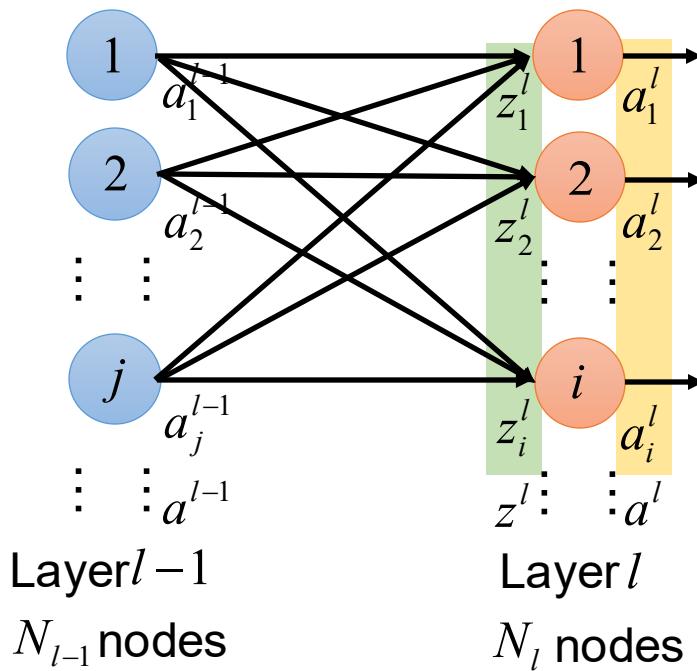
Output of a Layer



Output of a Layer – From a to z



Output of a Layer – From z to a

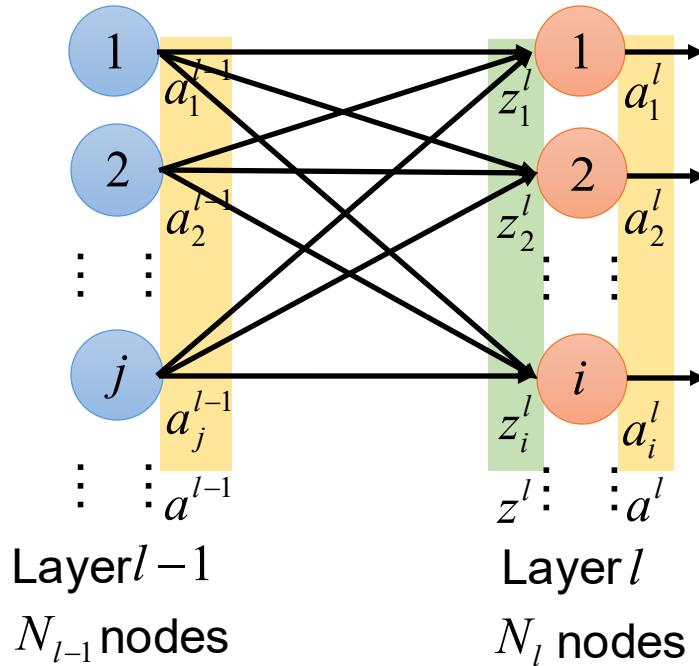


$$a_i^l = \sigma(z_i^l)$$

$$\begin{bmatrix} a_1^l \\ a_2^l \\ \vdots \\ a_i^l \\ \vdots \end{bmatrix} = \begin{bmatrix} \sigma(z_1^l) \\ \sigma(z_2^l) \\ \vdots \\ \sigma(z_i^l) \\ \vdots \end{bmatrix}$$

$$a^l = \sigma(z^l)$$

Output of a Layer



$$z^l = W^l a^{l-1} + b^l$$

$$a^l = \sigma(z^l)$$

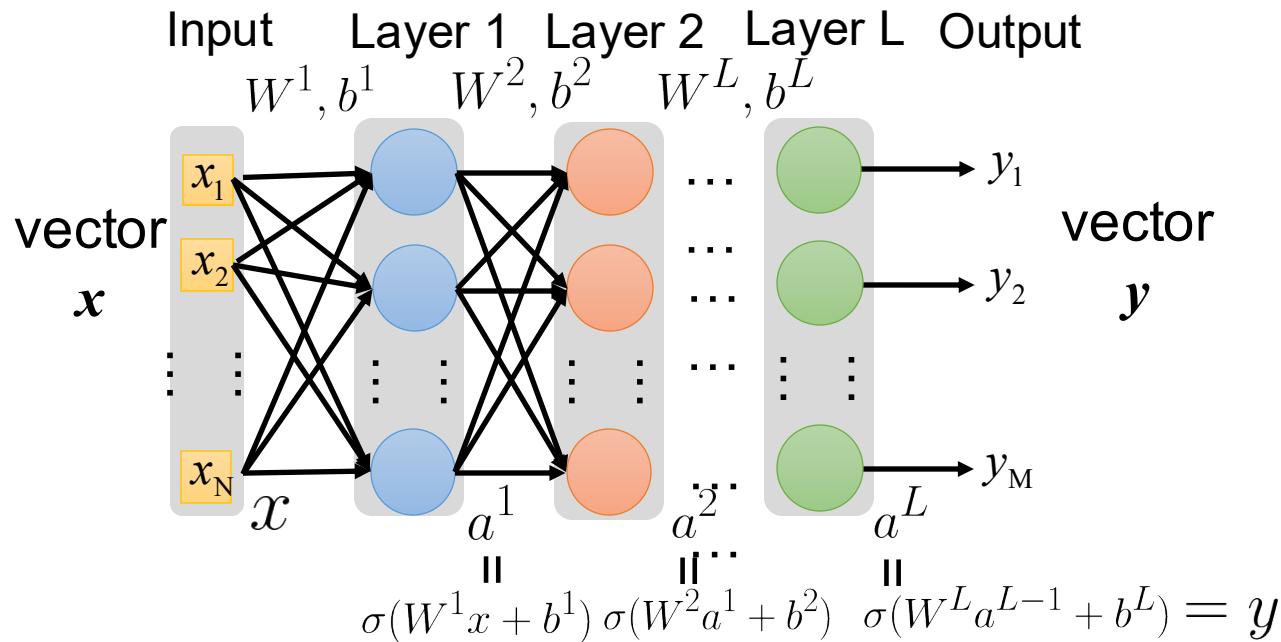


$$a^l = \sigma(W^l a^{l-1} + b^l)$$

DNN Formulation

$$f : R^N \rightarrow R^M$$

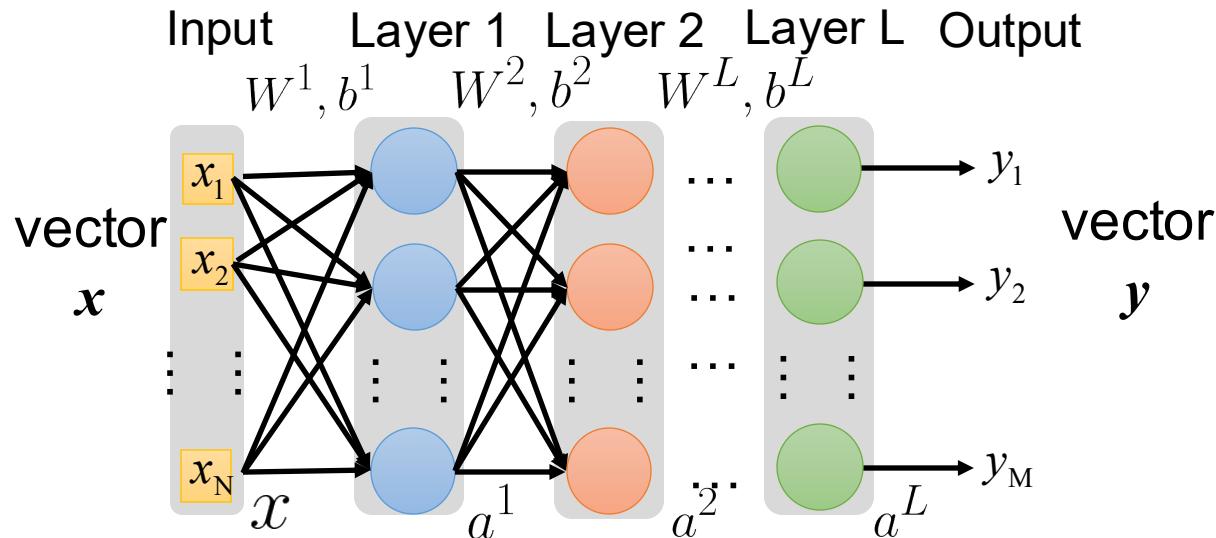
- Fully connected feedforward network



DNN Formulation

$$f : R^N \rightarrow R^M$$

- Fully connected feedforward network



$$y = f(x) = \sigma(W^L \cdots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \cdots + b^L)$$

Loss Function for Training

[Back to an earlier slide!](#)

x : “It makes too much noise.”

function input

\hat{y} : - (negative)
function output

Model: Hypothesis Function Set
 $f_1, f_2 \dots$

Training Data
 $\{(x_1, \hat{y}_1), (x_2, \hat{y}_2), \dots\}$



Training: Pick the best function f^*



“Best” Function f^*

A “Good” function: $f(x; \theta) \sim \hat{y} \rightarrow \|\hat{y} - f(x; \theta)\| \approx 0$

Define an example loss
function:

$$C(\theta) = \sum_k \|\hat{y}_k - f(x_k; \theta)\|$$

sum over the error of all training samples

Gradient Descent for DNNs



$$y = f(x) = \sigma(W^L \cdots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \cdots + b^L)$$

$$\theta = \{W^1, b^1, W^2, b^2, \dots, W^L, b^L\}$$

$$W^l = \begin{bmatrix} w_{11}^l & w_{12}^l & \dots \\ w_{21}^l & w_{22}^l & \dots \\ \vdots & & \dots \end{bmatrix} \quad b^l = \begin{bmatrix} \vdots \\ b_i^l \\ \vdots \end{bmatrix}$$

$$\nabla C(\theta) = \begin{bmatrix} \vdots \\ \frac{\partial C(\theta)}{\partial w_{ij}^l} \\ \vdots \\ \frac{\partial C(\theta)}{\partial b_i^l} \end{bmatrix}$$

Algorithm

Initialization: start at θ^0

while($\theta^{(i+1)} \neq \theta^i$)

{

 compute gradient at θ^i

 update parameters

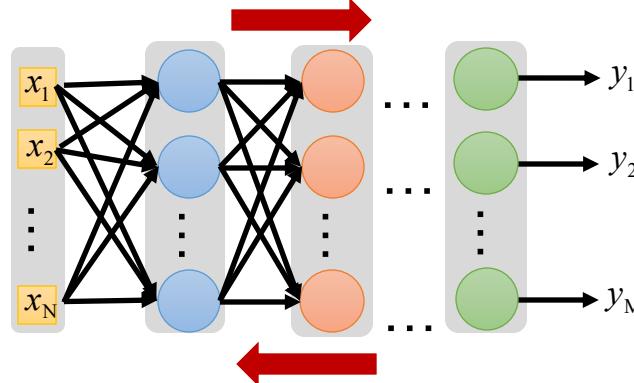
$\theta^{i+1} \leftarrow \theta^i - \eta \nabla_\theta C(\theta^i)$

}

To update weights efficiently, we use **backpropagation**.

Forward vs. Backward Propagation

- In a feedforward neural network
 - forward propagation
 - from input x to output y information flows forward through the network
 - during training, forward propagation can continue onward until it produces a scalar loss $C(\theta)$
 - back-propagation
 - allows the information from the cost to then flow backwards through the network, in order to compute the **gradient**
 - can be applied to any function



Chain Rule

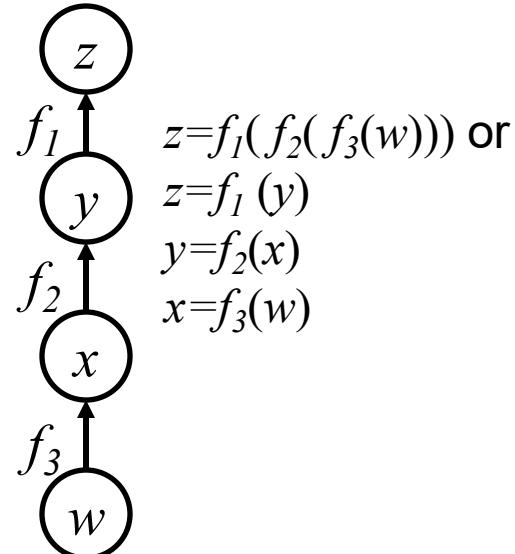
$$\frac{\partial z}{\partial w} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} \frac{\partial x}{\partial w}$$

$$= f'_1(y) f'_2(x) f'_3(w)$$

forward propagation for computing loss

$$= [f'_1(f_2(f_3(w))) | f'_2(f_3(w)) | f'_3(w)]$$

back-propagation of gradients for updating weights



Gradient Descent for DNNs



$$y = f(x) = \sigma(W^L \cdots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \cdots + b^L)$$

$$\theta = \{W^1, b^1, W^2, b^2, \dots, W^L, b^L\}$$

$$W^l = \begin{bmatrix} w_{11}^l & w_{12}^l & \dots \\ w_{21}^l & w_{22}^l & \dots \\ \vdots & & \dots \end{bmatrix} \quad b^l = \begin{bmatrix} \vdots \\ b_i^l \\ \vdots \end{bmatrix}$$

$$\nabla C(\theta) = \begin{bmatrix} \vdots \\ \frac{\partial C(\theta)}{\partial w_{ij}^l} \\ \vdots \\ \frac{\partial C(\theta)}{\partial b_i^l} \end{bmatrix}$$

Algorithm

Initialization: start at θ^0

while($\theta^{(i+1)} \neq \theta^i$)

{

 compute gradient at θ^i

 update parameters

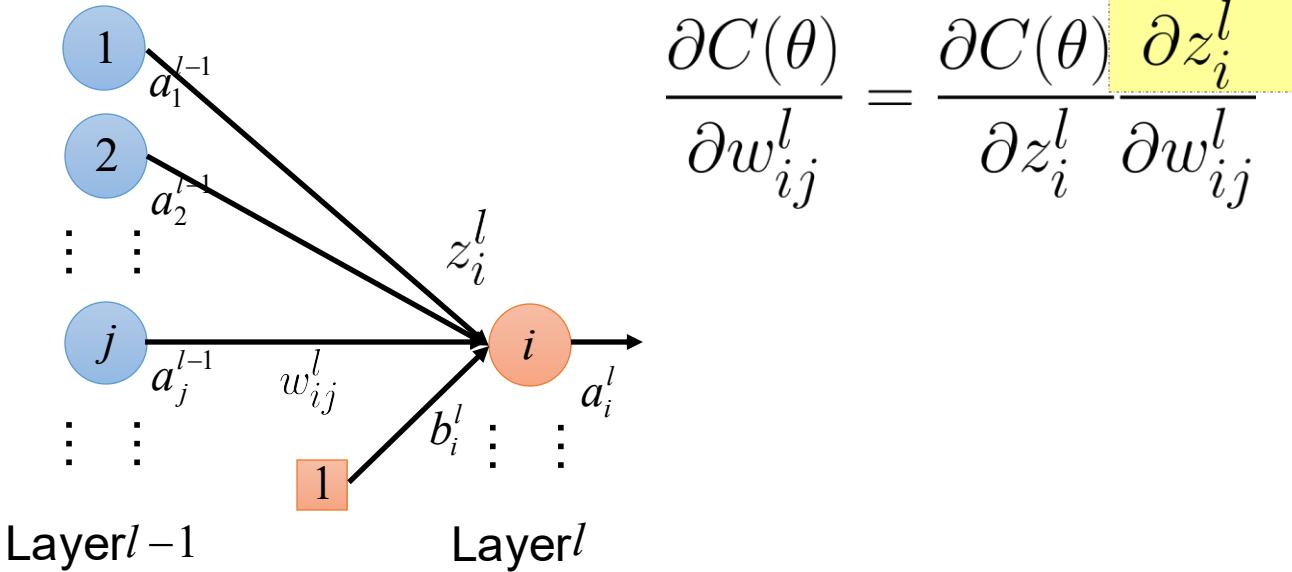
$\theta^{i+1} \leftarrow \theta^i - \eta \nabla_{\theta} C(\theta^i)$

}

To update weights efficiently, we use **backpropagation**.

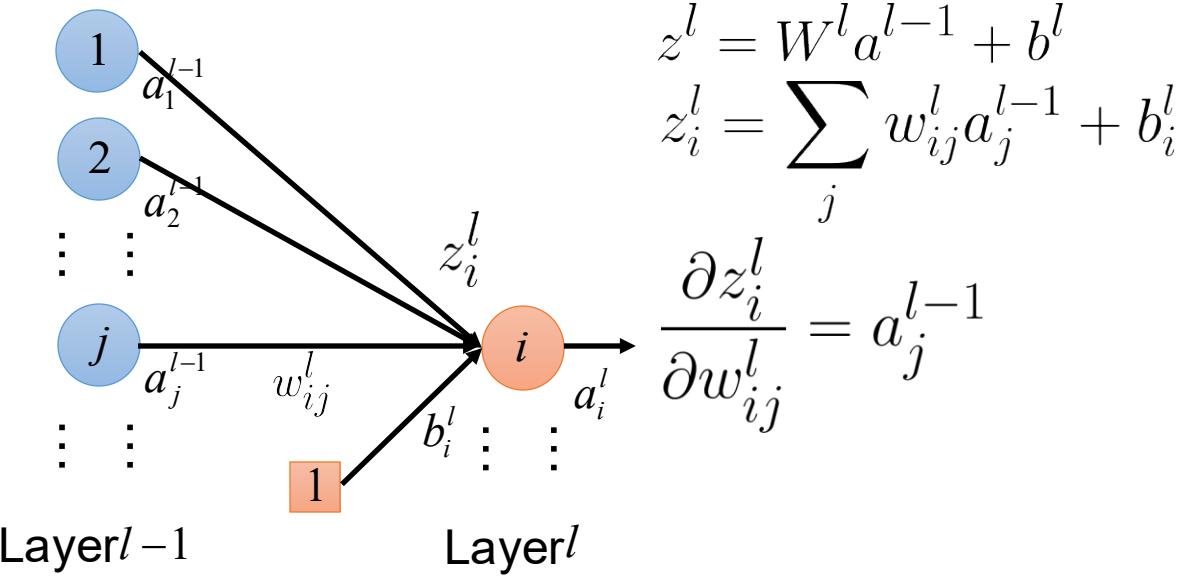
Gradient Descent for DNNs (cont.)

$$\partial C(\theta) / \partial w_{ij}^l$$



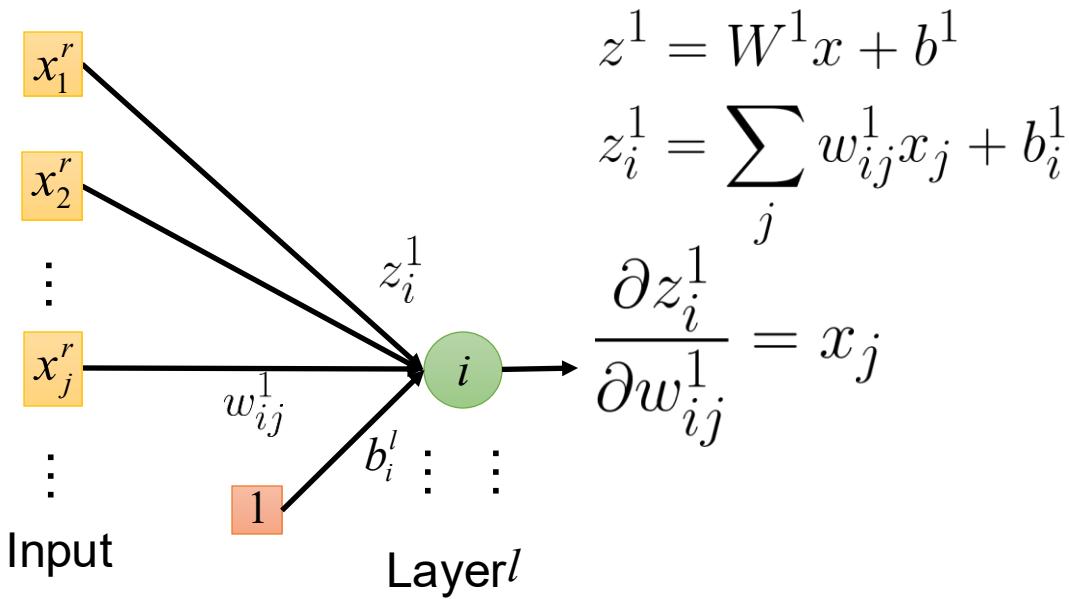
Gradient Descent for DNNs (cont.)

$$\partial z_i^l / \partial w_{ij}^l \quad (l > 1)$$



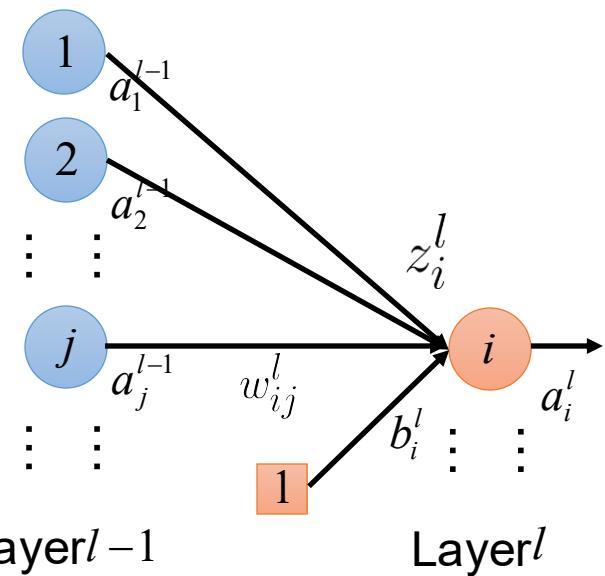
Gradient Descent for DNNs (cont.)

$$\partial z_i^l / \partial w_{ij}^l \quad (l = 1)$$



Gradient Descent for DNNs (cont.)

$$\partial C(\theta) / \partial w_{ij}^l$$

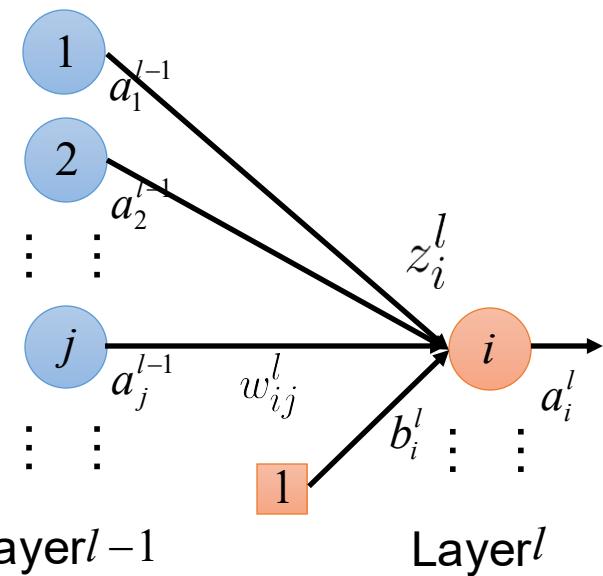


$$\frac{\partial C(\theta)}{\partial w_{ij}^l} = \frac{\partial C(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l}$$

$$\frac{\partial z_i^l}{\partial w_{ij}^l} = \begin{cases} a_j^{l-1} & , l > 1 \\ x_j & , l = 1 \end{cases}$$

Gradient Descent for DNNs (cont.)

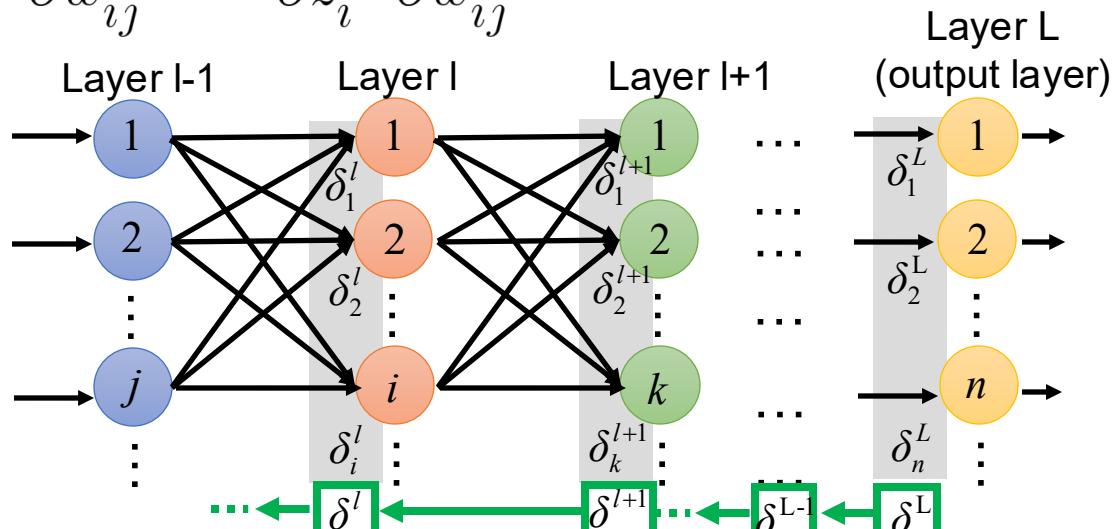
$$\partial C(\theta) / \partial w_{ij}^l$$



$$\frac{\partial C(\theta)}{\partial w_{ij}^l} = \boxed{\frac{\partial C(\theta)}{\partial z_i^l}} \frac{\partial z_i^l}{\partial w_{ij}^l}$$

Gradient Descent for DNNs (cont.)

$$\frac{\partial C(\theta)}{\partial w_{ij}^l} = \frac{\partial C(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l} \quad \delta_i^l : \begin{array}{l} \text{the propagated gradient} \\ \text{corresponding to the } l\text{-th layer} \end{array}$$



Idea: computing δ^l layer by layer (from δ^L to δ^1) is more efficient

Gradient Descent for DNNs (cont.)



$$\partial C(\theta) / \partial z_i^l = \delta_i^l$$

- Idea: from L to 1
- ① Initialization: compute δ^L
- ② Compute δ^l based on δ^{l+1}

Gradient Descent for DNNs (cont.)

$$\partial C(\theta) / \partial z_i^l = \delta_i^l$$

- Idea: from L to 1
- ① Initialization: compute δ^L
- ② Compute δ^l based on δ^{l+1}

$$\begin{aligned}\delta_i^L &= \frac{\partial C}{\partial z_i^L} \quad \Delta z_i^L \rightarrow \Delta a_i^L = \Delta y_i \rightarrow \Delta C \\ &= \boxed{\frac{\partial C}{\partial y_i}} \frac{\partial y_i}{\partial z_i^L}\end{aligned}$$

$\partial C / \partial y_i$ depends on the loss function

Gradient Descent for DNNs (cont.)



$$\partial C(\theta) / \partial z_i^l = \delta_i^l$$

- Idea: from L to 1

- ① Initialization: compute δ^L
- ② Compute δ^l based on δ^{l+1}

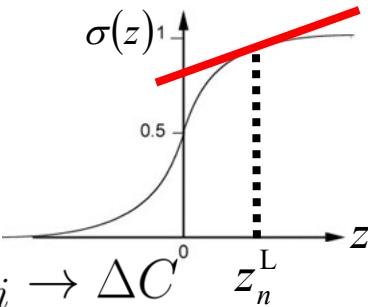
$$\delta_i^L = \frac{\partial C}{\partial z_i^L}$$

$$\Delta z_i^L \rightarrow \Delta a_i^L = \Delta y_i \rightarrow \Delta C^0$$

$$= \frac{\partial C}{\partial y_i} \frac{\partial y_i}{\partial z_i^L} = a_i^L = \sigma(z_i^L)$$

$$= \frac{\partial C}{\partial y_i} \sigma'(z_i^L)$$

$$\boxed{\delta^L = \sigma'(z^L) \odot \nabla C(y)}$$



$$\sigma'(z^L) = \begin{bmatrix} \sigma'(z_1^L) \\ \sigma'(z_2^L) \\ \vdots \\ \sigma'(z_i^L) \\ \vdots \end{bmatrix} \quad \nabla C(y) = \begin{bmatrix} \frac{\partial C}{\partial y_1} \\ \frac{\partial C}{\partial y_2} \\ \vdots \\ \frac{\partial C}{\partial y_i} \\ \vdots \end{bmatrix}$$

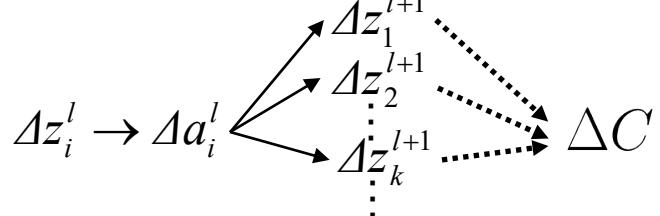
Gradient Descent for DNNs (cont.)

$$\partial C(\theta) / \partial z_i^l = \delta_i^l$$

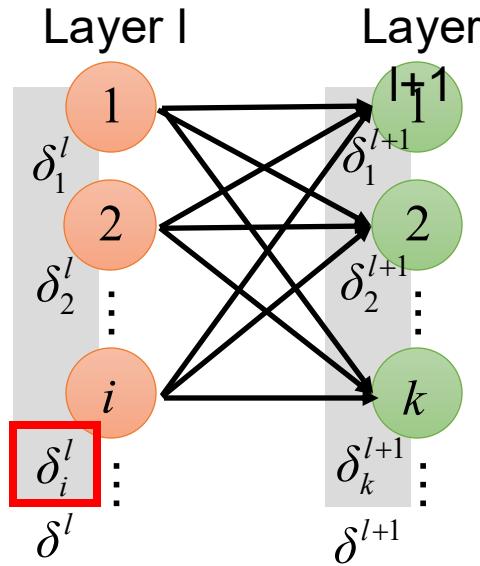
- Idea: from L to 1

① Initialization: compute δ^L

② Compute δ^l based on δ^{l+1}



$$\begin{aligned}\delta_i^l &= \frac{\partial C}{\partial z_i^l} = \sum_k \left(\frac{\partial C}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial a_i^l} \frac{\partial a_i^l}{\partial z_i^l} \right) \\ &= \frac{\partial a_i^l}{\partial z_i^l} \sum_k \left(\frac{\partial C}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial a_i^l} \right) \quad \delta_i^{l+1}\end{aligned}$$

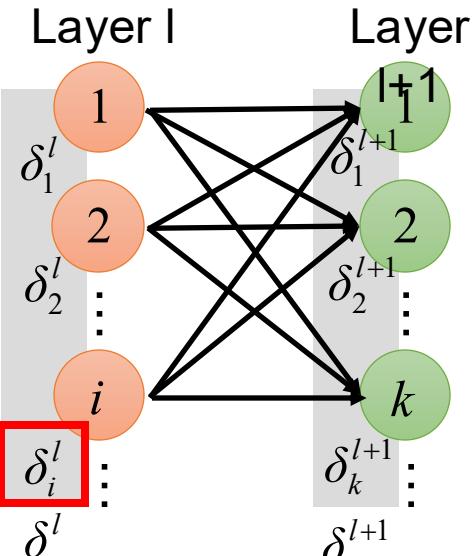
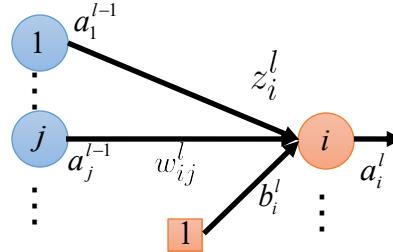


Gradient Descent for DNNs (cont.)

$$\partial C(\theta) / \partial z_i^l = \delta_i^l$$

- Idea: from L to 1
- ① Initialization: compute δ^L
 - ② Compute δ^l based on δ^{l+1}

$$\begin{aligned}\delta_i^l &= \frac{\partial a_i^l}{\partial z_i^l} \sum_k \frac{\partial z_k^{l+1}}{\partial a_i^l} \delta_k^{l+1} = \sum_k w_{ki}^{l+1} a_i^l + b_k^{l+1} \\ &= \sigma'(z_i) \sum_k \frac{\partial z_k^{l+1}}{\partial a_i^l} \delta_k^{l+1} \\ &= \sigma'(z_i) \sum_k w_{ki}^{l+1} \delta_k^{l+1}\end{aligned}$$

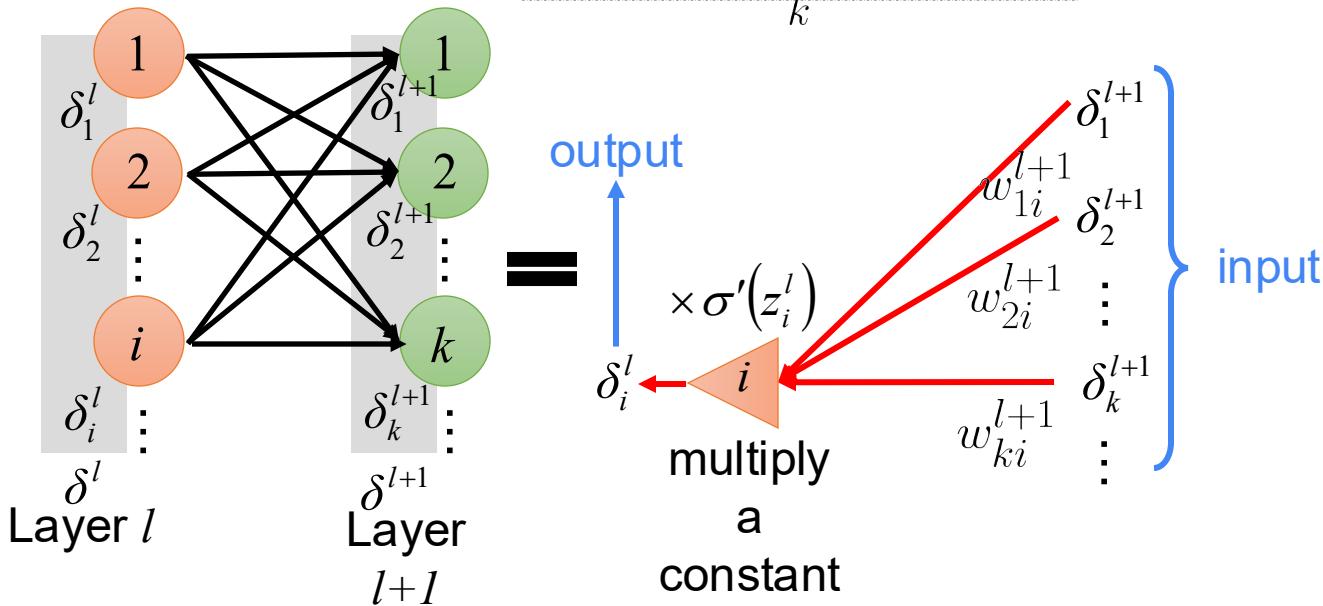


Gradient Descent for DNNs (cont.)

$$\partial C(\theta) / \partial z_i^l = \delta_i^l$$

- Rethink the propagation

$$\delta_i^l = \sigma'(z_i^l) \sum_k w_{ki}^{l+1} \delta_k^{l+1}$$



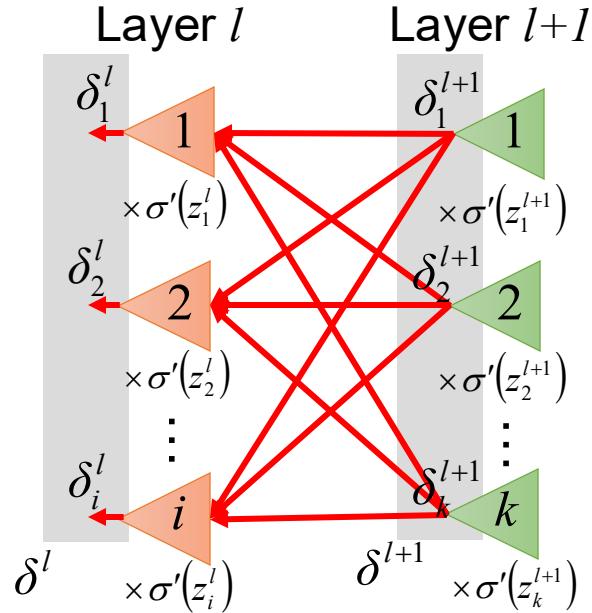
Gradient Descent for DNNs (cont.)

$$\partial C(\theta) / \partial z_i^l = \delta_i^l$$

$$\delta_i^l = \sigma'(z_i) \sum_k w_{ki}^{l+1} \delta_k^{l+1}$$

$$\sigma'(z^l) = \begin{bmatrix} \sigma'(z_1^l) \\ \sigma'(z_2^l) \\ \vdots \\ \sigma'(z_i^l) \\ \vdots \end{bmatrix}$$

$$\delta^l = \sigma'(z^l) \odot (W^{l+1})^T \delta^{l+1}$$



Gradient Descent for DNNs (cont.)

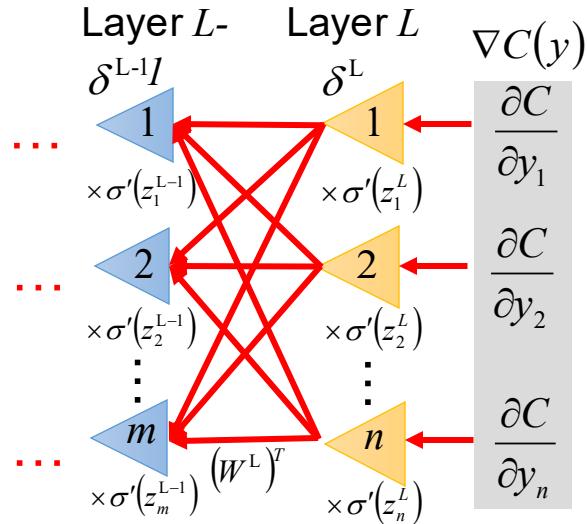
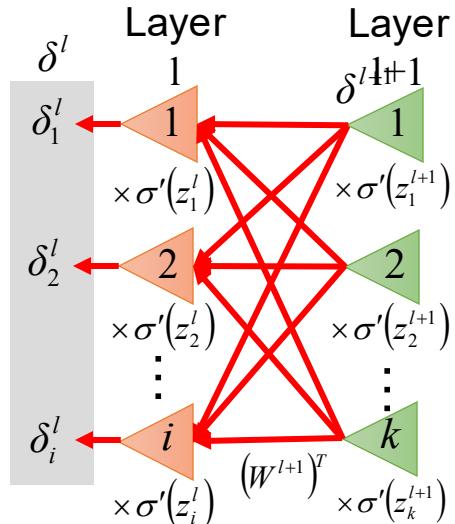
$$\frac{\partial C(\theta)}{\partial z_i^l} = \delta_i^l$$

$$\frac{\partial C(\theta)}{\partial w_{ij}^l} = \boxed{\frac{\partial C(\theta)}{\partial z_i^l}} \frac{\partial z_i^l}{\partial w_{ij}^l}$$

- Idea: from L to 1
- ① Initialization: compute δ^L
 - ② Compute δ^{l-1} based on δ^l

$$\delta^L = \sigma'(z^L) \odot \nabla C(y)$$

$$\delta^l = \sigma'(z^l) \odot (W^{l+1})^T \delta^{l+1}$$



Backpropagation

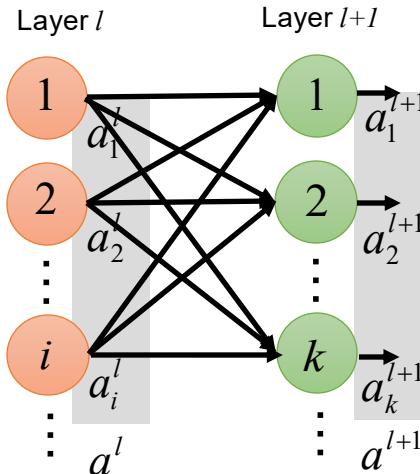
$$\frac{\partial C(\theta)}{\partial w_{ij}^l} = \frac{\partial C(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l}$$

$$\frac{\partial z_i^l}{\partial w_{ij}^l} = \begin{cases} a_j^{l-1}, & l > 1 \\ x_j, & l = 1 \end{cases}$$

Forward Pass

$$z^1 = W^1 x + b^1 \quad a^1 = \sigma(z^1)$$

$$z^l = W^l a^{l-1} + b^l \quad a^l = \sigma(z^l)$$



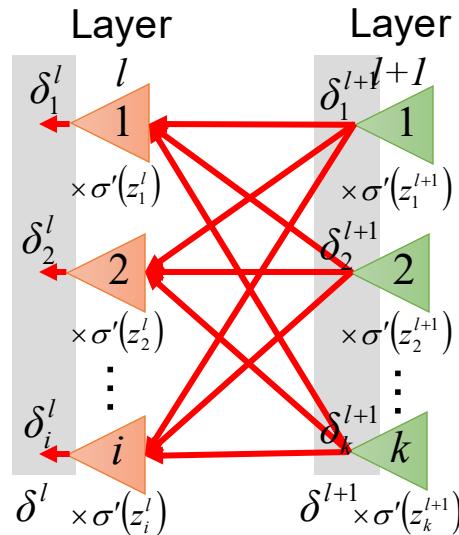
Backpropagation

$$\frac{\partial C(\theta)}{\partial w_{ij}^l} = \frac{\partial C(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l}$$

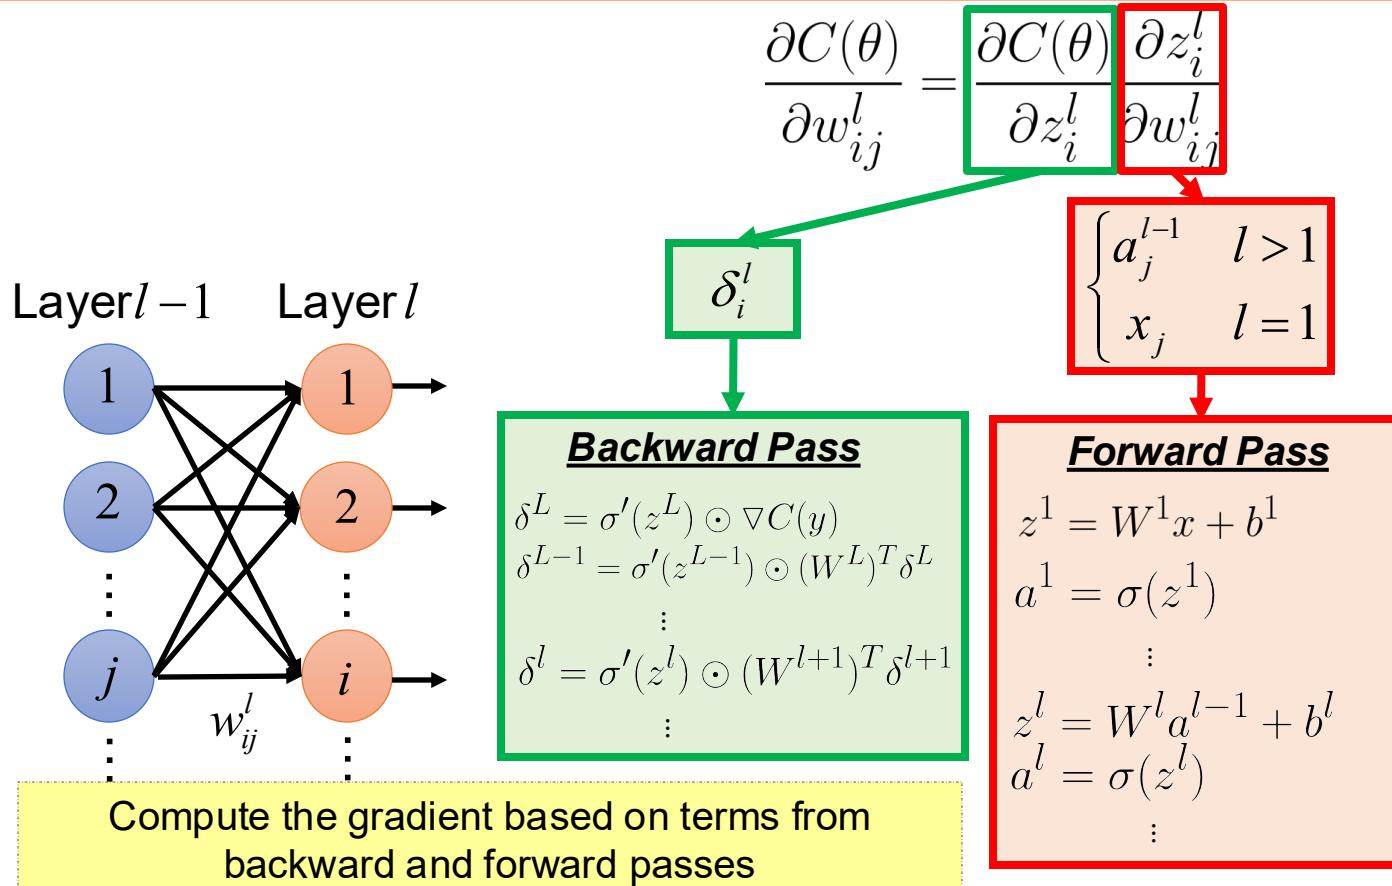
$$\frac{\partial C(\theta)}{\partial z_i^l} = \delta_i^l$$

Backward Pass

$$\begin{aligned}\delta^L &= \sigma'(z^L) \odot \nabla C(y) \\ \delta^{L-1} &= \sigma'(z^{L-1}) \odot (W^L)^T \delta^L \\ &\vdots \\ \delta^l &= \sigma'(z^l) \odot (W^{l+1})^T \delta^{l+1} \\ &\vdots\end{aligned}$$



Gradient Descent - Summary



Topics for Thursday

- Distributional Similarity
- Sparse Word Representations
- Word Embeddings and Word2Vec
- Language Modeling
- Unexpected things we learn with word embeddings