

# CS 546 – Advanced Topics in NLP

Dilek Hakkani-Tür



UNIVERSITY OF  
**ILLINOIS**  
URBANA-CHAMPAIGN



Siebel School of  
Computing  
and Data Science



# Topics for Today

## Sequence Modeling

- Convolutional Neural Networks
- Recurrent Neural Networks
- Example Applications

# Convolutional Neural Networks (CNNs)



- We have discussed models that deal with paired data: input words or utterances and output categories.
- Example: One-hot representations or word embeddings to represent input words
- Sometimes data exhibits rich structure, such as images, natural language.
  - Structure-less networks, i.e., MLPs, can fall short.
- CNNs: a type of NNs well-suited to detecting spatial substructure.

# Why CNNs? Example: Image Recognition



- Some patterns are much smaller than the whole image

A neuron does not have to see the whole image to discover the pattern.

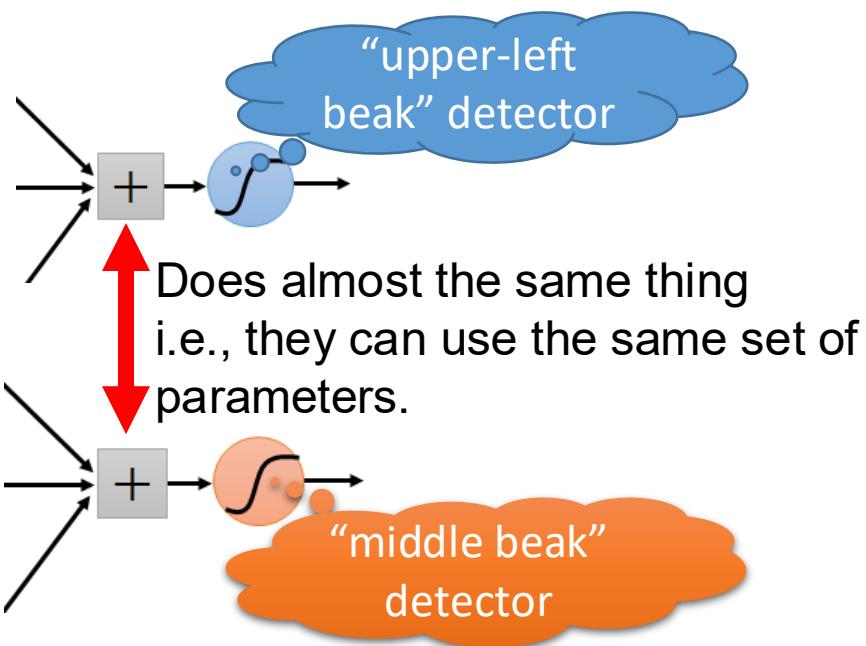
Connecting to small region with fewer parameters



# Why CNNs? Example: Image Recognition



- The same pattern can appear in different regions.



# Why CNNs? Example: Text Classification

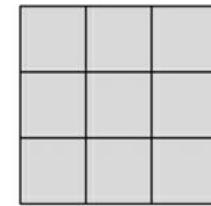
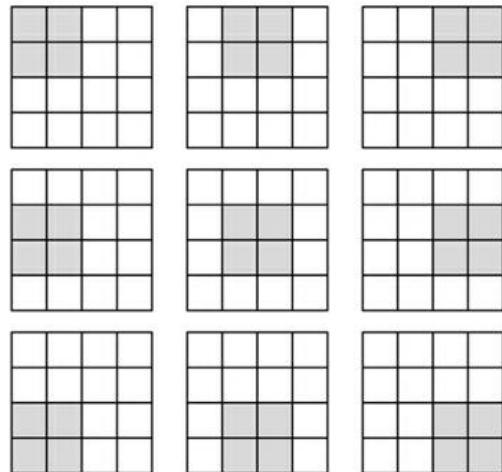
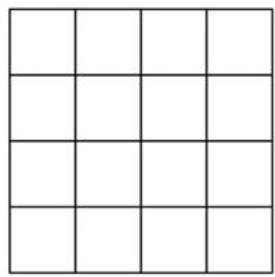


- Recognizing original nationalities from last names  
**O'Neill O'Shaughnessy**  
**Antonopoulos Kostopoulos Giannopoulos**
- Sentiment classification  
A **delicious** breakfast was served to us at Pillerago that morning.  
The brunch at Margoli, especially the scones, were **delicious**.

# Where's Waldo?



# Convolutional Layers

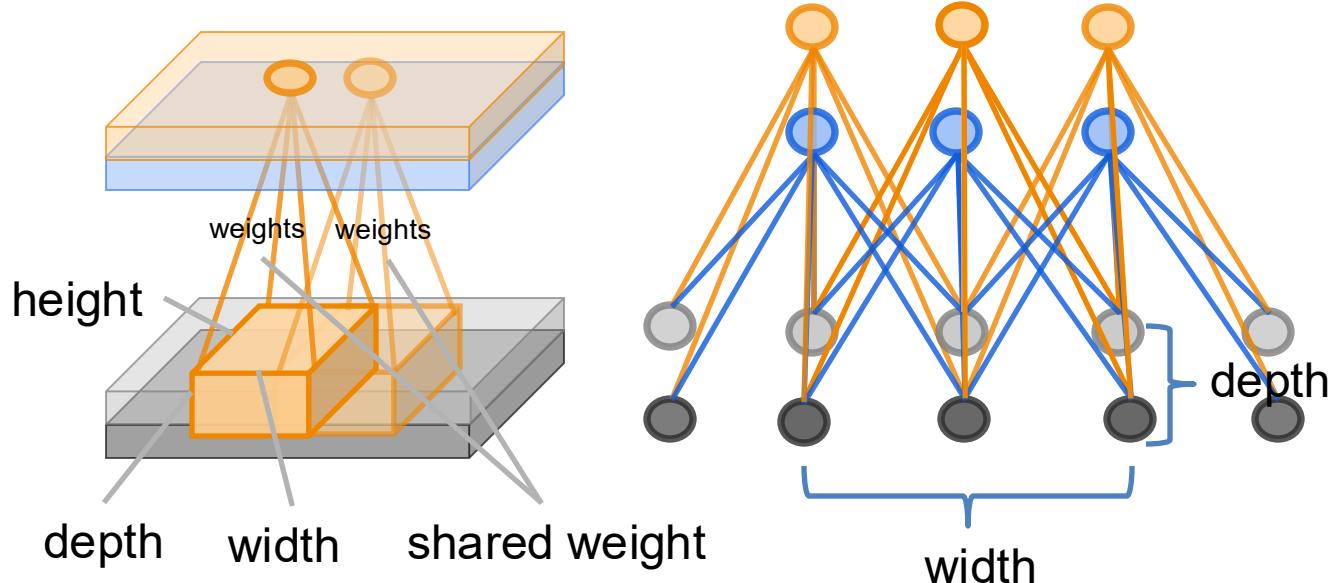


An input matrix  
and a convolutional kernel

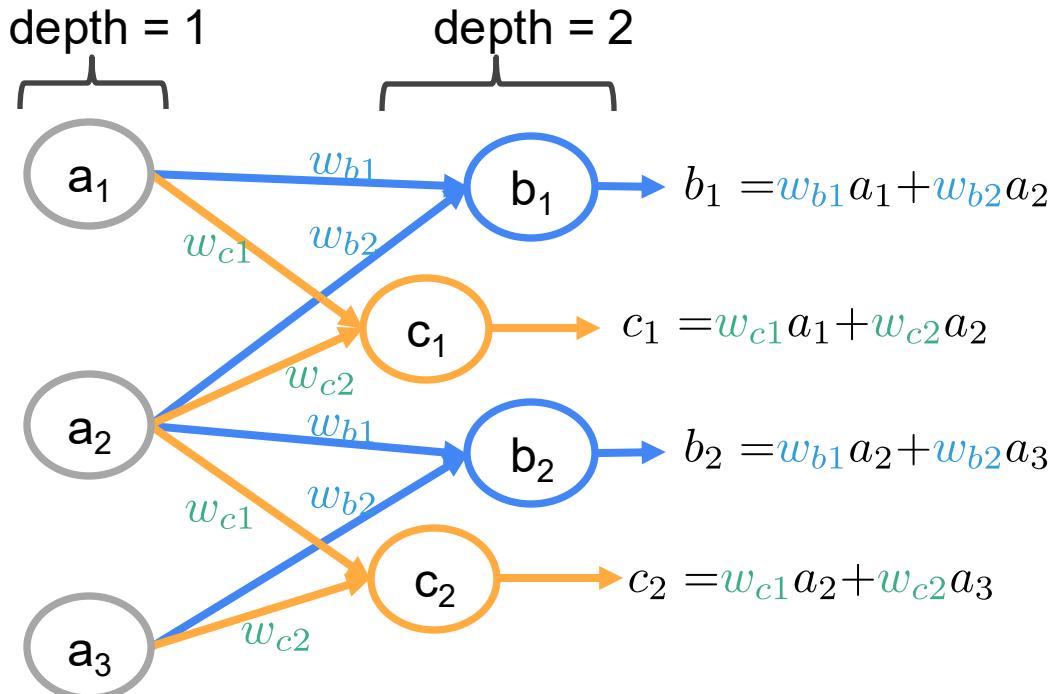
The kernel being applied to the  
input matrix

The output matrix  
(feature map)

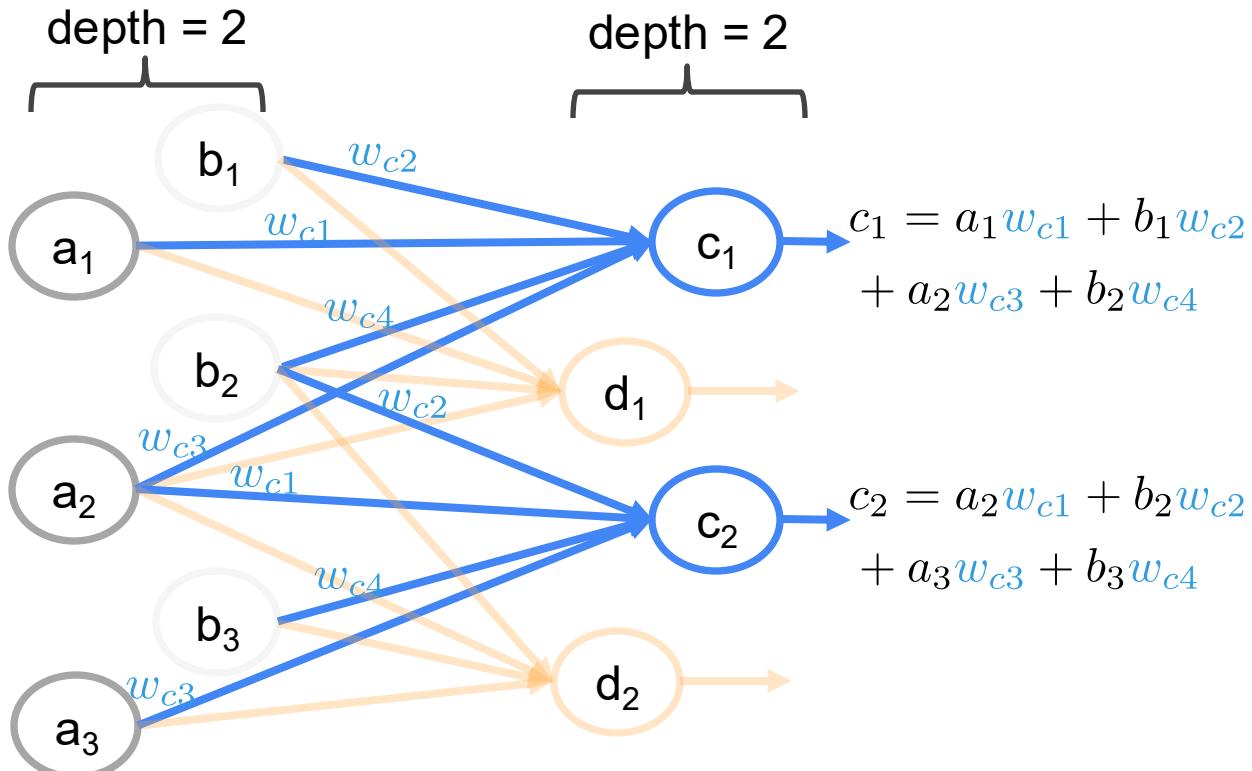
# Convolutional Layers (cont.)



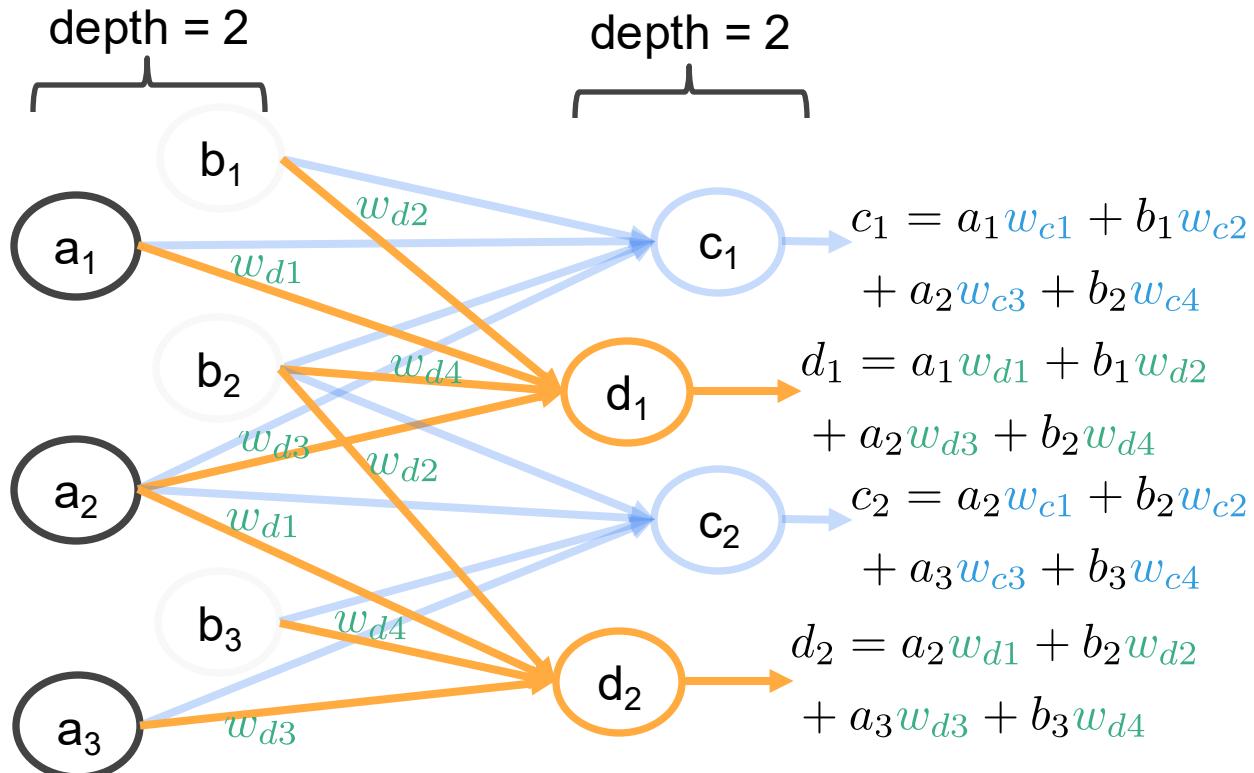
# Convolutional Layers (cont.)



# Convolutional Layers (cont.)

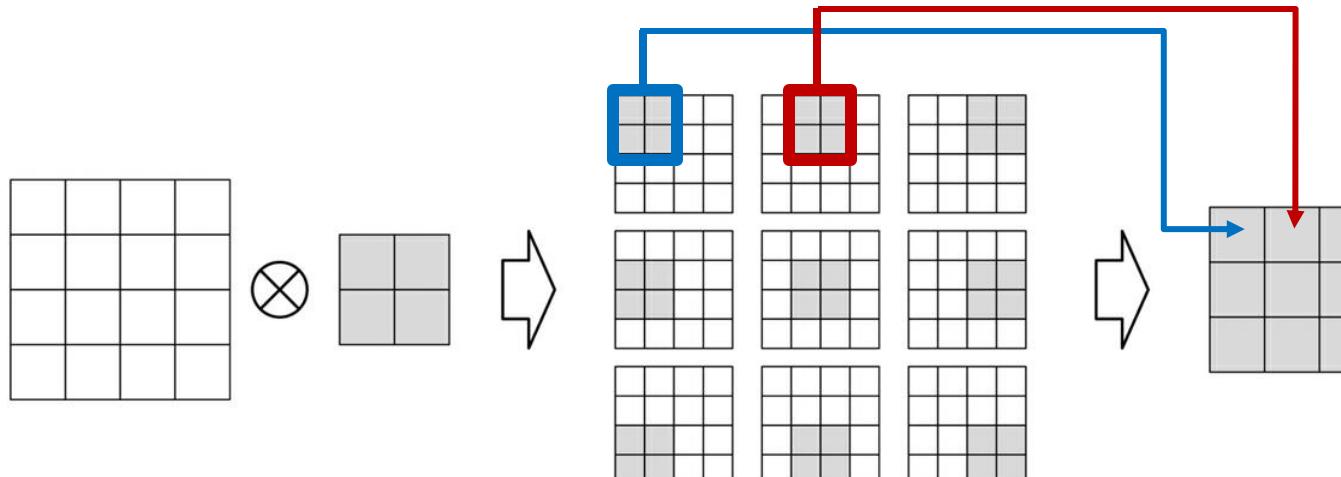


# Convolutional Layers (cont.)



# CNNs (cont.)

- CNNs: a type of NNs well-suited to detecting spatial substructure.



An input matrix  
and a convolutional kernel

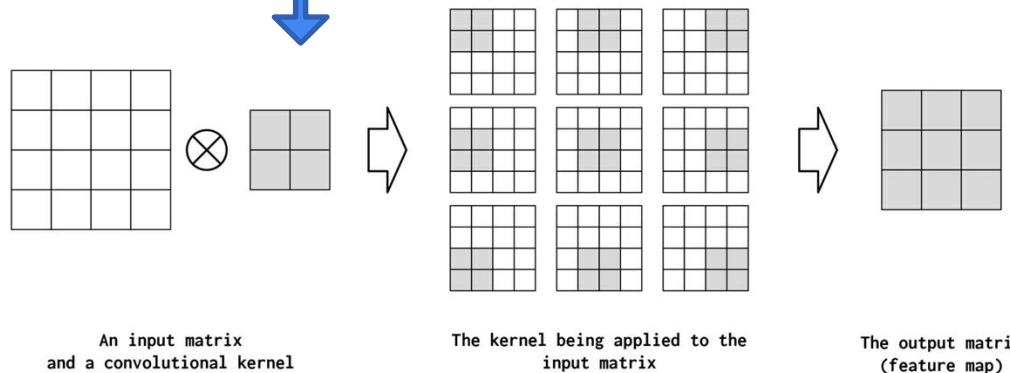
The kernel being applied to the  
input matrix

The output matrix  
(feature map)

# Hyper-parameters of CNNs

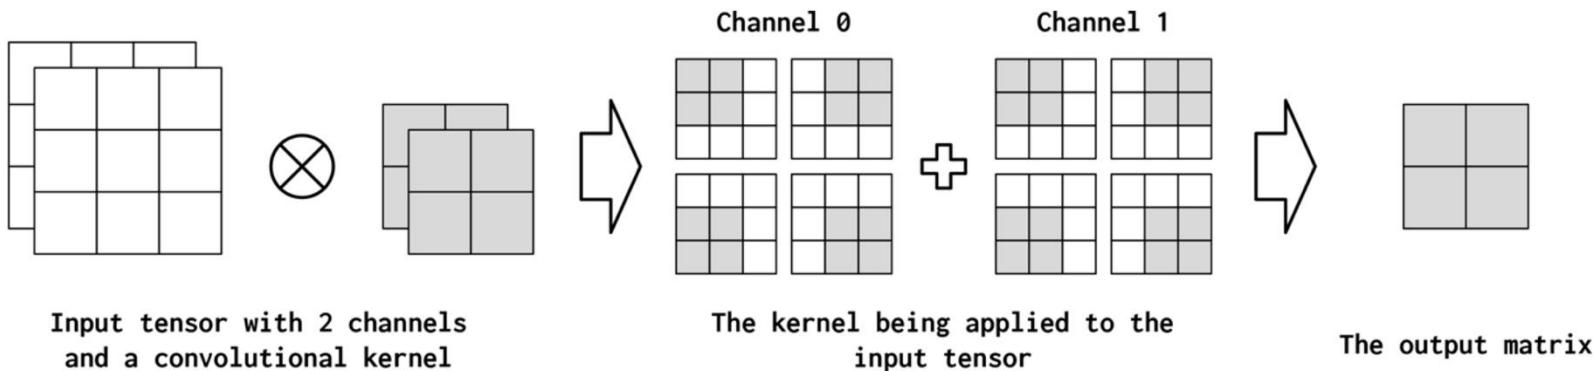


- Dimension of the Convolution Operation
- 2D convolution



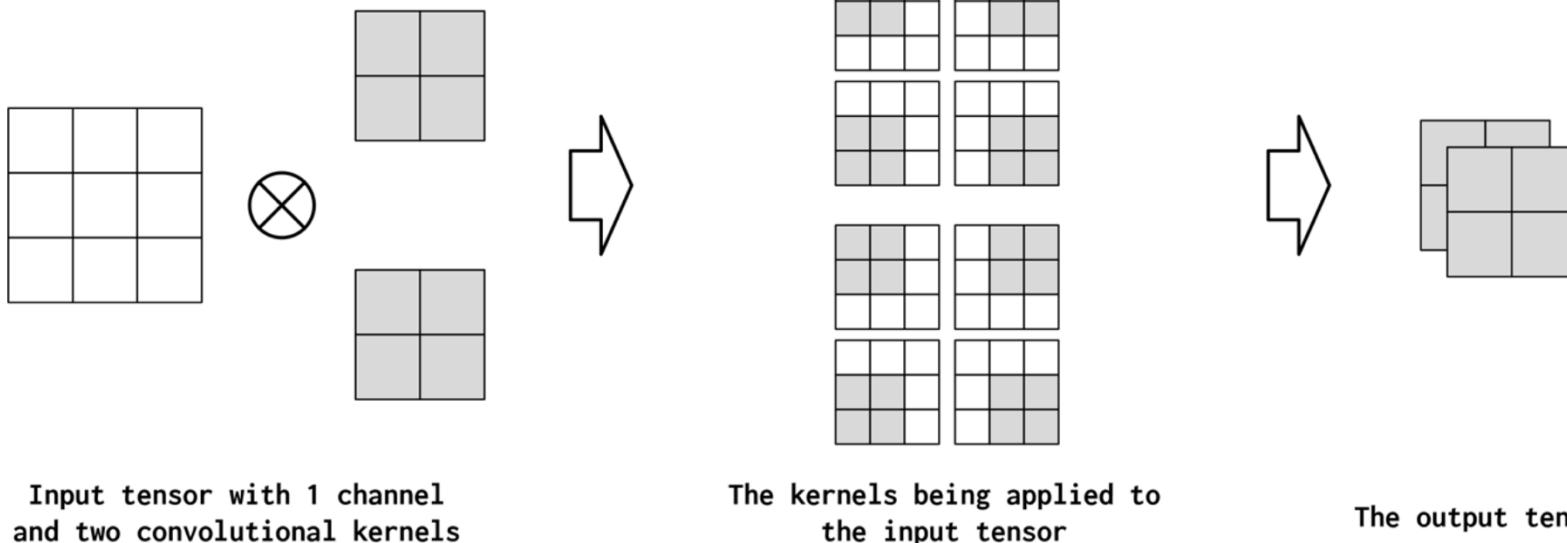
# Hyper-parameters of CNNs (cont.)

- **Channels**
- *input\_channels=2, output\_channels=1, kernel\_size=2, stride=1, padding=0*



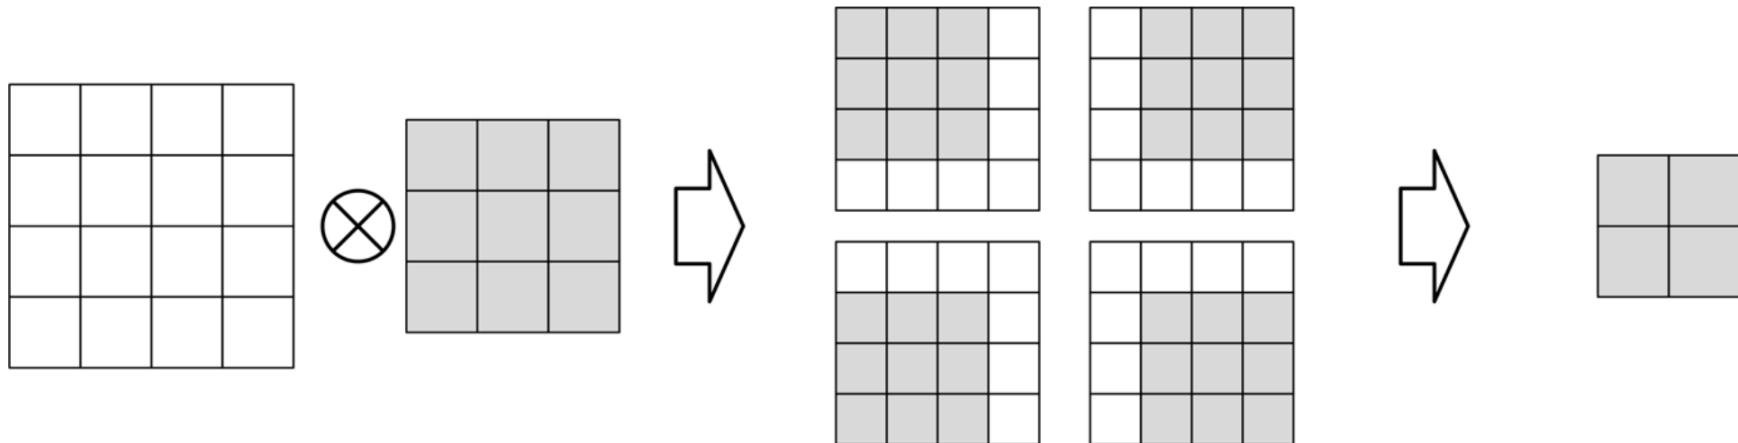
# Hyper-parameters of CNNs (cont.)

- **Channels**
- *input\_channels=1, output\_channels=2, kernel\_size=2, stride=1, padding=0*



# Hyper-parameters of CNNs (cont.)

## - Kernel Size



An input matrix  
and a convolutional kernel

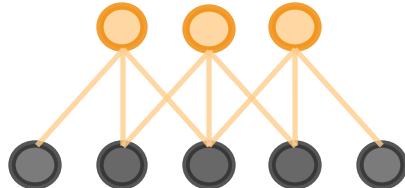
The kernel being applied to the  
input matrix

The output matrix

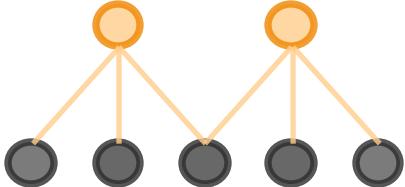
Kernel Size = 3, compare with the example 5 slides before!

# Hyper-parameters of CNNs (cont.)

## - Stride

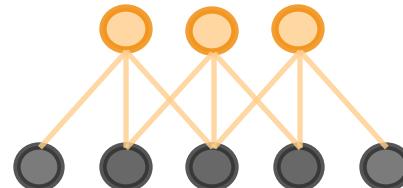


Stride = 1

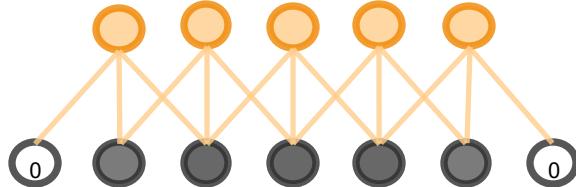


Stride = 2

## Padding



Padding = 0

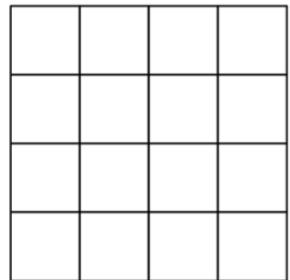


Padding = 1

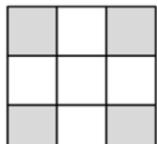
# Hyper-parameters of CNNs (cont.)



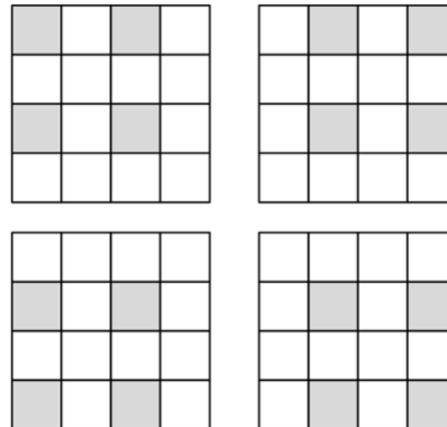
## - Dilation



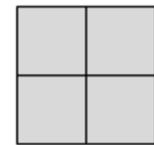
input matrix



Convolution  
kernel with  
dilation



The kernel being applied to the  
input matrix

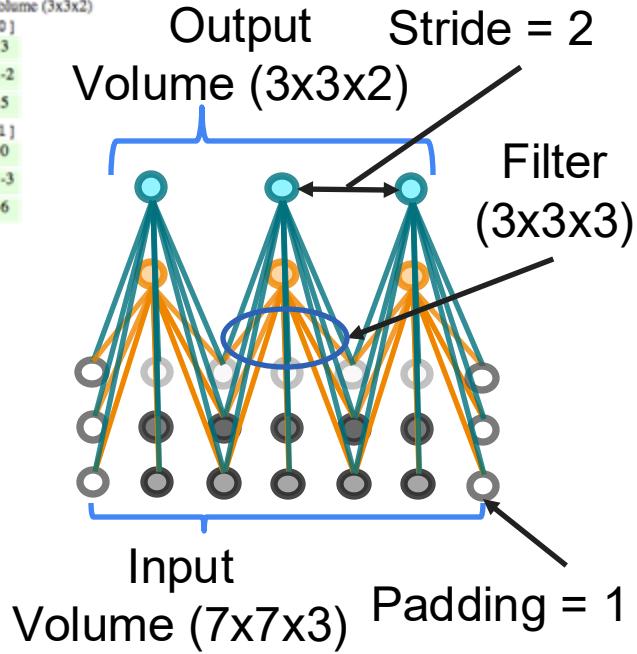


output matrix  
(feature map)

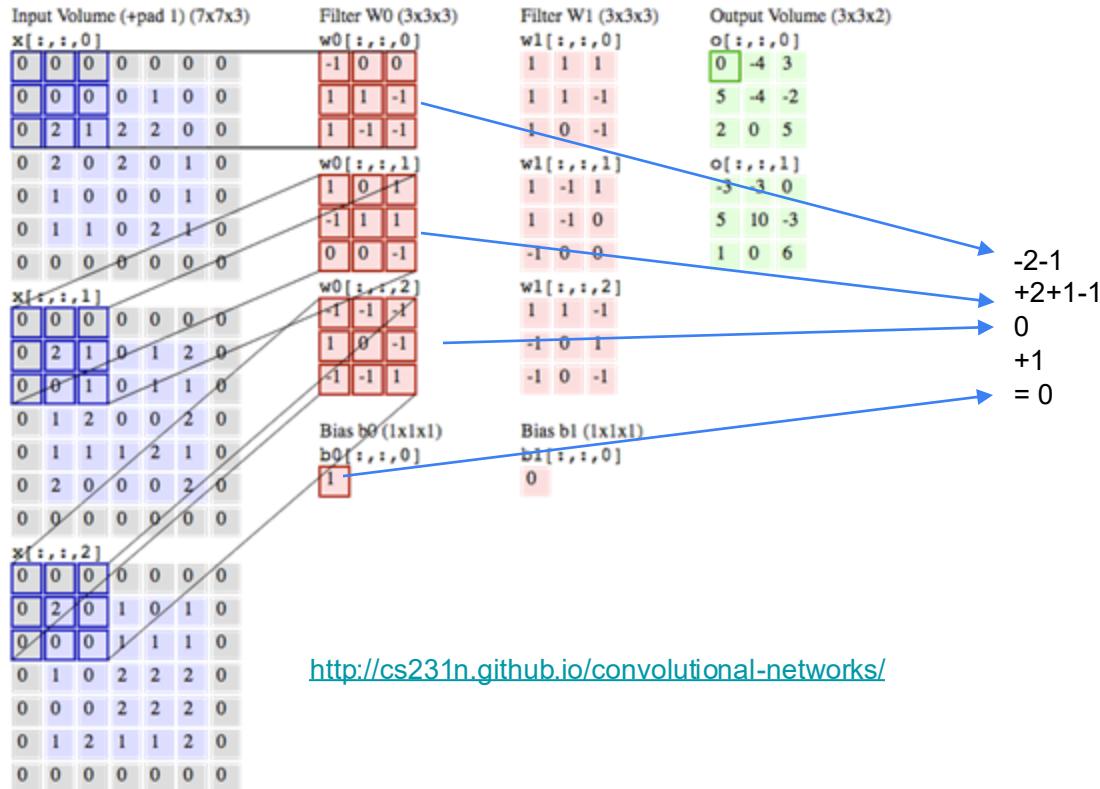
Dilation=2

# CNNs – Example Computation

Input Volume (+pad 1) (7x7x3)	Filter W0 (3x3x3)	Filter W1 (3x3x3)	Output Volume (3x3x2)
$x[z, z, 0]$	$w0[z, z, 0]$	$w1[z, z, 0]$	$o[z, z, 0]$
0 0 0 0 0 0 0 0 0	-1 0 0	1 1 1	0 -4 3
0 0 0 0 1 0 0 0	1 1 -1	1 1 -1	5 -4 -2
0 2 1 2 2 0 0 0	1 -1 -1	1 0 -1	2 0 5
0 2 0 2 0 1 0 0	$w0[z, z, 1]$	$w1[z, z, 1]$	$o[z, z, 1]$
0 1 0 0 0 1 0 0	1 0 1	1 -1 1	-3 -3 0
0 1 1 0 2 1 0 0	-1 1 1	1 -1 0	5 10 -3
0 0 0 0 0 0 0 0 0	0 0 -1	-1 0 0	1 0 6
$x[z, z, 1]$	$w0[z, z, 2]$	$w1[z, z, 2]$	
0 0 0 0 0 0 0 0 0	-1 -1 -1	1 1 -1	
0 2 1 0 1 2 0 0	1 0 -1	-1 0 1	
0 0 1 0 1 1 1 0	-1 -1 1	-1 0 -1	
0 1 2 0 0 2 0 0			
0 1 1 1 2 1 0 0			
0 2 0 0 0 2 0 0			
0 0 0 0 0 0 0 0 0			
$x[z, z, 2]$	$b0[1 \times 1 \times 1]$	$b1[1 \times 1 \times 0]$	
0 0 0 0 0 0 0 0 0	1	0	
0 2 0 1 0 1 1 0 0			
0 0 0 1 1 1 1 0 0			
0 1 2 0 2 2 2 0 0			
0 0 0 2 2 2 2 0 0			
0 1 2 1 1 1 2 0 0			
0 0 0 0 0 0 0 0 0			

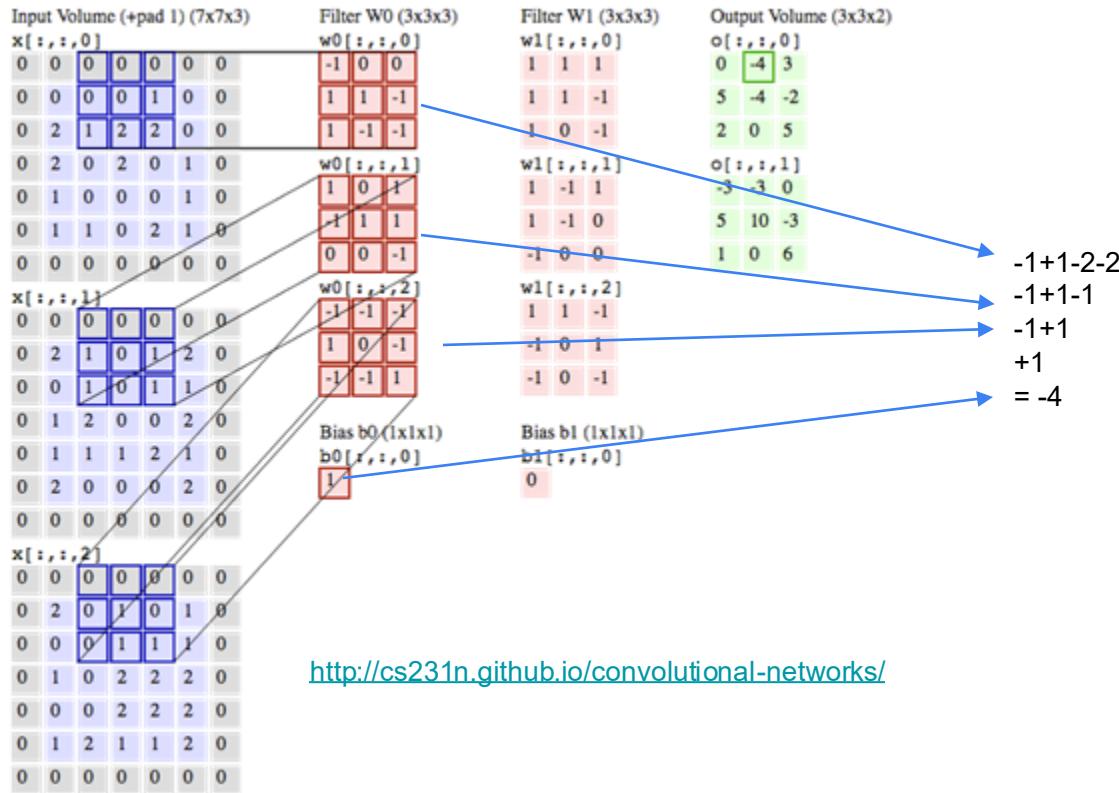


# CNNs – Example Computation (cont.)

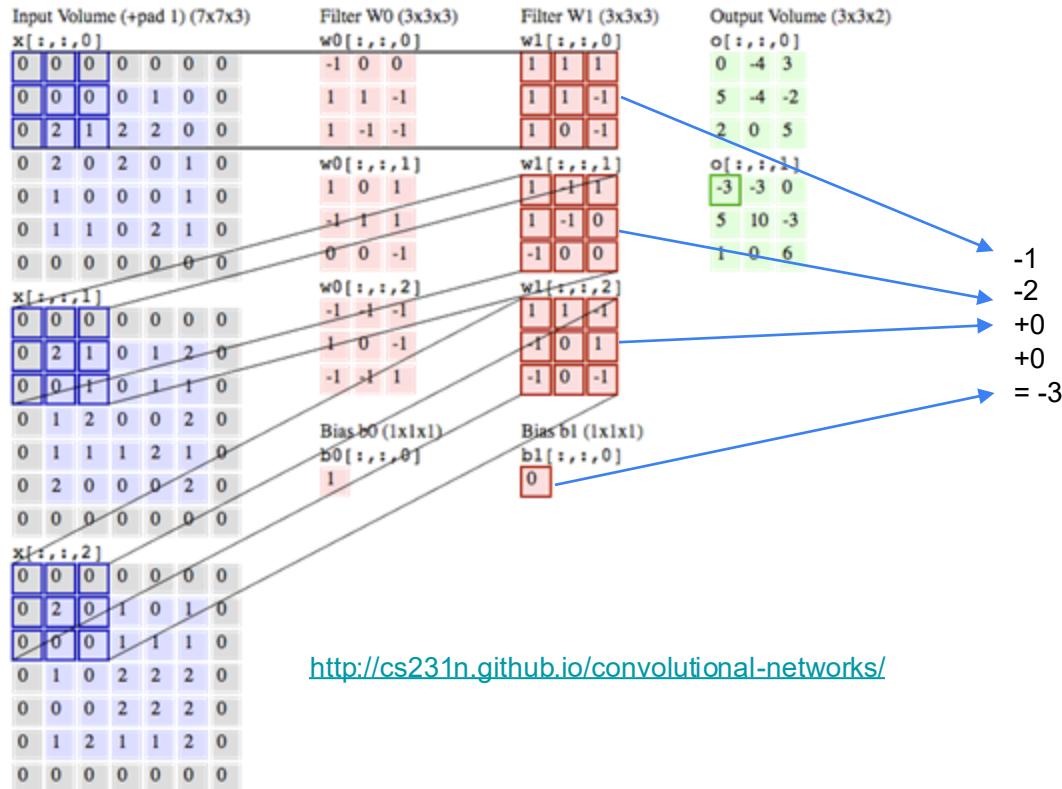


<http://cs231n.github.io/convolutional-networks/>

# CNNs – Example Computation (cont.)



# CNNs – Example Computation (cont.)



<http://cs231n.github.io/convolutional-networks/>

# Relationship with Convolution in Math



- The convolution between two functions, say  $f, g: \mathbb{R}^d \rightarrow R$  is defined as:

$$[f \circledast g](x) = \int_{\mathbb{R}^d} f(z)g(x - z)dz.$$

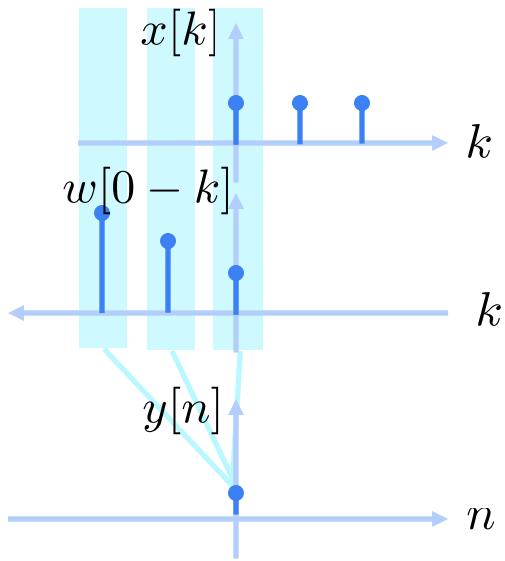
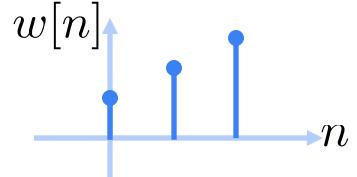
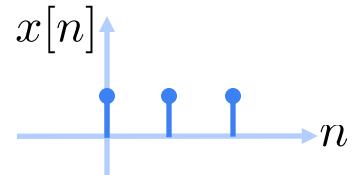
- Whenever we have discrete objects, the integral turns into a sum.

$$[f \circledast g](i) = \sum_a f(a)g(i - a).$$

# Relationship with Convolution in Math (cont.)



$$y[n] = \sum_k x[k]w[n-k]$$

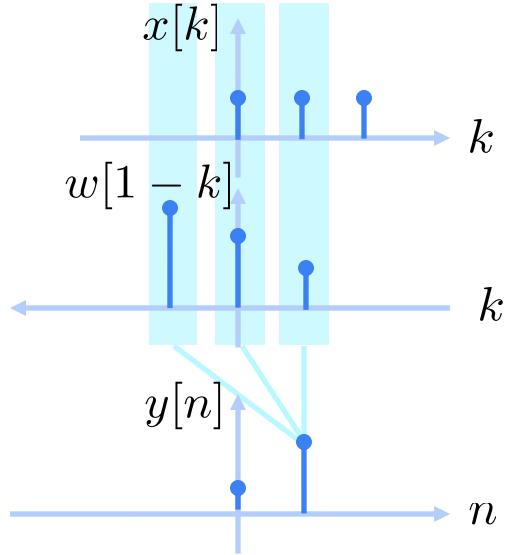
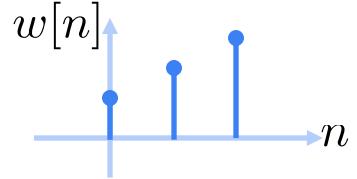
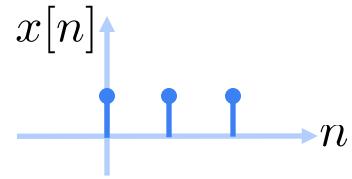


$$y[0] = x[-2]w[2] + x[-1]w[1] + x[0]w[0]$$

# Relationship with Convolution in Math (cont.)



$$y[n] = \sum_k x[k]w[n-k]$$

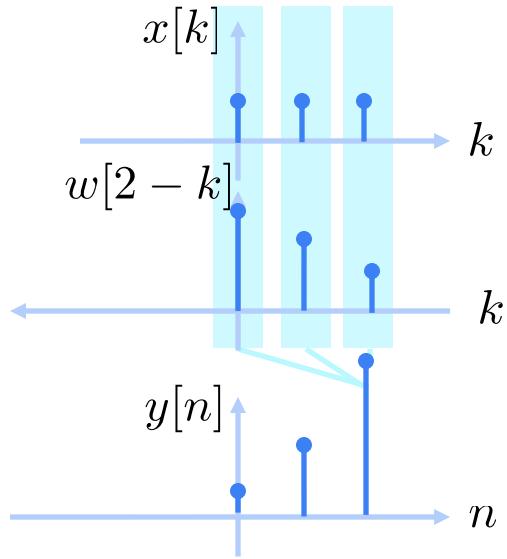
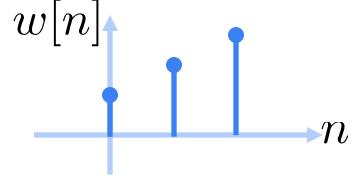
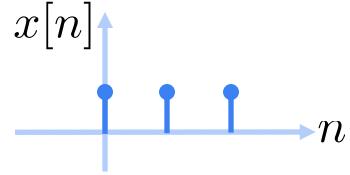


$$y[1] = x[-1]w[2] + x[0]w[1] + x[2]w[0]$$

# Relationship with Convolution in Math (cont.)



$$y[n] = \sum_k x[k]w[n-k]$$

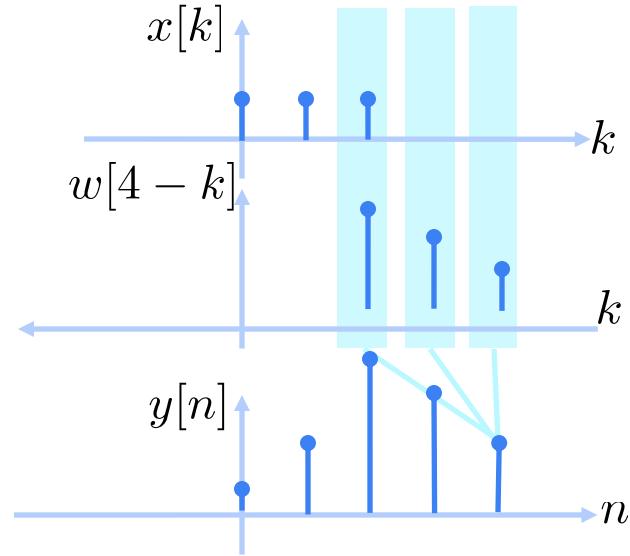
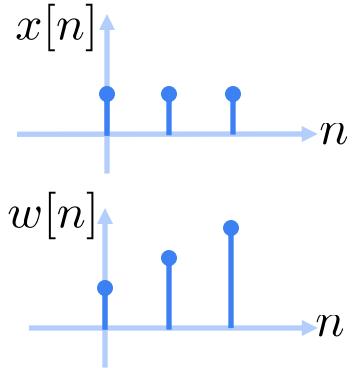


$$y[2] = x[0]w[2] + x[1]w[1] + x[2]w[0]$$

# Relationship with Convolution in Math (cont.)



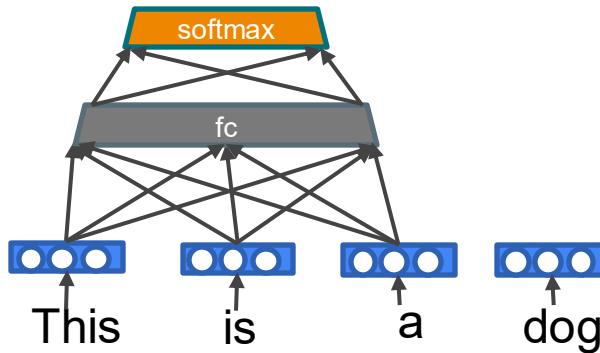
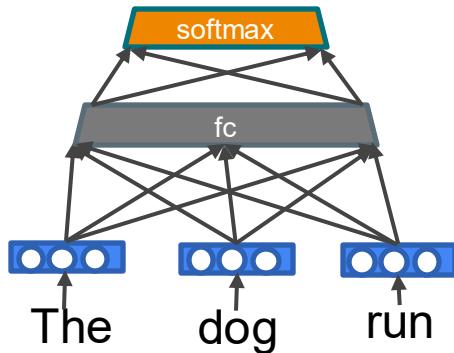
$$y[n] = \sum_k x[k]w[n - k]$$



$$y[4] = x[2]w[2] + x[3]w[1] + x[4]w[0]$$

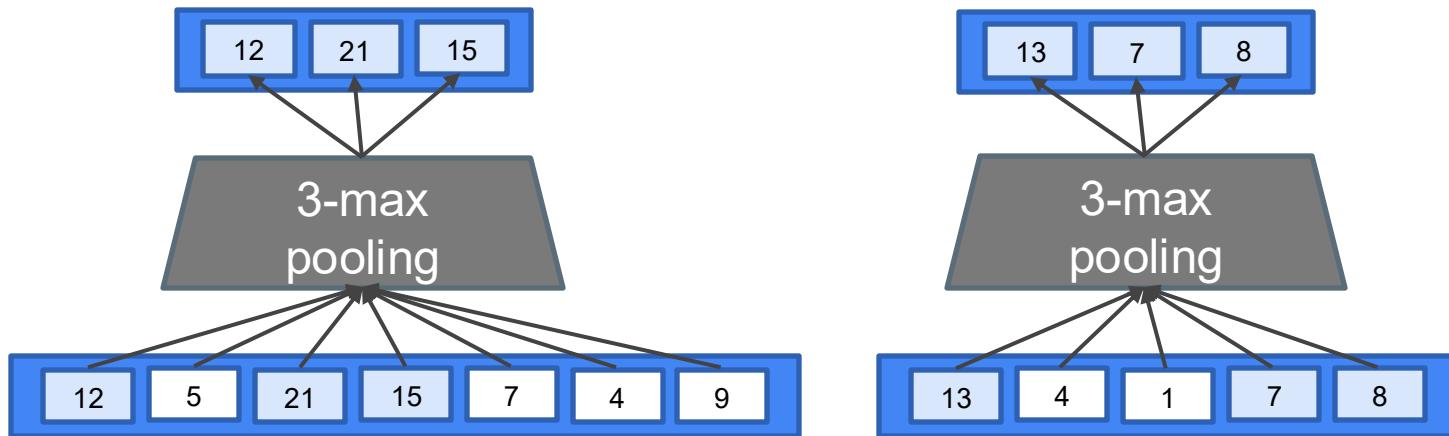
# Various Input Sizes

- Fully-connected layer and softmax layer
  - need fixed-size input



# k-max Pooling

- choose the k-max values
- preserve the order of input values
- variable-size input, fixed-size output



# Pooling Layer

1	3	2	4
5	7	6	8
0	0	3	3
5	5	0	0

Maximum  
Pooling



Average  
Pooling



$$\text{Max}(1,3,5,7) = 7$$

7	8
5	3

$$\text{Avg}(1,3,5,7) = 4$$

$$\text{Max}(0,0,5,5) = 5$$

4	5
2.5	1.5

# Topics for Today



## Sequence Modeling

- Convolutional Neural Networks
- Recurrent Neural Networks
- Example Applications

# Recurrent Neural Networks (RNNs)

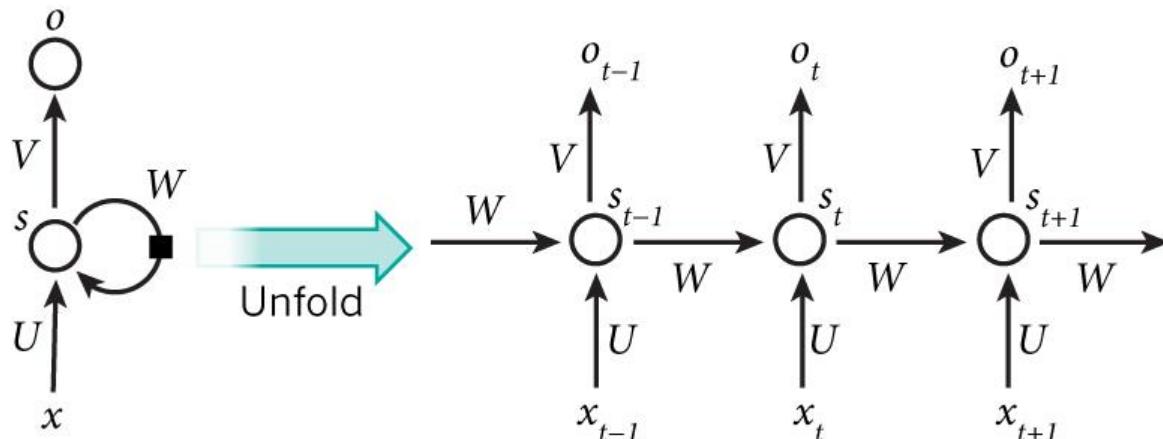


- Idea: condition the neural network on all previous words and tie the weights at each time step
- Assumption: **temporal** information matters

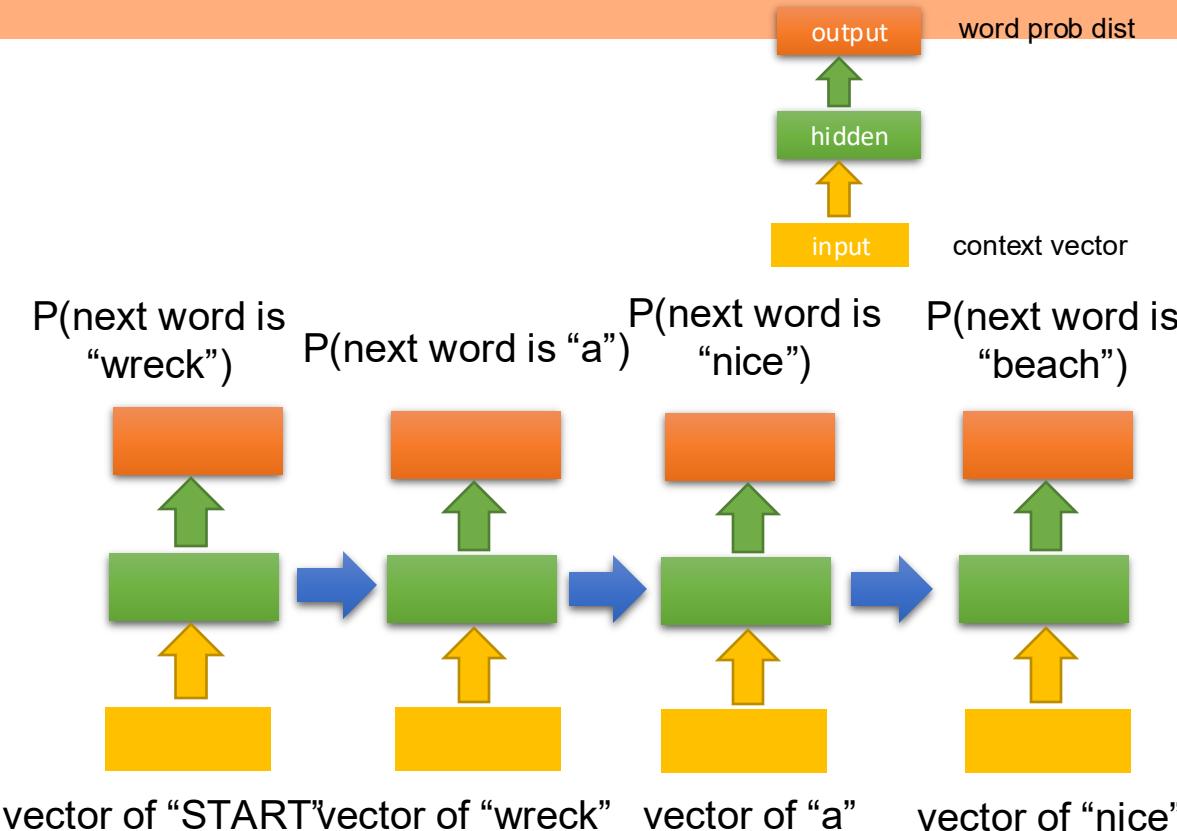
# RNNs – Definition

$$s_t = \sigma(W s_{t-1} + U x_t) \quad \sigma(\cdot): \text{tanh, ReLU}$$

$$o_t = \text{softmax}(V s_t)$$



# RNN Language Modeling (RNN-LM)



Idea: pass the information from the previous hidden layer to leverage all contexts

# RNN-LM Formulation

- At each time step,  
 $h_t = \sigma(Wh_{t-1} + Ux_t)$   
 $\hat{y}_t = \text{softmax}(Vh_t)$

$$P(x_{t+1} = w_j \mid x_1, \dots, x_t) = \hat{y}_{t,j}$$



$W$



$U$



vector of the current  
word

probability of the next word



$V$



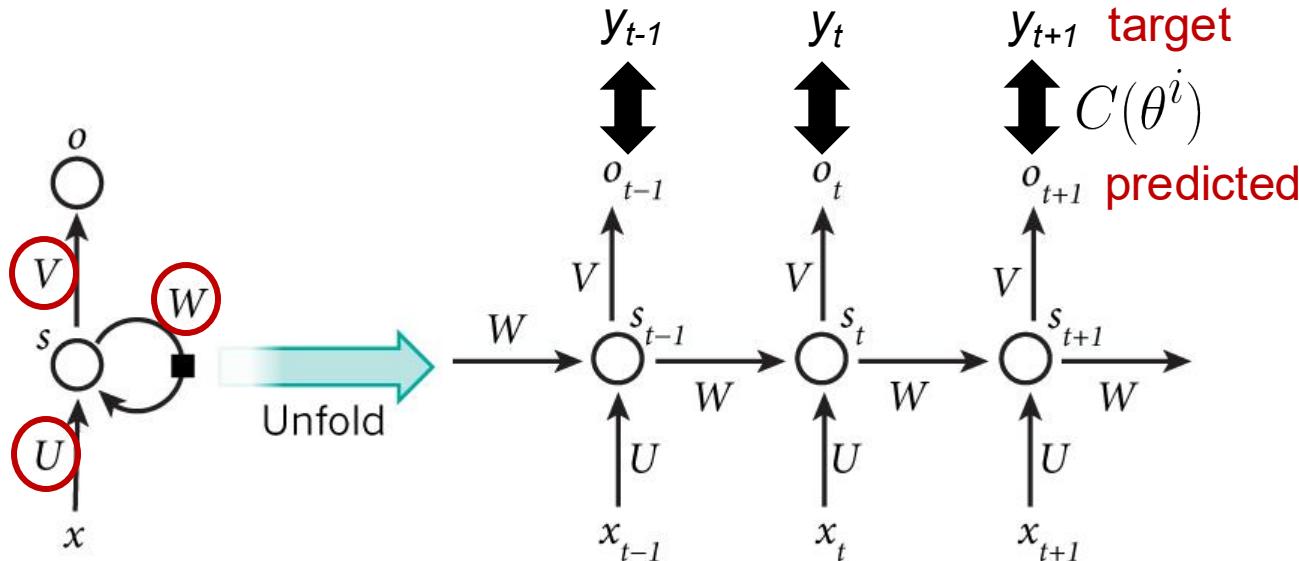
$U$



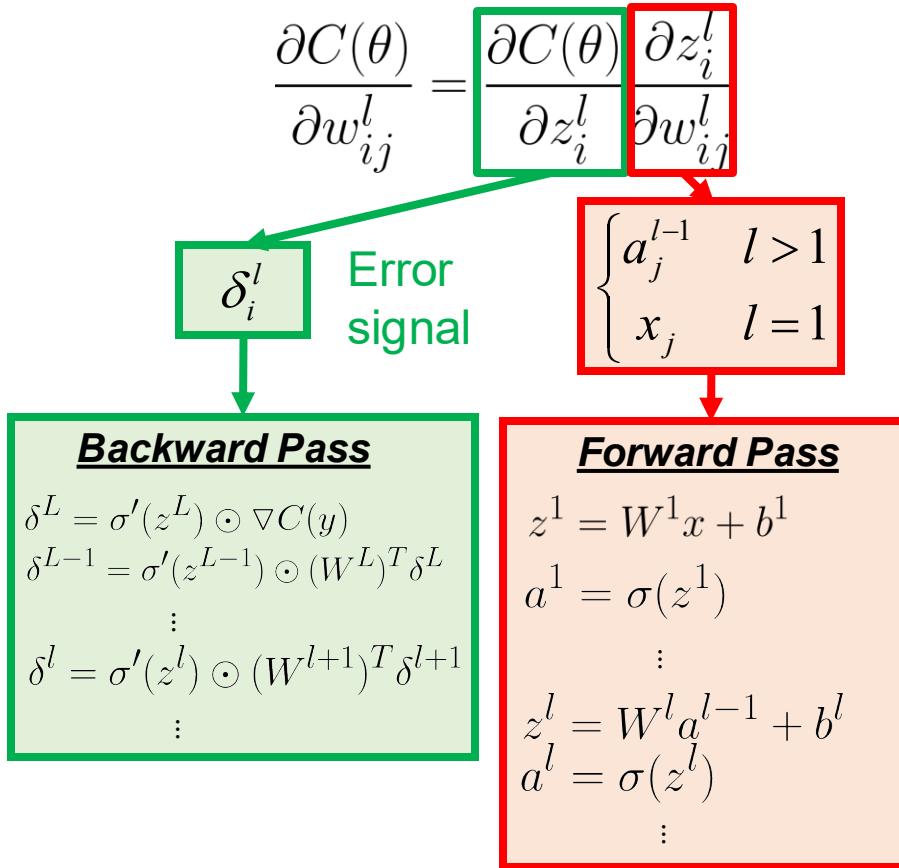
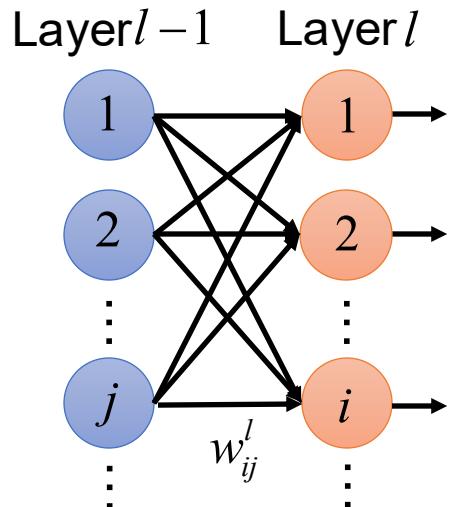
# RNNs – Model Training

- All model parameters  $\theta = \{U, V, W\}$  can be updated by

$$\theta^{i+1} \leftarrow \theta^i - \eta \nabla_{\theta} C(\theta^i)$$

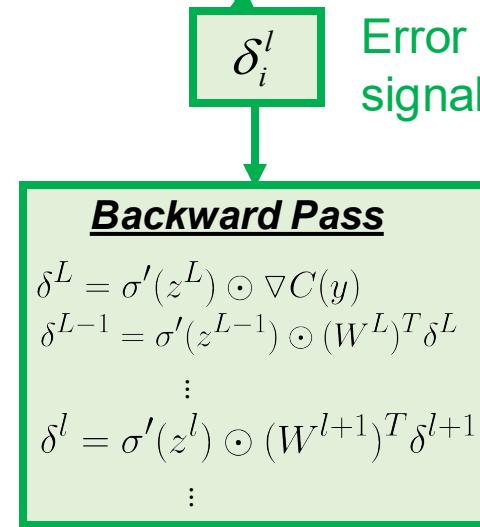
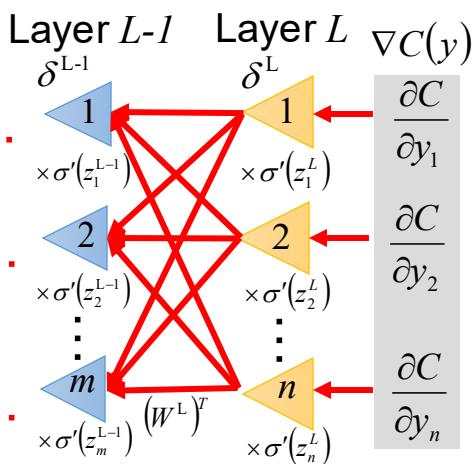
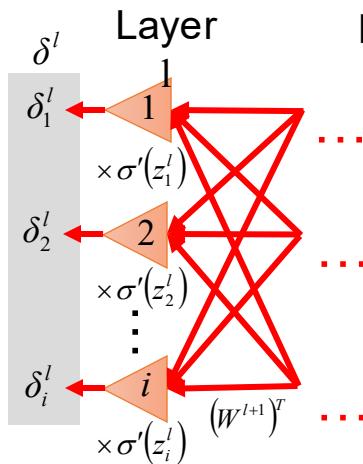


# Backpropagation



# Backpropagation (cont.)

$$\frac{\partial C(\theta)}{\partial w_{ij}^l} = \boxed{\frac{\partial C(\theta)}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l}}$$

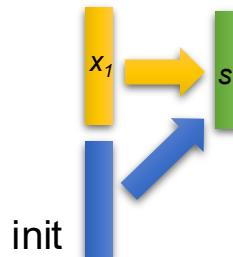


# Backpropagation Through Time (BPTT)

- Unfold



- Input: init,  $x_1, x_2, \dots, x_t$
- Output:  $o_t$
- Target:  $y_t$



$$\frac{\partial C}{\partial o_1}$$

$$\frac{\partial C}{\partial o_2}$$

$$\vdots$$

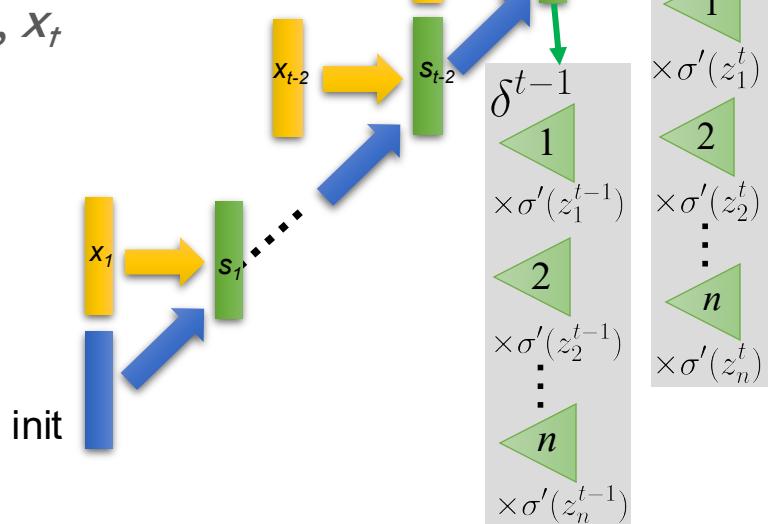
$$\frac{\partial C}{\partial o_n}$$

# BPTT (cont.)

## - Unfold

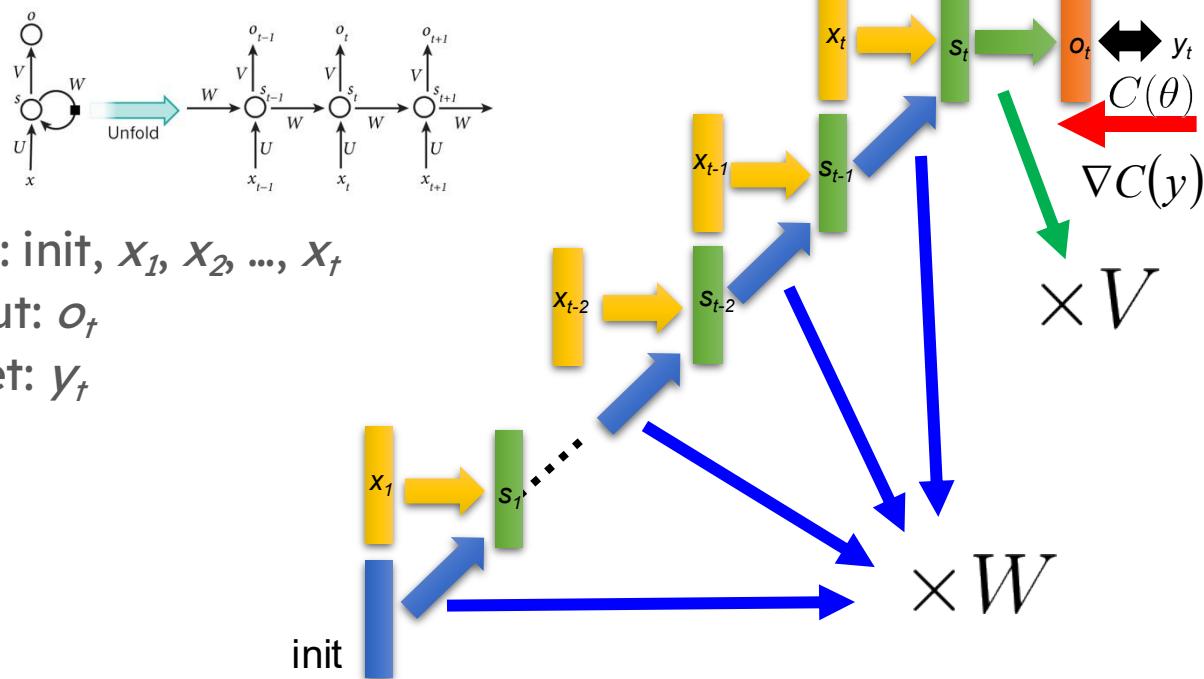


- Input: init,  $x_1, x_2, \dots, x_t$
- Output:  $o_t$
- Target:  $y_t$



# BPTT (cont.)

## - Unfold

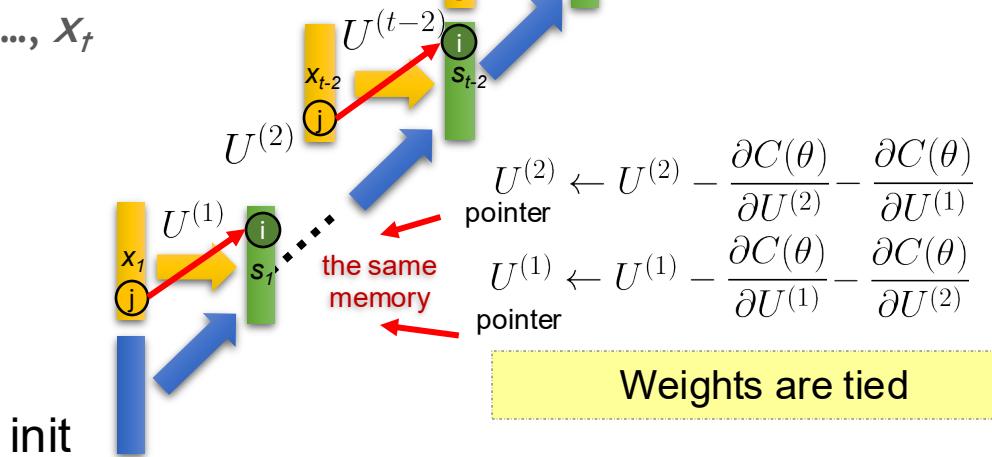


# BPTT (cont.)

## - Unfold



- Input: init,  $x_1, x_2, \dots, x_t$
- Output:  $o_t$
- Target:  $y_t$

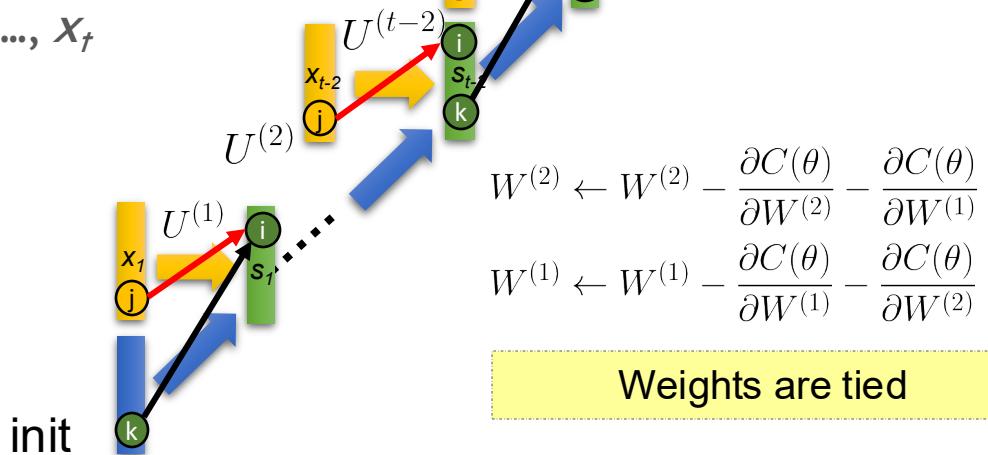


# BPTT (cont.)

## - Unfold



- Input: init,  $x_1, x_2, \dots, x_t$
- Output:  $o_t$
- Target:  $y_t$



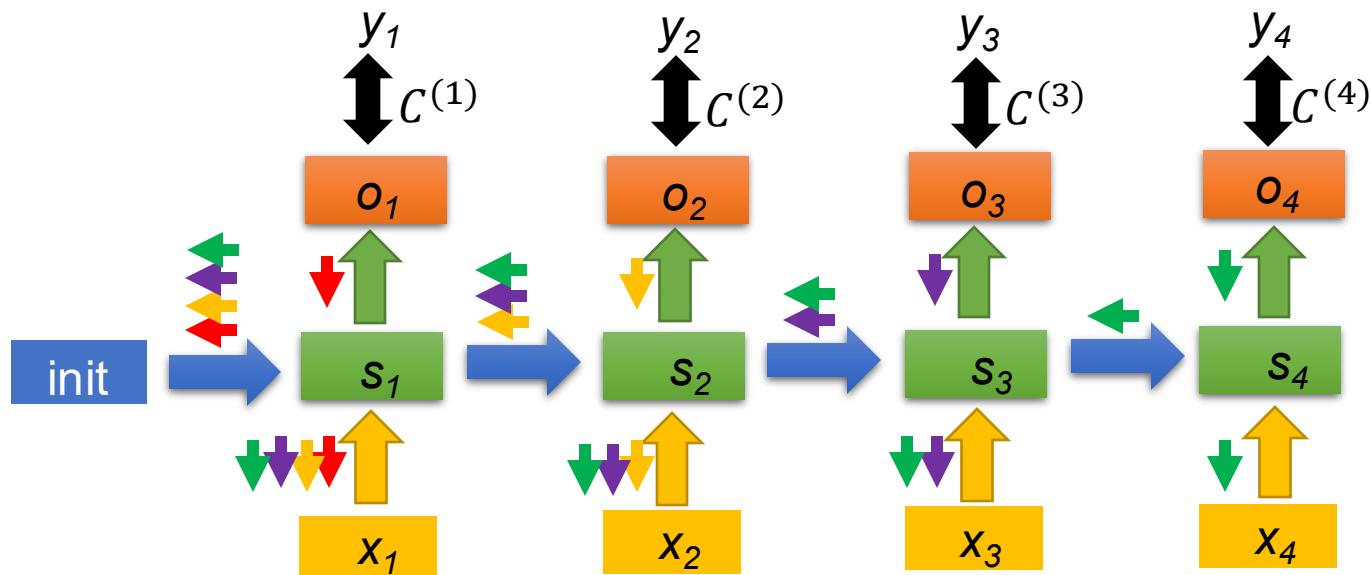
# BPTT (cont.)

Forward Pass:

Compute  $s_1, s_2, s_3, s_4 \dots$

Backward Pass:

For  $C^{(4)}$  For  $C^{(3)}$   
For  $C^{(2)}$  For  $C^{(1)}$



# Training Issues with RNN

- The gradient is a product of Jacobian matrices, each associated with a step in the forward computation
- Multiply the same matrix at each time step during backprop

$$\delta^l = \sigma'(z^l) \odot (W^{l+1})^T \delta^{l+1}$$

The gradient becomes very small or very large quickly  
→ vanishing or exploding gradient

# Clipping: A Solution to Exploding Gradients



Idea: control the gradient value to avoid exploding

---

**Algorithm 1** Pseudo-code for norm clipping

---

```
 $\hat{g} \leftarrow \frac{\partial \mathcal{E}}{\partial \theta}$ 
if  $\|\hat{g}\| \geq threshold$  then
     $\hat{g} \leftarrow \frac{threshold}{\|\hat{g}\|} \hat{g}$ 
end if
```

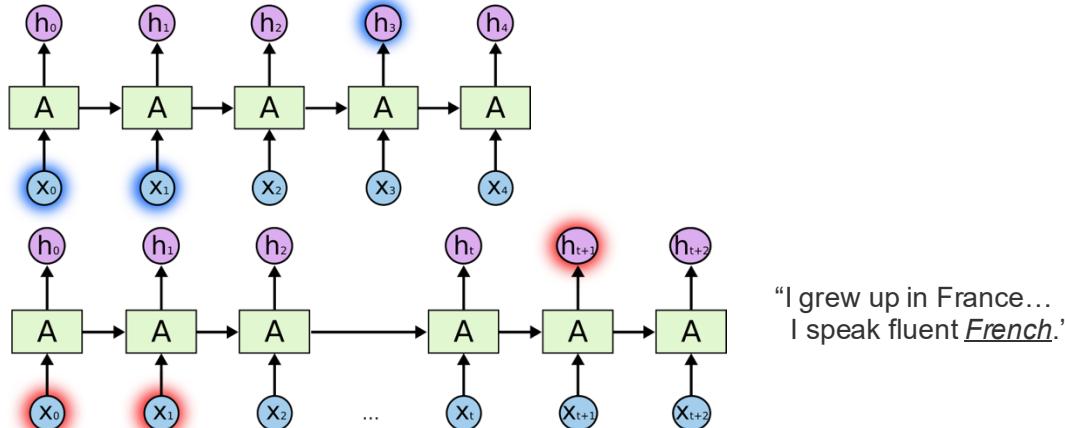
---

Parameter setting: values from half to ten times the average can still yield convergence

# Gating Mechanisms: Solution to Vanishing Gradients

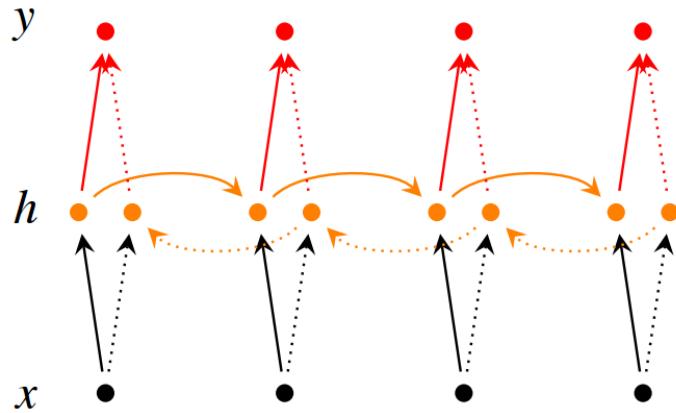


- RNN models temporal sequence information
  - can handle “long-term dependencies” *in theory*



Issue: RNN cannot handle such “long-term dependencies” in practice due to vanishing gradient  
→ apply the gating mechanism to directly encode the long-distance information

# RNN Extensions – Bidirectional RNNs



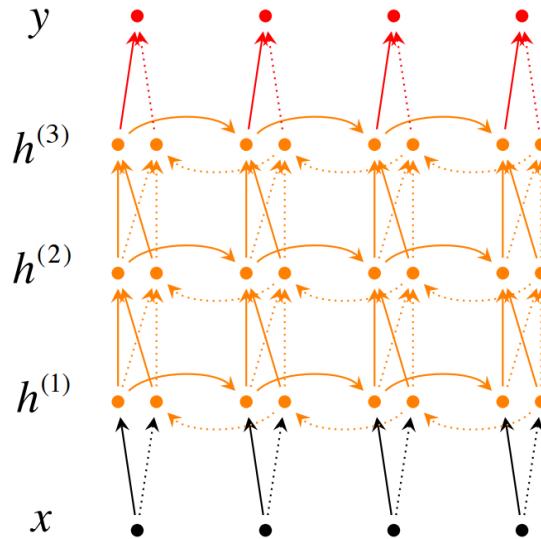
$$\vec{h}_t = f(\vec{W}x_t + \vec{V}\vec{h}_{t-1} + \vec{b})$$

$$\overleftarrow{h}_t = f(\overleftarrow{W}x_t + \overleftarrow{V}\overleftarrow{h}_{t+1} + \overleftarrow{b})$$

$$y_t = g(U[\vec{h}_t; \overleftarrow{h}_t] + c)$$

$h = [\vec{h}; \overleftarrow{h}]$  represents (summarizes) the past and future around a single token

# RNN Extensions – Deep Bidirectional RNNs



$$\overset{\rightarrow}{h}_t^{(i)} = f(\overset{\rightarrow}{W}^{(i)} h_t^{(i-1)} + \overset{\rightarrow}{V}^{(i)} \overset{\rightarrow}{h}_{t-1}^{(i)} + \overset{\rightarrow}{b}^{(i)})$$

$$\overset{\leftarrow}{h}_t^{(i)} = f(\overset{\leftarrow}{W}^{(i)} h_t^{(i-1)} + \overset{\leftarrow}{V}^{(i)} \overset{\leftarrow}{h}_{t+1}^{(i)} + \overset{\leftarrow}{b}^{(i)})$$

$$y_t = g(U[\overset{\rightarrow}{h}_t^{(L)}; \overset{\leftarrow}{h}_t^{(L)}] + c)$$

Each memory layer passes an intermediate representation to the next



# Topics for Today

## Sequence Modeling

- Convolutional Neural Networks
- Recurrent Neural Networks
- Example Applications

# Example: CNNs for Text Classification

I  
like  
this  
movie  
very  
much  
!

0.6	0.5	0.2	-0.1	0.4
0.8	0.9	0.1	0.5	0.1
0.4	0.6	0.1	-0.1	0.7
---	---	---	---	---
---	---	---	---	---
---	---	---	---	---
---	---	---	---	---

0.2	0.1	0.2	0.1	0.1
0.1	0.1	0.4	0.1	0.1



- Example depiction from: <http://www.joshuakim.io/understanding-how-convolutional-neural-network-cnn-perform-text-classification-with-word-embeddings/>

# Example: CNNs for Text Classification (cont.)



I  
like  
this  
movie  
very  
much  
!

0.6	0.5	0.2	-0.1	0.4
0.8	0.9	0.1	0.5	0.1
0.4	0.6	0.1	-0.1	0.7
---	---	---	---	---
---	---	---	---	---
---	---	---	---	---
---	---	---	---	---

0.2	0.1	0.2	0.1	0.1
0.1	0.1	0.4	0.1	0.1

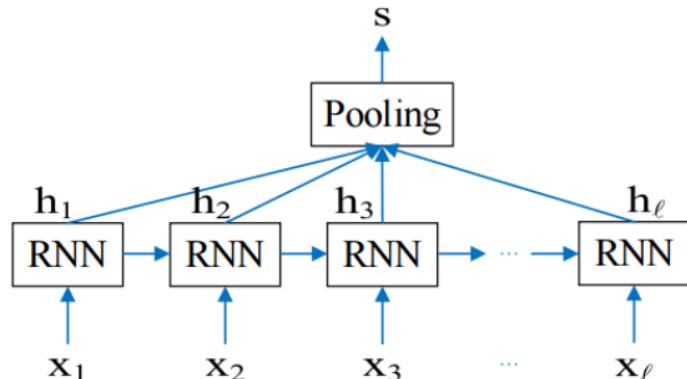
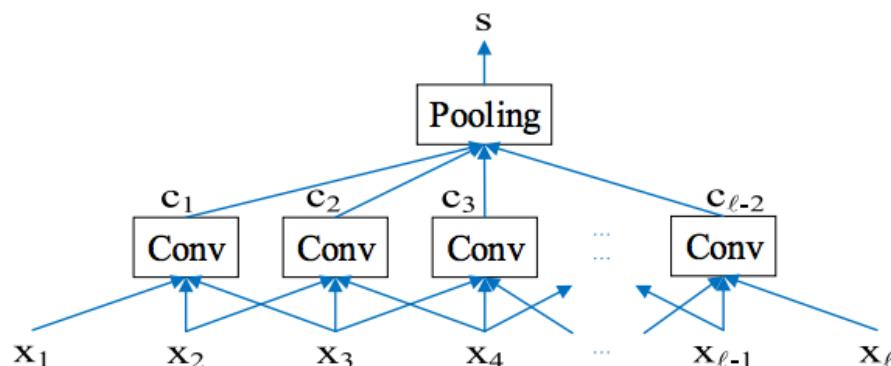


- Example depiction from: <http://www.joshuakim.io/understanding-how-convolutional-neural-network-cnn-perform-text-classification-with-word-embeddings/>

# Example: CNNs and RNNs for Text Classification

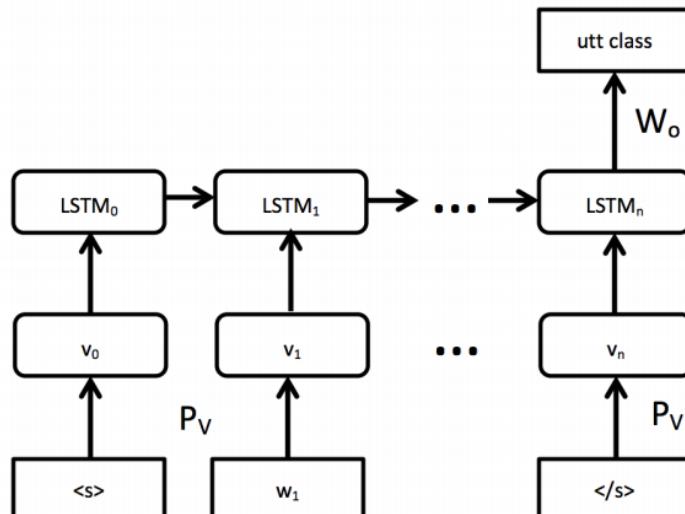
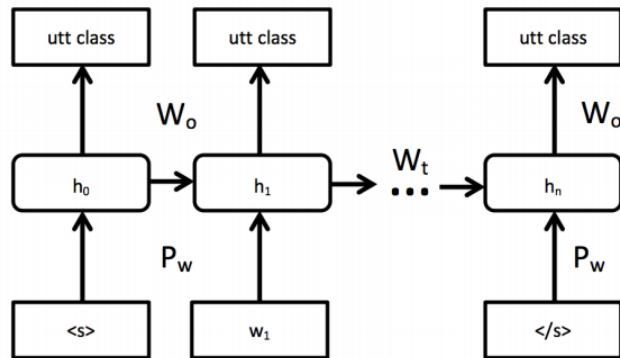


- ([Lee & Dernoncourt, NAACL, 2016](#))
- Dialogue Act Classification



# Example: RNNs for Text Classification

- ([Ravuri and Stolcke, Interspeech, 2015](#))
- Addressee Detection



# Example: RNNs for Sequence Tagging

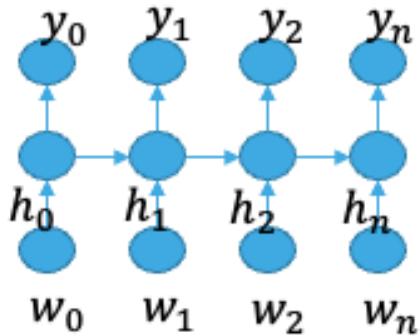


- (Mesnil et al., IEEE TASLP, 2015)
  - Slot tagging

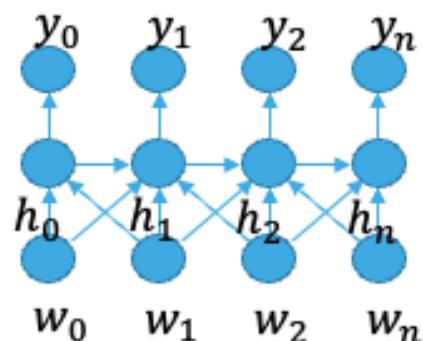


<START> send an email to bob about fishing this weekend <END>

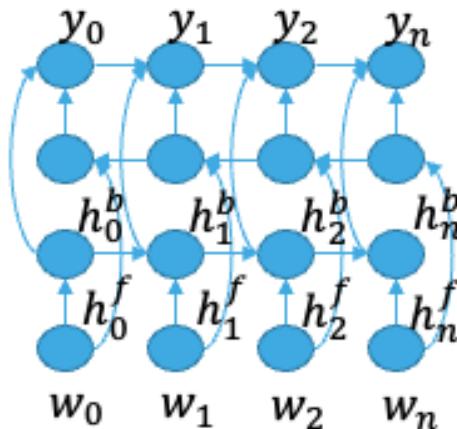
↓      ↓      ↓      ↓      ↓      ↓      ↓      ↓      ↓      ↓  
 O      O      O      O      O      B-contact name   I-contact name   I-subject   I-subject   send\_email



(a) RNN



(b) RNN-LA

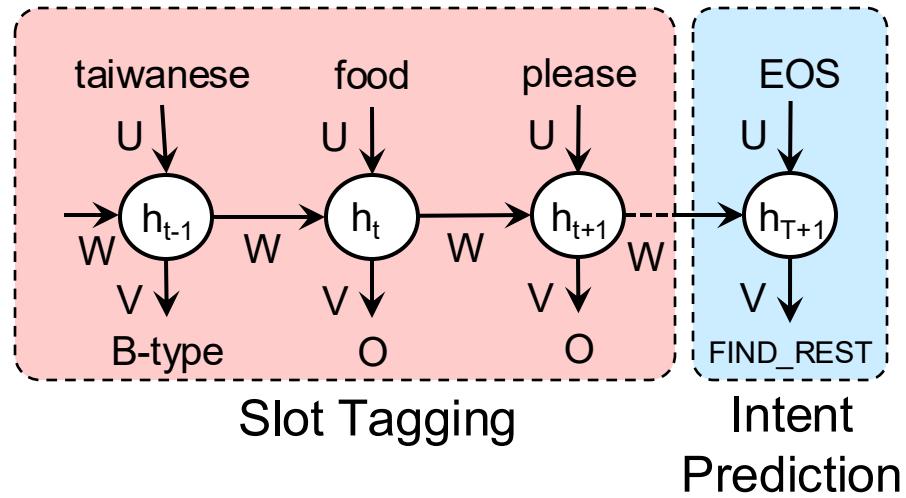


(c) bi-RNN

# Example: RNNs for Joint Utterance Classification and Sequence Tagging



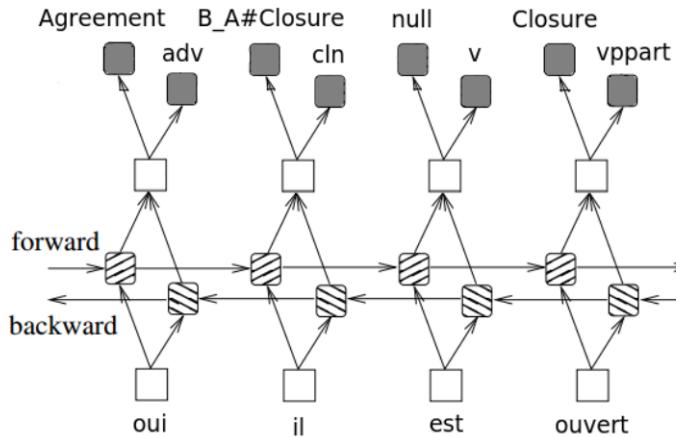
- ([Hakkani-Tür et al., Interspeech, 2016](#))
- Slot filling (or tagging) and intent prediction in the same output sequence



# Example: RNNs for Multi-Task Sequence Tagging



- (Tafforeau et al., Interspeech, 2016)
  - Goal: exploit data from domains/tasks with a lot of data to improve ones with less data
  - Lower layers are shared across domains/tasks
  - Output layer is specific to task



# Topics for Thursday

## Model Architectures and Contextual Embeddings

- Long-Short Term Memory (LSTM)
- Gated Recurrent Units (GRU)
- Example Sequence Classification Tasks
- Elmo and Contextual Embeddings

# Midterm 1



- September 30<sup>th</sup>
- In class
- True/False and multiple-choice questions from content we discussed in class.

# Preparing Final Project Proposals



- Final Project Proposal team sign up deadline: Sept 23<sup>rd</sup>
- Spreadsheet to sign up project teams:  
[https://docs.google.com/spreadsheets/d/1EJ\\_5Xby0mRhHFmSRSmxlv6Gws4Qs5T8P5JUiKYKAZcA/edit?usp=sharing](https://docs.google.com/spreadsheets/d/1EJ_5Xby0mRhHFmSRSmxlv6Gws4Qs5T8P5JUiKYKAZcA/edit?usp=sharing)
- Reach out to me or TAs soon if you need help with project ideas and teaming.