

# Reinforcement Learning in the Era of LLMs

Ishika Agarwal  
CS 546: Advanced NLP



# Introduction



RL is a very hot topic right now!

A screenshot of a Google Scholar search results page. The search query in the bar is "'reinforcement learning' llms". A large black arrow points from the text "RL is a very hot topic right now!" above to the search bar. The results show approximately 19,200 results found in 0.06 seconds. The first result is a paper titled "Deepseek-r1 incentivizes reasoning in LLMs through reinforcement learning" by D Guo, D Yang, H Zhang, J Song, P Wang, Q Zhu, R Xu, et al., published in Nature in 2025. The second result is a paper titled "Rlfh deciphered: A critical analysis of reinforcement learning from human feedback for LLMs" by S Chaudhari, P Aggarwal, V Murahari, et al., published in ACM Computing in 2025.

Google Scholar

"reinforcement learning" llms

Articles

About 19,200 results (0.06 sec)

Any time

Since 2025

Since 2024

Since 2021

Custom range...

Sort by relevance

Sort by date

Deepseek-r1 incentivizes reasoning in LLMs through reinforcement learning  
D Guo, D Yang, H Zhang, J Song, P Wang, Q Zhu, R Xu... - Nature, 2025 - nature.com  
... Recent breakthroughs, exemplified by large language models (LLMs) 1,2 and chain-of-... that  
the reasoning abilities of LLMs can be incentivized through pure reinforcement learning (RL), ...  
☆ Save ♫ Cite Cited by 69 Related articles All 6 versions

Rlfh deciphered: A critical analysis of reinforcement learning from human  
feedback for LLMs  
S Chaudhari, P Aggarwal, V Murahari... - ACM Computing ..., 2025 - dl.acm.org

# Introduction



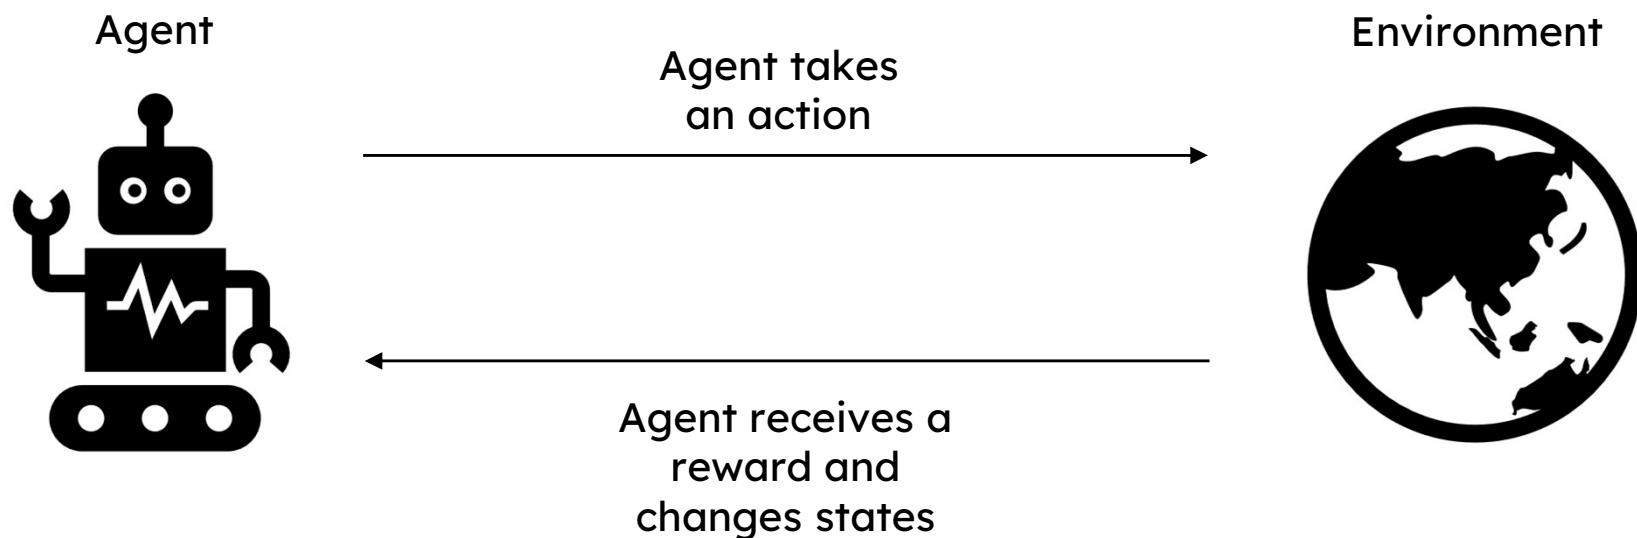
## Model development:

1. Pre-trained model (learns how to speak)
2. Instruction-tuned model (learns how to be useful)
  - a. SFT
3. Generalizable/task specific models
  - a. RL

# What is RL?



- Learning from positive/negative reinforcement



# Agenda



- I. Basics
- II. Common RL Algorithms
- III. RL Applications

# Agenda



## I. Basics

- I. Terminology
- II. Bellman equations
- III. Policy gradient

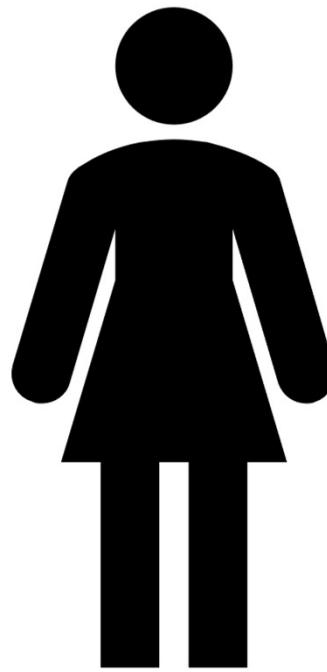
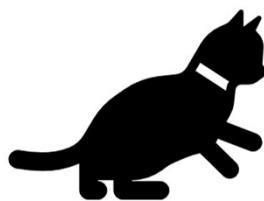
## II. Common RL Algorithms

## III. RL Applications

# Terminology



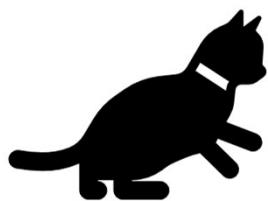
Suppose you are training your cat...



# Terminology



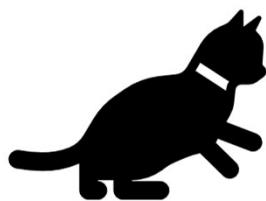
Your cat is an **agent!**



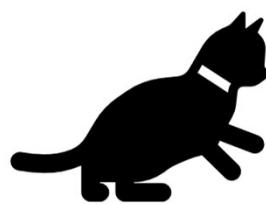
# Terminology



Your cat is an agent! And its **state** is *sitting*.



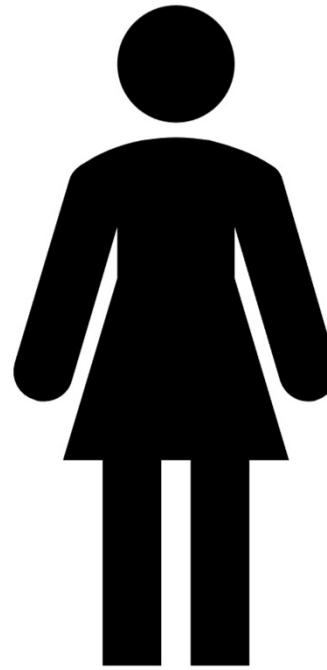
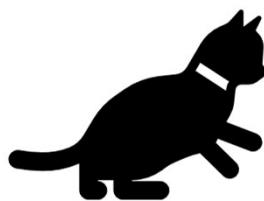
# Terminology



# Terminology



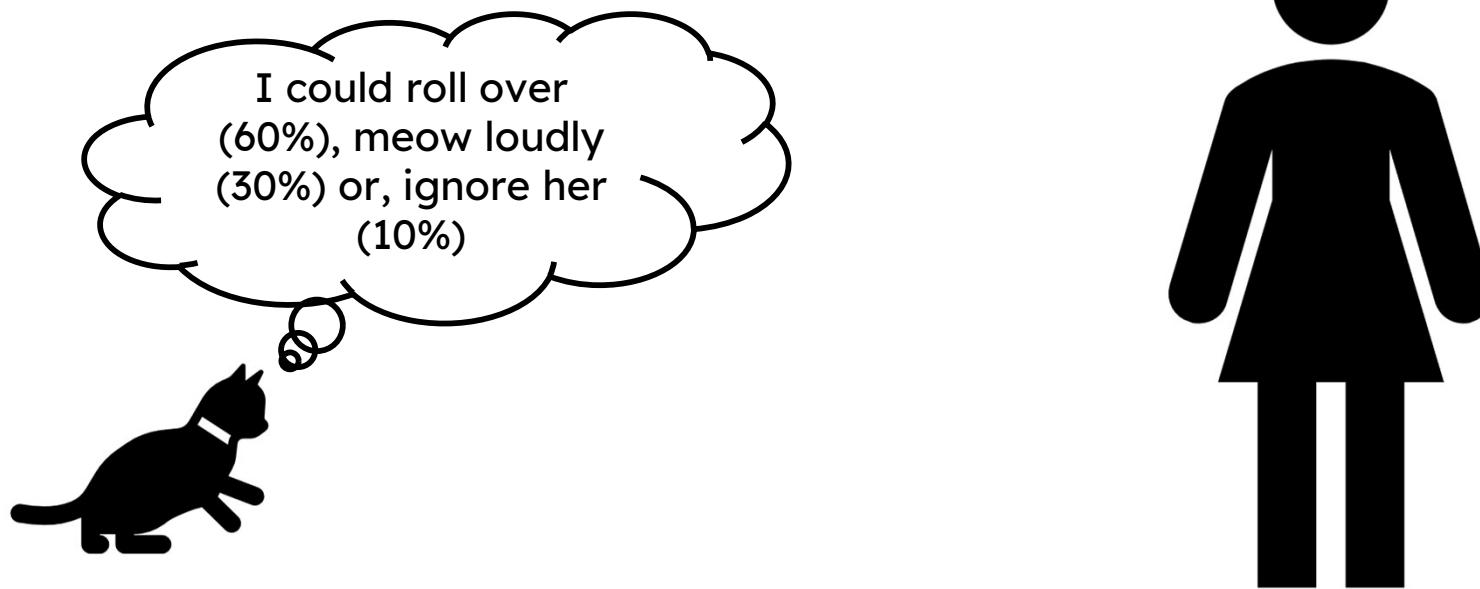
Whatever your cat does next is an **action**.



# Terminology



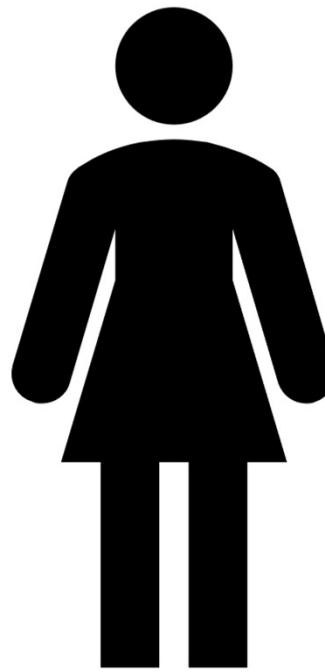
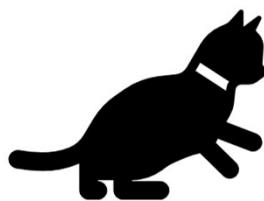
Whatever your cat does next is an action. It will sample its next action from its **policy**.



# Terminology



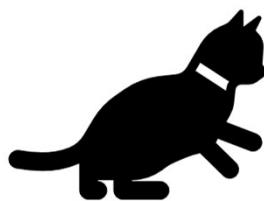
Your cat makes the **action** of rolling over.



# Terminology



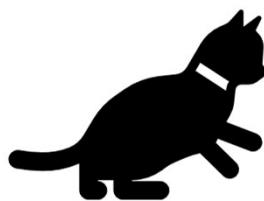
Cat receives a positive **reward**.  
It realizes that next time, it  
should roll over with higher likelihood.



# Terminology



Your cat will take a series of actions: a **trajectory**.



# Terminology - recap



Agent	The learner or decision maker
Environment	The (contained) area the agent makes decisions in
State	The representation of the current situation
Action	The choice the agent makes
Reward	The feedback the agent receives
Policy	The current mapping from states to actions
Trajectory	A set of actions

# Terminology – draw parallels to language



Agent	The learner or decision maker	Language model
Environment	The (contained) area the agent makes decisions in	Language
State	The representation of the current situation	The current context of the model ( $s_{<t}$ )
Action	The choice the agent makes	The token $s_t$ generated at index $t$
Reward	The feedback the agent receives	Outcome or process reward
Policy	The current mapping from states to actions	Language model
Trajectory	A set of actions	Also, $s_{<t}$

# Bellman Equations



Two things that make RL work (and difficult...):

1. Q function
2. Value function

# Bellman Equations

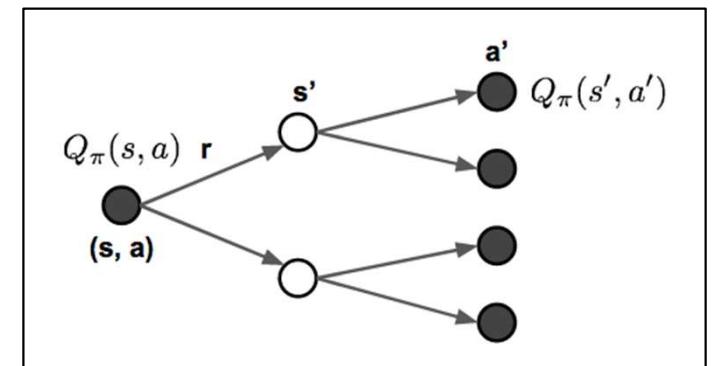


Two things that make RL work (and difficult...):

## 1. Q function

- What is the expected reward of taking action  $a$  in state  $s$ ?
- $Q_\pi(s, a)$ : the current reward plus the expected reward of trajectory
- $$Q_\pi(s, a) = r(s, a) + \gamma \sum_{s'} P(s'|a, s) \sum_{a'} \pi(a'|s') Q_\pi(s', a')$$

## 2. Value function



<https://lilianweng.github.io/posts/2018-02-19-rl-overview/>

# Bellman Equations



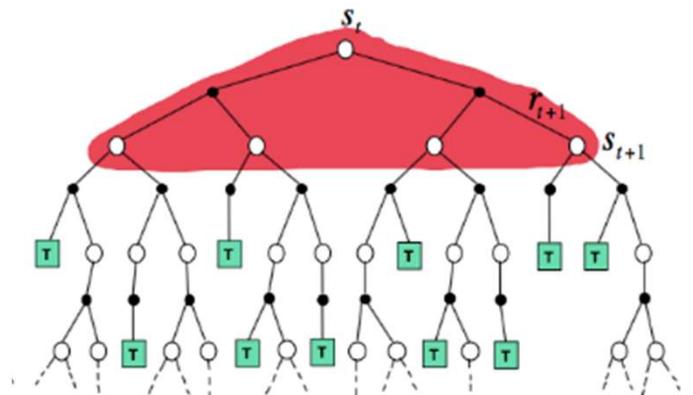
Two things that make RL work (and difficult...):

## 1. Q function

- What is the expected reward of taking action  $a$  in state  $s$ ?
- $Q_\pi(s, a)$ : the current reward plus the expected reward of trajectory
- $$Q_\pi(s, a) = r(s, a) + \gamma \sum_{s'} P(s'|a, s) \sum_{a'} \pi(a'|s') Q_\pi(s', a')$$

## 2. Value function

- What is the expected reward of being in state  $s$ ?
- $V_\pi(s)$ : the average reward of the trajectory
- $$V_\pi(s) = \sum_a \pi(a|s) [r(s, a) + \gamma \sum_{s'} P(s'|a, s) V_\pi(s')]$$



<https://davidstarsilver.wordpress.com/wp-content/uploads/2025/04/lecture-4-model-free-prediction-.pdf>

# Policy Gradient



- In RL we want to maximize the expected return:

$$J(\pi) = E_{\tau \sim \pi} [R(\tau)]$$

- The policy gradient comes out to be:

$$\nabla_{\theta} J(\pi_{\theta}) = E_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) R(\tau) \right]$$

- $\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$  → where do I go to make this action **more likely**?
- $R(\tau)$  → how **good** is this trajectory?

# Policy Gradient



- In RL we want to maximize the expected return:

$$J(\pi) = E_{\tau \sim \pi} [R(\tau)]$$

- The policy gradient comes out to be:

$$\nabla_{\theta} J(\pi_{\theta}) = E_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) R(\tau) \right]$$

- $\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$  → where do I go to make this action **more likely**?
- $R(\tau)$  → how **good** is this trajectory?
- **Intuition check!!** If there are **good** actions with **low probability**, the policy gradient will be...?
  - Bigger or smaller?

# Agenda



## I. Basics

- I. Terminology
- II. Bellman equations
- III. Policy gradient

## II. Common RL Algorithms

- I. PPO
- II. DPO
- III. GRPO

## III. RL Applications

# Vanilla Policy Optimization



- Receive a reward at the end of the entire trajectory

$$L_{VPG} = -E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t | s_t) R(\tau) \right]$$

- Once you and your cat are done playing for an hour, the cat receives a treat/no treat

# Vanilla Policy Optimization



- Receive a reward at the end of the entire trajectory

$$L_{VPG} = -E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t | s_t) R(\tau) \right]$$

- Once you and your cat are done playing for an hour, the cat receives a treat/no treat
- Problems:
  - Rewards are **very sparse!**
  - Can we give **rewards for each action** instead...?

# Proximal Policy Optimization (PPO)



- Approximate the reward at each step + future steps:

$$L_{PPO} = -E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T r_t(\theta) A(s_t, a_t), \right]$$

- $A(s_t, a_t)$ : advantage function
  - Tells you how much better  $a_t$  is for  $s_t$  compared to other actions
- $r_t(\theta)$  is NOT reward, it is a ratio:  $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$ 
  - Measures how drastically the new policy is changing compared to the old one

# Proximal Policy Optimization (PPO)



- Approximate the reward at each step + future steps:

$$L_{PPO} = -E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T \min(r_t(\theta)A(s_t, a_t), \text{CLIP}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)A(s_t, a_t)) \right]$$

- $A(s_t, a_t)$ : advantage function
  - Tells you how much better  $a_t$  is for  $s_t$  compared to other actions
- $r_t(\theta)$  is NOT reward, it is a ratio:  $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$ 
  - Measures how drastically the new policy is changing compared to the old one
  - Clipping ensures the model doesn't stray too far from the previous model

# Proximal Policy Optimization (PPO)



- Approximate the reward at each step + future steps:

$$L_{PPO} = -E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T \min(r_t(\theta)A(s_t, a_t), CLIP(r_t(\theta), 1 - \epsilon, 1 + \epsilon)A(s_t, a_t)) \right]$$

- $A(s_t, a_t)$ : advantage function
  - Tells you how much better  $a_t$  is for  $s_t$  compared to other actions
- $r_t(\theta)$  is NOT reward, it is a ratio:  $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$ 
  - Measures how drastically the new policy is changing compared to the old one
  - Clipping ensures the model doesn't stray too far from the previous model

# Advantage function...?



$$L_{PPO} = -E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T r_t(\theta) A(s_t, a_t) \right]$$

- $A(s_t, a_t) = Q(s_t, a_t) - V(s_t)$ 
  - The expected return of taking action  $a_t$  minus the expected return of being in state  $s_t$
  - Intuitively:
    - How much better is action  $a_t$  at  $s_t$  compared to other actions
    - For a given trajectory, what if I took action  $a_t$  at  $s_t$  instead? Would it be better?
  - **Intuition check!!** What does it mean when the advantage is...
    - Negative?
    - Zero?

# Advantage function...?



$$L_{PPO} = -E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T r_t(\theta) A(s_t, a_t) \right]$$

- $A(s_t, a_t) = Q(s_t, a_t) - V(s_t)$ 
  - The expected return of taking action  $a_t$  minus the expected return of being in state  $s_t$
  - $Q(s_t, a_t)$  is given (“yay!” when the cat rolls over)
  - $V(s_t)$  is estimated with a neural network
    - This could be expensive if the state space gets really large

# Advantage function...?



$$L_{PPO} = -E_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^T r_t(\theta) A(s_t, a_t) \right]$$

- $A(s_t, a_t) = Q(s_t, a_t) - V(s_t)$ 
  - The expected return of taking action  $a_t$  minus the expected return of being in state  $s_t$
  - $Q(s_t, a_t)$  is given (“yay!” when the cat rolls over)
  - $V(s_t)$  is estimated with a neural network
    - This could be **expensive** if the state space gets really large
- Can the model autonomously figure out which actions are better than others?
  - Preference pairs
  - Group optimization

# Direct Preference Optimization (DPO)



ChatGPT 4o ▾

What's your answer?

A

Which response do you prefer?  
Your choice will help make ChatGPT better.

Response 1

Correct.

**Explanation:** The main categories of Microsoft 365 plans are Home, Business, and Enterprise.

**Next Question:**

**Question 4:** Which Microsoft 365 Enterprise plan includes advanced security and compliance features?

- A. E1
- B. E2
- C. E3
- D. E5

What's your answer?

Response 2

Correct.

**Explanation:** The main categories of Microsoft 365 plans are Home, Business, and Enterprise.

**Next Question:**

**Question 4:** Which Microsoft 365 Enterprise plan includes advanced security and compliance features?

- A. E1
- B. E2
- C. E3
- D. E5

What's your answer?

Message ChatGPT

ChatGPT can make mistakes. Check important info.

# Direct Preference Optimization (DPO)



- Alternative to RLHF:

$$L_{DPO} = -E_{x,y^+,y^-}[\log \sigma(\beta \frac{\pi_\theta(y^+|x)}{\pi_{\theta_{ref}}(y^+|x)} - \beta \frac{\pi_\theta(y^-|x)}{\pi_{\theta_{ref}}(y^-|x)})]$$

- $x, y^+, y^-$ : input, desired response, undesired response
- $\beta$  controls how far  $\pi_\theta$  can move from  $\pi_{\theta_{ref}}$
- Move in the direction of favoring  $y^+$  and away from favoring  $y^-$ 
  - Your cat learns to  $y^+ = \text{"meow at you"}$  ( $y^- = \text{"jump at you"}$ ) when you come home

# Group Relative Policy Optimization (GRPO)



ChatGPT 4o ▾



What's your answer?

A

Which response do you prefer?

Your choice will help make ChatGPT better.

## Response 1

Response 1

Correct.

**Explanation:** The main categories of Microsoft 365 plans are Home, Business, and Enterprise.

Next Question:

**Question 4:** Which Microsoft 365 Enterprise plan includes advanced security and compliance features?

- A. E1
- B. E2
- C. E3
- D. E5

What's your answer?

## Response 2

Response 2

Correct.

**Explanation:** The main categories of Microsoft 365 plans are Home, Business, and Enterprise.

Next Question:

**Question 4:** Which Microsoft 365 Enterprise plan includes advanced security and compliance features?

- A. E1
- B. E2
- C. E3
- D. E5

What's your answer?

## Response n-1

Response 1

Correct.

**Explanation:** The main categories of Microsoft 365 plans are Home, Business, and Enterprise.

Next Question:

**Question 4:** Which Microsoft 365 Enterprise plan includes advanced security and compliance features?

- A. E1
- B. E2
- C. E3
- D. E5

What's your answer?

## Response n

Response 2

Correct.

**Explanation:** The main categories of Microsoft 365 plans are Home, Business, and Enterprise.

Next Question:

**Question 4:** Which Microsoft 365 Enterprise plan includes advanced security and compliance features?

- A. E1
- B. E2
- C. E3
- D. E5

What's your answer?

Message ChatGPT



ChatGPT can make mistakes. Check important info.

# Group Relative Policy Optimization (GRPO)



- PPO-style loss over a group of responses

$$L_{GRPO} = \frac{1}{G} \sum_{i=1}^G L_{PPO}$$

- $\tilde{A}(s_t, a_t) = R_t - \bar{R}$ 
  - Samples  $G$  completions
  - Scores each of them
  - Advantage of completion  $i$  is the difference in the reward of  $i$  and the mean reward

# Group Relative Policy Optimization (GRPO)



- PPO-style loss over a group of responses

$$L_{GRPO} = \frac{1}{G} \sum_{i=1}^G \sum_{t=0}^T \min \left( r_t(\theta) \tilde{A}(s_t, a_t), CLIP(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \tilde{A}(s_t, a_t) \right)$$

- $\tilde{A}(s_t, a_t) = R_t - \bar{R}$ 
  - Samples  $G$  completions
  - Scores each of them
  - Advantage of completion  $i$  is the difference in the reward of  $i$  and the mean reward

# Recap



## PPO:

- Improves the RL policy without deviating too much (ensuring no performance degradation)

## DPO:

- Eliminates the need for an explicit value function (expected rewards)
- Advantage based on preference pairs

## GRPO:

- Also eliminates the need for an explicit value function
- Advantage based on relative rewards within a group

# Agenda



## I. Basics

- I. Terminology
- II. Bellman equations
- III. Policy gradient

## II. Common RL Algorithms

- I. PPO
- II. DPO
- III. GRPO

## III. RL Applications

- I. PPO: RLHF
- II. DPO: Reasoning
- III. GRPO: Self-Correction

---

## Training language models to follow instructions with human feedback

---

Long Ouyang\*   Jeff Wu\*   Xu Jiang\*   Diogo Almeida\*   Carroll L. Wainwright\*

Pamela Mishkin\*   Chong Zhang   Sandhini Agarwal   Katarina Slama   Alex Ray

John Schulman   Jacob Hilton   Fraser Kelton   Luke Miller   Maddie Simens

Amanda Askell<sup>†</sup>              Peter Welinder              Paul Christiano\*<sup>†</sup>

Jan Leike\*              Ryan Lowe\*

OpenAI

# RLHF



Step 1

**Collect demonstration data, and train a supervised policy.**

A prompt is sampled from our prompt dataset.

Explain the moon landing to a 6 year old

A labeler demonstrates the desired output behavior.



Some people went to the moon...

This data is used to fine-tune GPT-3 with supervised learning.



Step 2

**Collect comparison data, and train a reward model.**

A prompt and several model outputs are sampled.

Explain the moon landing to a 6 year old

A Explain gravity...  
B Explain war...  
C Moon is natural satellite of...  
D People went to the moon...

A labeler ranks the outputs from best to worst.

D > C > A = B

This data is used to train our reward model.

RM

D > C > A = B

Step 3

**Optimize a policy against the reward model using reinforcement learning.**

A new prompt is sampled from the dataset.

Write a story about frogs

The policy generates an output.



PPO



Once upon a time...



The reward model calculates a reward for the output.

$r_k$

The reward is used to update the policy using PPO.

## Step 1

**Collect demonstration data, and train a supervised policy.**

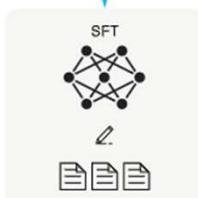
A prompt is sampled from our prompt dataset.



A labeler demonstrates the desired output behavior.



This data is used to fine-tune GPT-3 with supervised learning.



## Step 2

**Collect comparison data, and train a reward model.**

A prompt and several model outputs are sampled.



A labeler ranks the outputs from best to worst.



This data is used to train our reward model.



## Step 3

**Optimize a policy against the reward model using reinforcement learning.**

A new prompt is sampled from the dataset.



The policy generates an output.



Once upon a time...



The reward model calculates a reward for the output.

The reward is used to update the policy using PPO.

$r_k$

## Step 1

### Collect demonstration data, and train a supervised policy.

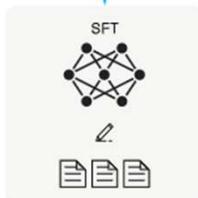
A prompt is sampled from our prompt dataset.



A labeler demonstrates the desired output behavior.



This data is used to fine-tune GPT-3 with supervised learning.



## Step 2

### Collect comparison data, and train a reward model.

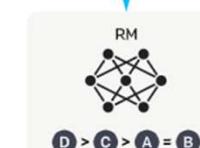
A prompt and several model outputs are sampled.



A labeler ranks the outputs from best to worst.



This data is used to train our reward model.



## Step 3

### Optimize a policy against the reward model using reinforcement learning.

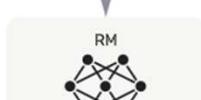
A new prompt is sampled from the dataset.



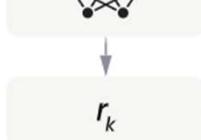
The policy generates an output.



The reward model calculates a reward for the output.



The reward is used to update the policy using PPO.



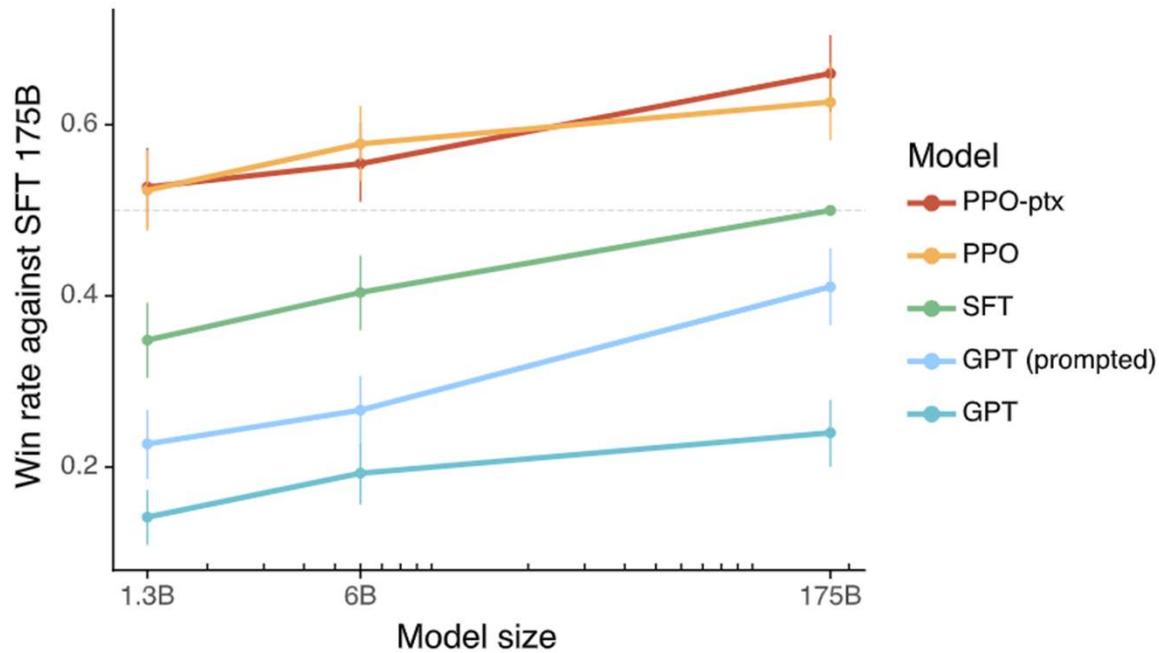


Figure 1: Human evaluations of various models on our API prompt distribution, evaluated by how often outputs from each model were preferred to those from the 175B SFT model. Our InstructGPT models (PPO-ptx) as well as its variant trained without pretraining mix (PPO) significantly outperform the GPT-3 baselines (GPT, GPT prompted); outputs from our 1.3B PPO-ptx model are preferred to those from the 175B GPT-3. Error bars throughout the paper are 95% confidence intervals.

---

# Chain of Preference Optimization: Improving Chain-of-Thought Reasoning in LLMs

---

Xuan Zhang<sup>\*12</sup>, Chao Du<sup>†1</sup>, Tianyu Pang<sup>1</sup>, Qian Liu<sup>1</sup>, Wei Gao<sup>2</sup>, Min Lin<sup>1</sup>

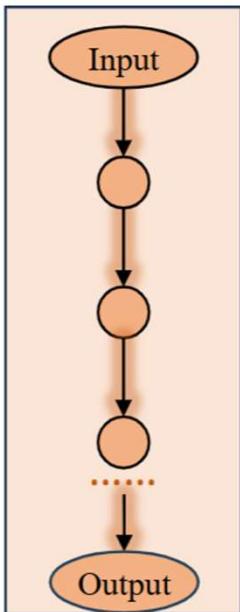
<sup>1</sup>Sea AI Lab, Singapore

<sup>2</sup>School of Computing and Information Systems, Singapore Management University  
xuanzhang.2020@phdcs.smu.edu.sg; weigao@smu.edu.sg;  
{duchao, liuqian, tianyupang, linmin}@sea.com

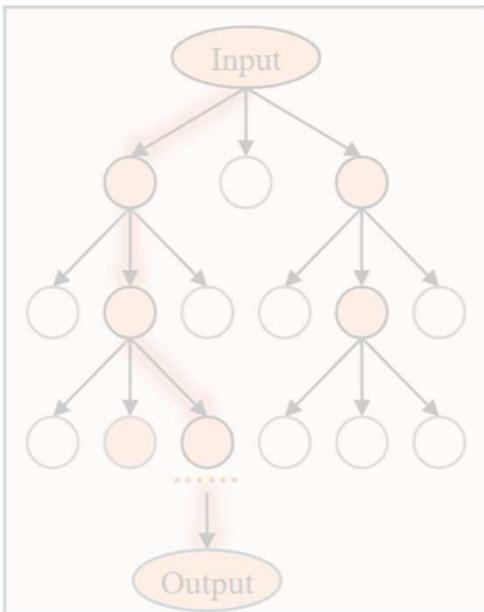
# Chain of Preference Optimization



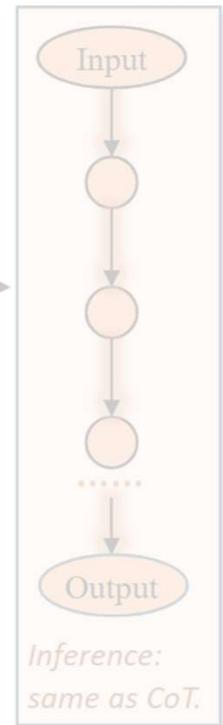
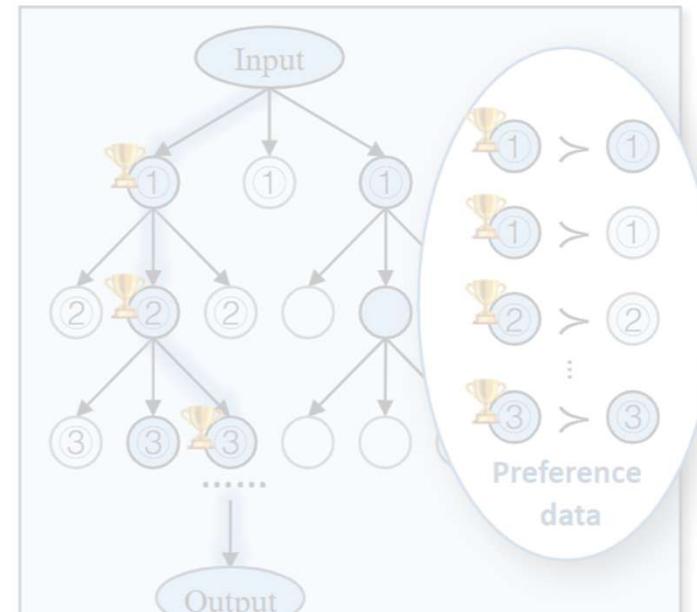
(a) Chain of Thought (CoT)



(b) Tree of Thought (ToT)



(c) Chain of Preference Optimization (CPO)



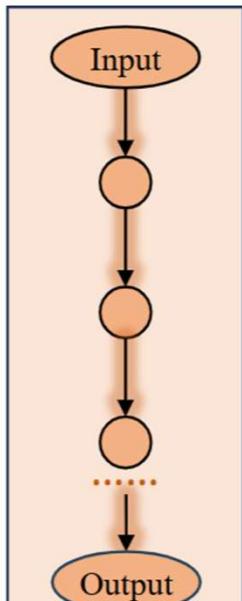
Selected thoughts.  
 Pruned thoughts.

Training.  
 Inference.

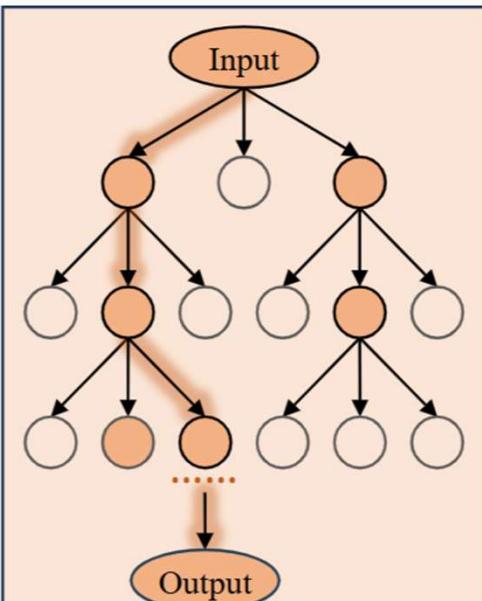
# Chain of Preference Optimization



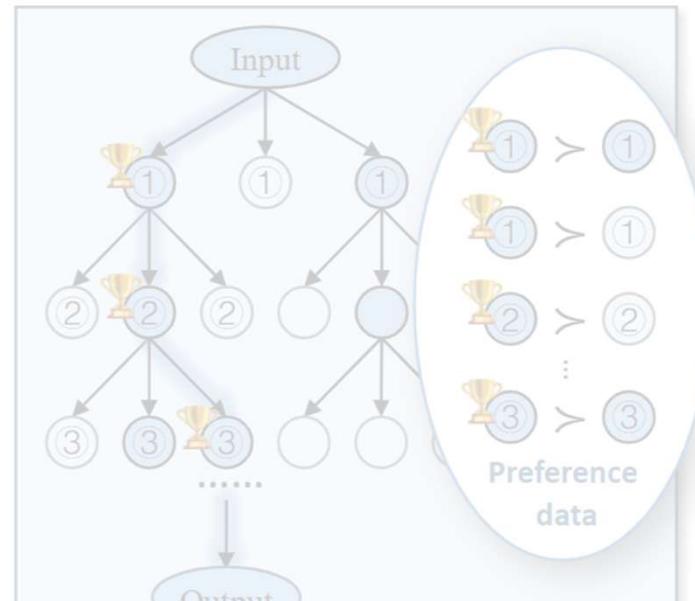
(a) Chain of Thought (CoT)



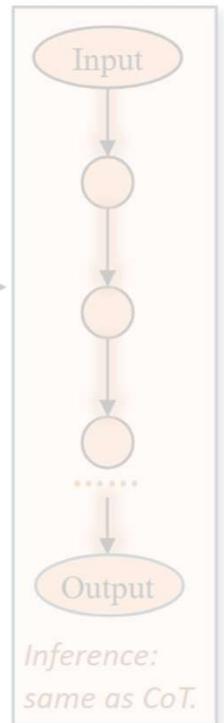
(b) Tree of Thought (ToT)



(c) Chain of Preference Optimization (CPO)



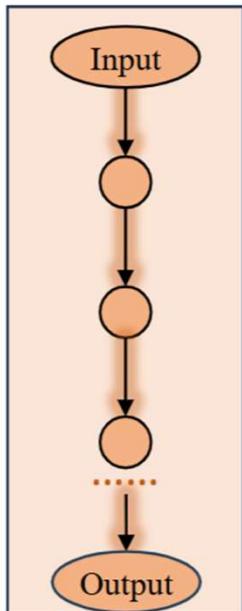
*Collecting preference data (of each step)  
through ToT and then tuning LLM via DPO.*



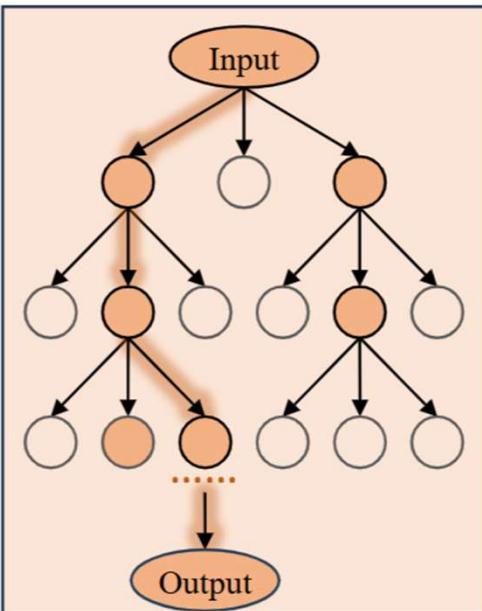
# Chain of Preference Optimization



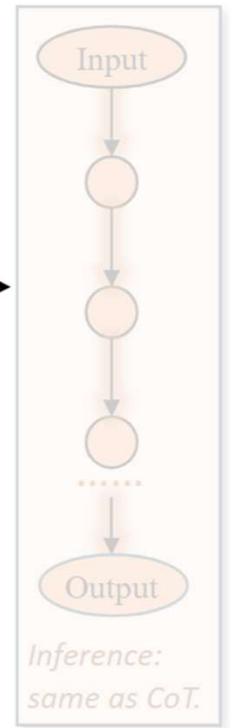
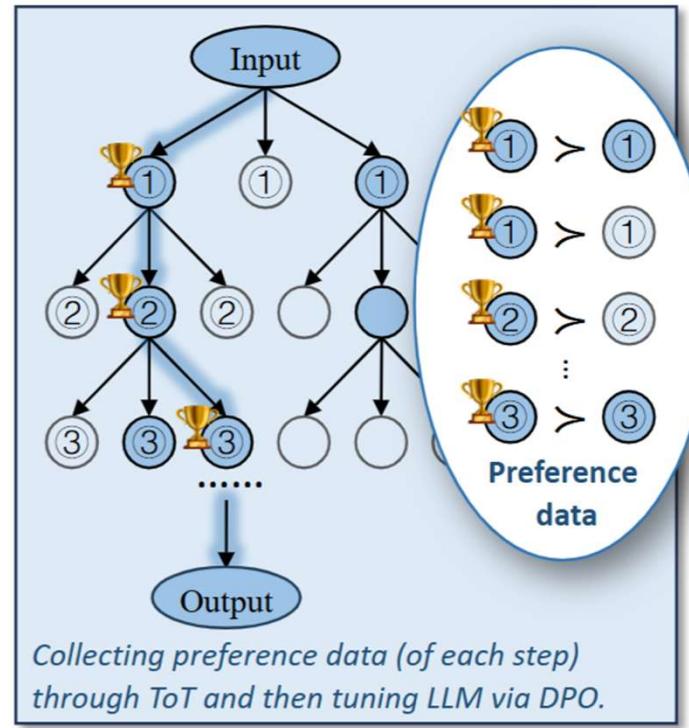
(a) Chain of Thought (CoT)



(b) Tree of Thought (ToT)



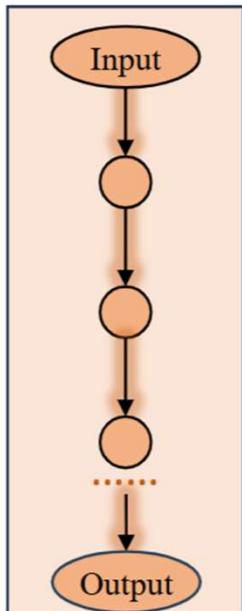
(c) Chain of Preference Optimization (CPO)



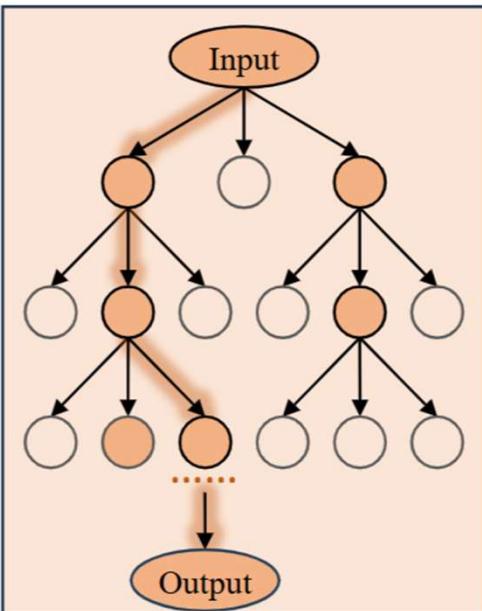
# Chain of Preference Optimization



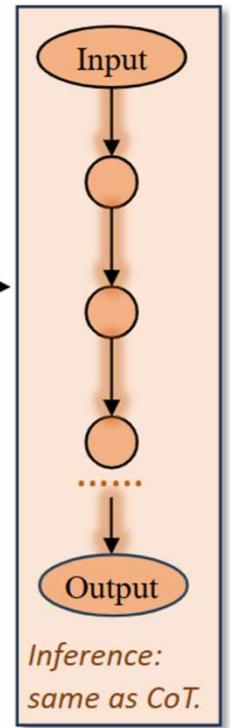
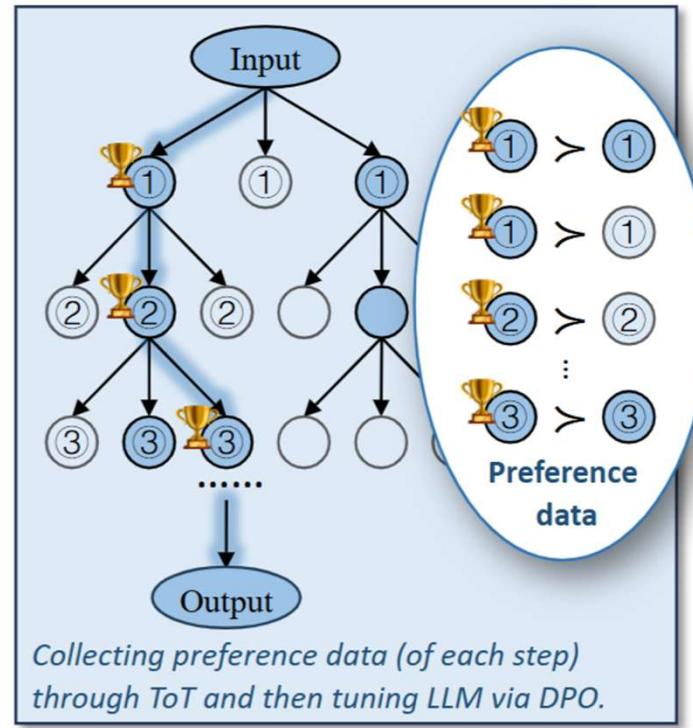
(a) Chain of Thought (CoT)



(b) Tree of Thought (ToT)



(c) Chain of Preference Optimization (CPO)



# Chain of Preference Optimization



Table 1: Experimental results for ToT, CoT, TS-SFT, and our proposed CPO across complex task including question answering, fact verification, and arithmetic reasoning are presented. \* mean significantly better than the best baseline (TS-SFT) with  $p < 0.01$ . **Bold** denotes the best method and the second best if the top method is ToT.

		ToT [8]		CoT [1]		TS-SFT [11]		CPO (ours)	
		Acc. (%)↑	Latency (s/ins.)↓	Acc. (%)↑	Latency (s/ins.)↓	Acc. (%)↑	Latency (s/ins.)↓	Acc. (%)↑	Latency (s/ins.)↓
<b>LLaMA2-7B</b>									
<i>Question Answering</i>	Bam.	<b>33.6</b>	1168.4	29.6	37.2	30.4	36.5	<b>32.0*</b>	38.2
	2Wiki.	28.6	847.6	26.3	35.7	27.6	35.5	<b>29.7*</b>	35.7
	Hot.	23.0	1100.7	21.0	45.5	22.7	44.8	<b>24.0*</b>	41.1
<i>Fact Verification</i>	FVR.	47.3	2087.1	45.8	33.8	47.5	34.0	<b>53.2*</b>	36.8
	FVRS.	47.5	2539.5	44.3	40.6	46.0	40.4	<b>49.0*</b>	41.2
	Vita.	50.7	2639.3	47.3	35.9	51.0	40.1	<b>52.7*</b>	40.1
<i>Arithmetic</i>	SVA.	42.7	1861.1	37.7	33.3	43.1	30.2	<b>46.0*</b>	32.1
<i>Average Performance</i>		39.1	1749.1	36.0	37.4	38.3	37.4	<b>40.9*</b>	37.9

---

# Maximizing Confidence Alone Improves Reasoning

---

**Mihir Prabhudesai\***  
Carnegie Mellon University

**Lili Chen\***  
Carnegie Mellon University

**Alex Ippoliti\***  
Carnegie Mellon University

**Katerina Fragkiadaki**  
Carnegie Mellon University

**Hao Liu**  
Carnegie Mellon University

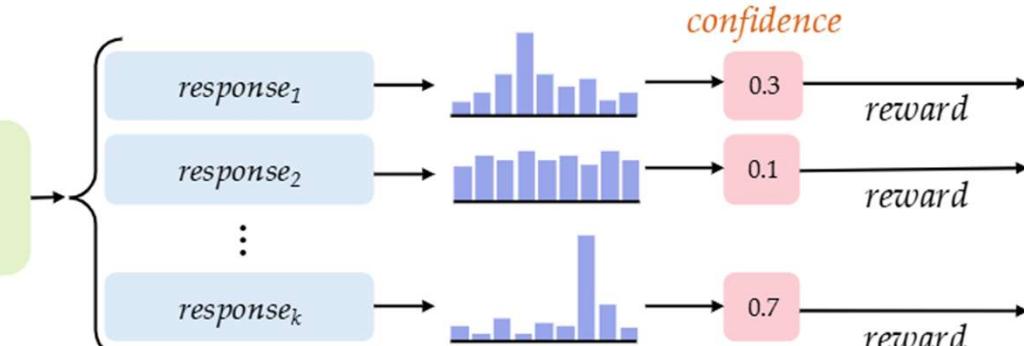
**Deepak Pathak**  
Carnegie Mellon University

# Maximizing Confidence

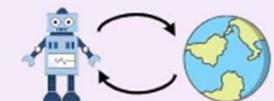


Reinforcement Learning via  
Entropy Minimization

*prompt* → language model



Reinforcement  
Learning



# Maximizing Confidence



Give confidence (negative entropy) as a reward to GRPO:

- Multiply the probability of a token with its log probability
  - $p_t(v) \log p_t(v)$  where  $v$  is a token at position  $t$
- Sum across all tokens in the vocabulary at position  $t$ 
  - $\sum_{v \in V} p_t(v) \log p_t(v)$  where  $V$  is the vocabulary
  - Most of them will be close to 0
- Average across tokens in the sequence:
  - $R(y_{pred}) = \frac{1}{T} \sum_{t=1}^T \sum_{v \in V} p_t(v) \log p_t(v)$
  - This is the final reward function!

# Maximizing Confidence

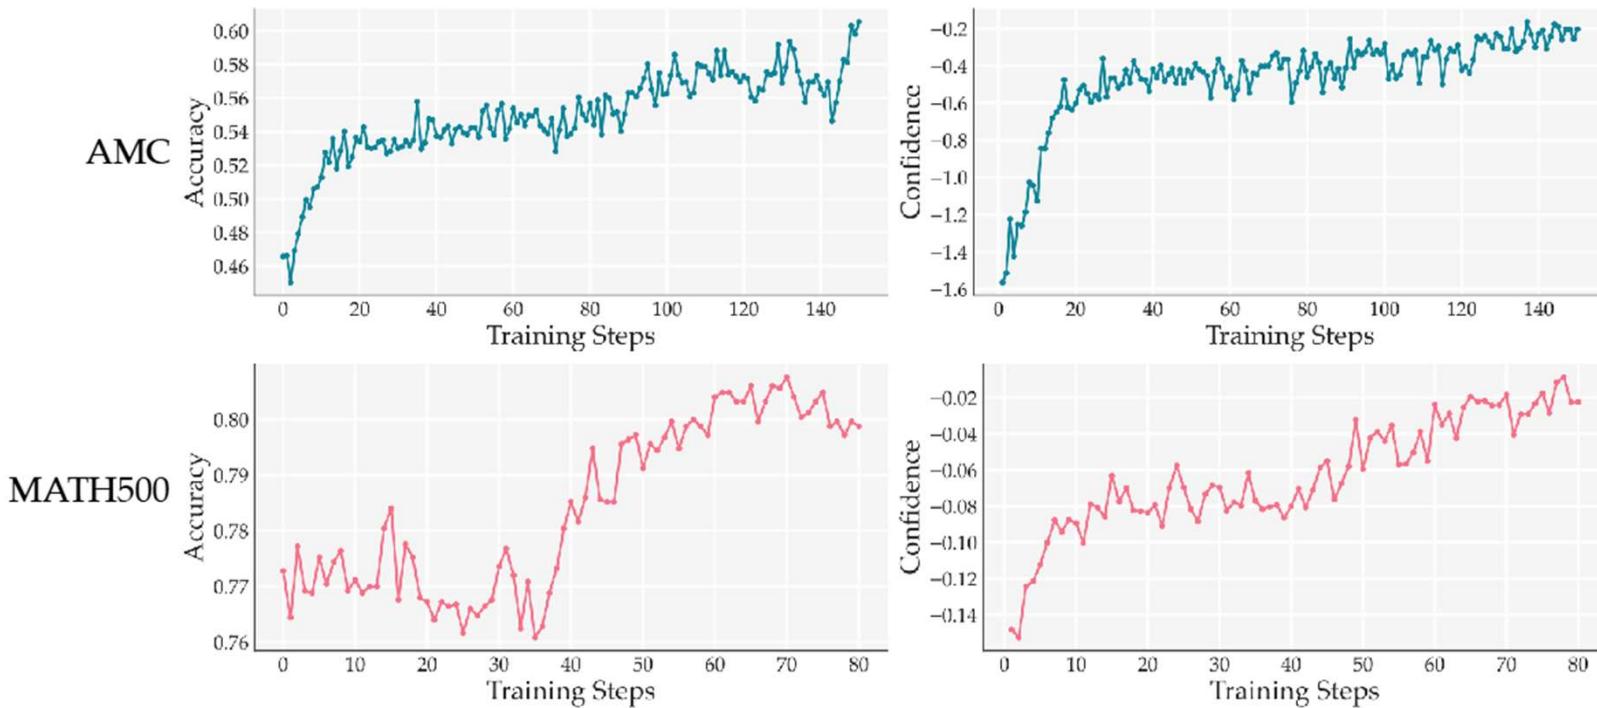


Figure 3: Accuracy and confidence over the course of training. The trends indicate that accuracy and confidence are indeed highly correlated and therefore it is natural to use confidence as a reward.

# Conclusion



# Conclusion



## Welcome to the Era of Experience

David Silver, Richard S. Sutton\*

### Abstract

We stand on the threshold of a new era in artificial intelligence that promises to achieve an unprecedented level of ability. A new generation of agents will acquire superhuman capabilities by learning predominantly from experience. This note explores the key characteristics that will define this upcoming era.

# Conclusion



LLMs are changing a lot of how we use RL:

- **Infinite state and action spaces**
- **Human-data to experience-data**
  - Human asks a question, agent responds, human gives feedback
  - Agent interacts with its environment and receives signals from the environment

# Conclusion



In the era of experience,

- Agents will not solely require on human feedback to improve
- Example:
  - A health and wellness agent prescribes a diet and exercise regimen
  - Human reward: “yes this is working” / “no this isn’t working”
  - Experience reward: heart rate, sleep patterns, blood work results, etc.
- We can achieve agents that are better than humans

# Conclusion



In the era of experience,

- Agents will not solely require on human feedback to improve
- Example:
  - A health and wellness agent prescribes a diet and exercise regimen
  - Human reward: “yes this is working” / “no this isn’t working”
  - Experience reward: heart rate, sleep patterns, blood work results, etc.
- We can achieve agents that are better than humans

My opinion: easier to hype, harder to adopt; we’re living in exciting times!

# Reinforcement Learning in the Era of LLMs

Ishika Agarwal  
CS 546: Advanced NLP

