



# NoSQL Overview & MongoDB Basics

**Abdu Alawini**

University of Illinois at Urbana-Champaign

CS411: Database Systems

**February 8, 2021**

Some slides were adopted from S. Davidson and Z. Ives with permission



# Learning Objectives

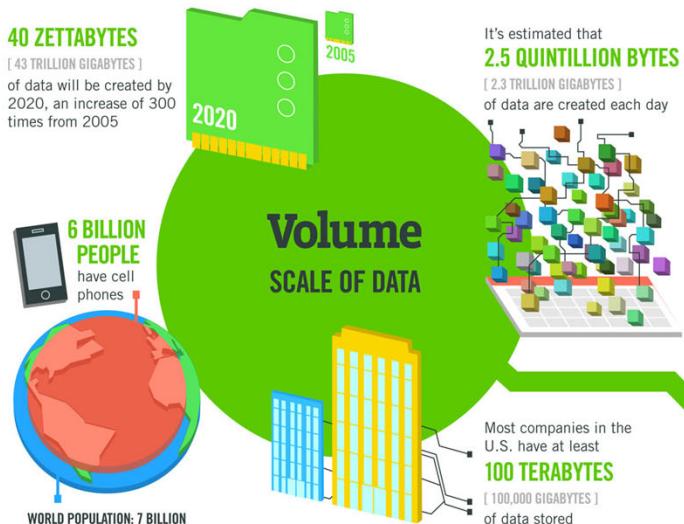
After this lecture, we will:

- Introduce the NoSQL paradigm
- Discuss the trade-offs between relational and non-relational(NoSQL) databases
- Introduce MongoDB, a document-oriented database
- Learn MongoDB simple queries.



# Why NoSQL?

- Databases are no longer one-size-fits-all
- The needs of modern applications do not always match what relational databases provide.
- Every large web platform (e.g. Google, Facebook, LinkedIn) has developed some sort of custom solution to scale.



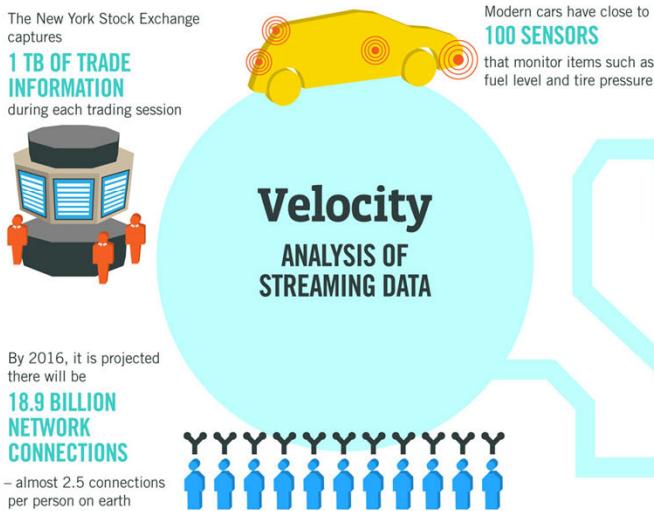
# The FOUR V's of Big Data

From traffic patterns and music downloads to web history and medical records, data is recorded, stored, and analyzed to enable the technology and services that the world relies on every day. But what exactly is big data, and how can these massive amounts of data be used?

As a leader in the sector, IBM data scientists break big data into four dimensions: **Volume**, **Velocity**, **Variety** and **Veracity**.

Depending on the industry and organization, big data encompasses information from multiple internal and external sources such as transactions, social media, enterprise content, sensors and mobile devices. Companies can leverage data to adapt their products and services to better meet customer needs, optimize operations and infrastructure, and find new sources of revenue.

By 2015  
**4.4 MILLION IT JOBS**  
will be created globally to support big data, with 1.9 million in the United States



Sources: McKinsey Global Institute, Twitter, Cisco, Gartner, EMC, SAS, IBM, MEPTEC, QAS

As of 2011, the global size of data in healthcare was estimated to be

**150 EXABYTES**  
[ 161 BILLION GIGABYTES ]

**30 BILLION PIECES OF CONTENT**

are shared on Facebook every month



## Variety DIFFERENT FORMS OF DATA

400 MILLION TWEETS

are sent per day by about 200 million monthly active users



**1 IN 3 BUSINESS LEADERS**

don't trust the information they use to make decisions

**27% OF RESPONDENTS**

in one survey were unsure of how much of their data was inaccurate

## Veracity UNCERTAINTY OF DATA



Poor data quality costs the US economy around  
**\$3.1 TRILLION A YEAR**





## “Big Data” is two problems

- The **analysis** problem
  - How to extract useful info, using modeling, ML and stats.
- The **storage** problem
  - How to store and manipulate huge amounts of data to facilitate fast queries and analysis
- Problems with traditional (relational) storage
  - Not flexible
  - Hard to partition, i.e. place different segments on different machines
- NoSQL solutions address these problems.



## The analysis problem...

- So far, we've focused on the storage problem – flexible schemas. There is also the analysis problem, which requires **scalability**.
- Relational databases typically scale by getting bigger servers
  - Scaling across multiple servers is complicated
- NoSQL solutions are all about scaling across multiple servers (e.g. cloud instances)
  - Data can be automatically distributed across nodes/servers
  - “Map-reduce” spreads computation across a cluster



# Types of NoSQL solutions

- Key-value stores:



- Column-oriented:



- Document:



- Graph: Neo4j





# Relational-NoSQL Trade-offs

Fundamentally, there are several different trade-offs

- Schema vs. no schema
  - Schema → performance, no schema → flexibility but parse overhead  
(can have partial schemas like in XML)
- Replication, data partitioning
  - Replicas mean faster queries, slower (consistent) updates
- Level of abstraction
  - High-level queries – parsing, optimization, etc. – vs. low-level operations
- Consistency
  - What does the database do on concurrent updates, especially when distributed?

# Outline

- ✓ NoSQL Introduction
- MongoDB
- Model and simple queries



# Overview of MongoDB®

- MongoDB is an example of a document-oriented NoSQL solution
- The query language is limited, and oriented around “collection” (relation) at a time processing
- The power of the solution lies in the distributed, parallel nature of query processing
  - Replication and sharding



## MongoDB Data Model

- A MongoDB deployment hosts several databases
  - A **database** holds a set of collections
    - A **collection** holds a set of documents
      - A **document** is a set of key-value pairs

RDBMS	MongoDB
Table	Collection
Row(s)	JSON Document
Index	Index
Join	Embedding & Linking
Partition	Shard
Partition Key	Shard Key

# Basic data types

- Null
- Boolean
- Integer (32- and 64-bit)
- Floating point
- String
- Date
- ObjectId
- Code (JavaScript)
- Array
- Embedded document

# Sample Document

```
mydoc = {  
    _id: 1, Always indexed, automatically assigned unless provided  
    name: { first: "John", last: "Backus" },  
    birthyear: 1924,  
    contribs: [ "Fortran", "ALGOL", "Backus-Naur Form", "FP" ],  
    awards: [ { award_id: "NMSoo1",  
                year: 1975 },  
              { award_id: "TA99",  
                year: 1977 } ]  
}
```

Array of documents



# Core MongoDB operations

- CRUD: *create, read, update, and delete*
- Insert
  - One at a time: `db.people.insert(mydoc)`
  - New (version 3.2):  
`db.collection.insertOne()`,  
`db.collection.insertMany()`
- Delete
  - Documents that match some predicate, e.g. to remove the document in the previous slide:  
`db.people.deleteOne({“_id”: 1})`  
`db.people.deleteMany({birthyear: 1924})`
  - All documents in a collection: `db.people.deleteMany()`
    - The collection still remains, with indexes
  - Remove a collection (faster): `db.people.drop()`



## Core MongoDB operations, cont.

- Update documents in a collection
  - db.collection.updateOne(), db.collection.updateMany()

```
db.people.updateMany( {birthyear: 1924}, {$set: {birthyear: 1925}})
```

```
db.people.updateMany( {birthyear: 1924}, {$set: {type: "Deceased"}})
```

- \$rename operator: change property name.

```
db.people.update({}, { $rename : {"birthyear": "birth" }})
```

# Querying

- Use `find( )` function and a query document
- Ranges, set inclusion, inequalities using `$` conditionals
- Complex queries using `$where` clause
- Queries return a database cursor
- Meta-operations on cursor include skipping some number of results, limiting the number of results returned, sorting results.



## Another sample document

```
d={  
    _id : ObjectId("4c4ba5c0672c685e5e8aabf3"),  
    author : "Kevin",  
    date : new Date("February 2, 2012"),  
    text : "About MongoDB...",  
    birthyear: 1980,  
    tags : [ "tech", "databases" ]  
}
```

```
> db.posts.insert(d)
```



# Find

Return entire collection in posts:

```
db.posts.find()
```

Return posts that match condition (conjunction):

```
db.posts.find({author: "Kevin", birthyear: 1980})
```

```
{ _id : ObjectId("4c4ba5c0672c685e5e8aabf3"), author : "Kevin",  
date : Date("February 2, 2012"), birthyear: 1980,  
text : "About MongoDB...", tags : [ "tech", "databases" ]}
```

# “Pretty” format

- If you want to be able to read the result:

```
db.posts.find({author: "Kevin", birthyear: 1980}).pretty()
```

```
{  
    _id : ObjectId("4c4ba5c0672c685e5e8aabf3"),  
    author : "Kevin",  
    date : Date("February 2, 2012"),  
    birthyear: 1980,  
    text : "About MongoDB...",  
    tags : [ "tech", "databases" ]  
}
```



## Specifying which keys to return

```
db.people.find({}, {name:1, contribs:1})
```

```
{  
  _id: 1,  
  name: { first: "John", last: "Backus" },  
  contribs: [ "Fortran", "ALGOL", "Backus-Naur Form", "FP" ]  
}
```

```
db.people.find({}, {_id: 0, name:1})
```

```
{  
  name: { first: "John", last: "Backus" }  
}
```



# Ranges, Negation, OR-clauses

- Comparison operators: \$lt, \$lte, \$gt, \$gte
  - db.posts.find({birthyear: {\$gte: 1970, \$lte: 1990}})
- Negation: \$ne
  - db.posts.find({birthyear: {\$ne: 1982}})
- Or queries: \$in (single key), \$or (different keys)
  - db.posts.find({birthyear: {\$in: [1982, 1985]}})
  - db.posts.find({\$or: [{birthyear: 1982}, {author: "John"}], name: "abdu"})

# Arrays

- `db.posts.find({tags: "tech"})`
  - Print complete information about posts which are tagged “tech”
- `db.posts.find({tags: {$all: ["tech", "databases"]}}, {author:l, tags:l})`
  - Print author and tags of posts which are tagged with both “tech” and “databases” (among other things)
  - Contrast this with:  
`db.posts.find({tags: ["databases", "tech"]})`



# Querying Embedded Documents

- `db.people.find({"name.first": "John"})`
    - Finds all people with first name John
  - `db.people.find({"name.first": "John", "name.last": "Smith"})`
    - Finds all people with first name John and last name Smith.
    - Contrast with
- `db.people.find({"name": {"first": "John", "last": "Smith"}})`

# Sample Document

```
mydoc = {  
    _id: 1,  
    name: { first: "John", last: "Backus" },  
    birthyear: 1924,  
    contribs: [ "Fortran", "ALGOL", "Backus-Naur Form", "FP" ],  
    awards: [ { award_id: "NMSoo1",  
                year: 1975 },  
              { award_id: "TA99",  
                year: 1977 } ]  
}
```

# Summary

- NoSQL solutions address the needs of modern applications
- MongoDB is an example of a document-oriented solution
  - The query language is oriented around “collection” at a time processing
- The power of many of these solutions lies in the distributed, parallel nature of query processing
  - Replication and sharding