



# **Indexing: Hash Tables**

**Abdu Alawini**

University of Illinois at Urbana-Champaign

CS411: Database Systems

# Learning Objectives

After this lecture, you should be able to:

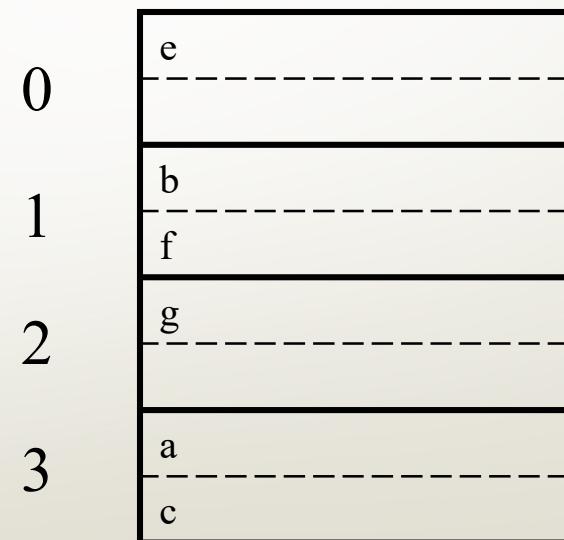
- Describe how to search, insert and delete keys from
  - Secondary storage Hash Table (HT)
  - Extensible HT
  - Linear HT

# Hash Tables

- Secondary storage hash tables are much like main memory ones
- Recall basics:
  - There are  $B$  *buckets*
  - A hash function  $h(k)$  maps a key  $k$  to  $\{0, 1, \dots, B-1\}$
  - Store in bucket  $h(k)$  a pointer to record with key  $k$
- Secondary storage: bucket = block
  - Store in the block of bucket  $h(k)$  any record with key  $k$
  - use overflow blocks when needed

# Hash Table Example

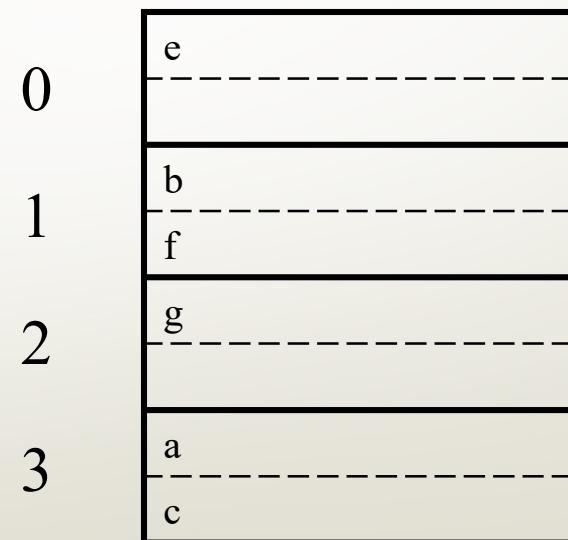
- Assume 1 bucket (block) stores 2 records
- $h(e)=0$
- $h(b)=h(f)=1$
- $h(g)=2$
- $h(a)=h(c)=3$



# Searching in a Hash Table

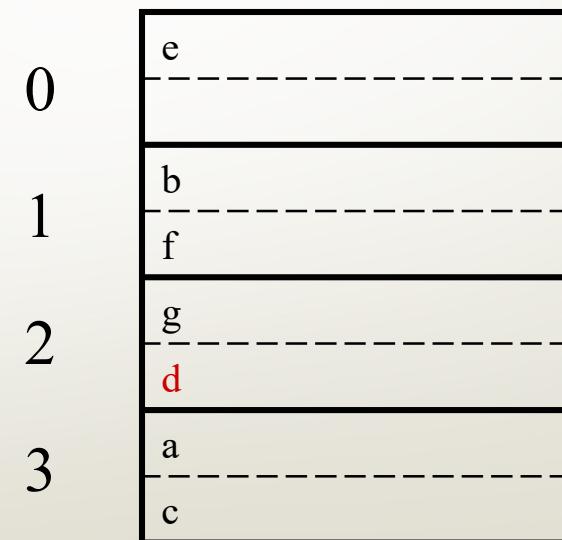
- Search for a:
- Compute  $h(a)=3$
- Read bucket (block) 3
- 1 disk access

Main memory may have an array of pointers (to buckets) accessible by bucket number.



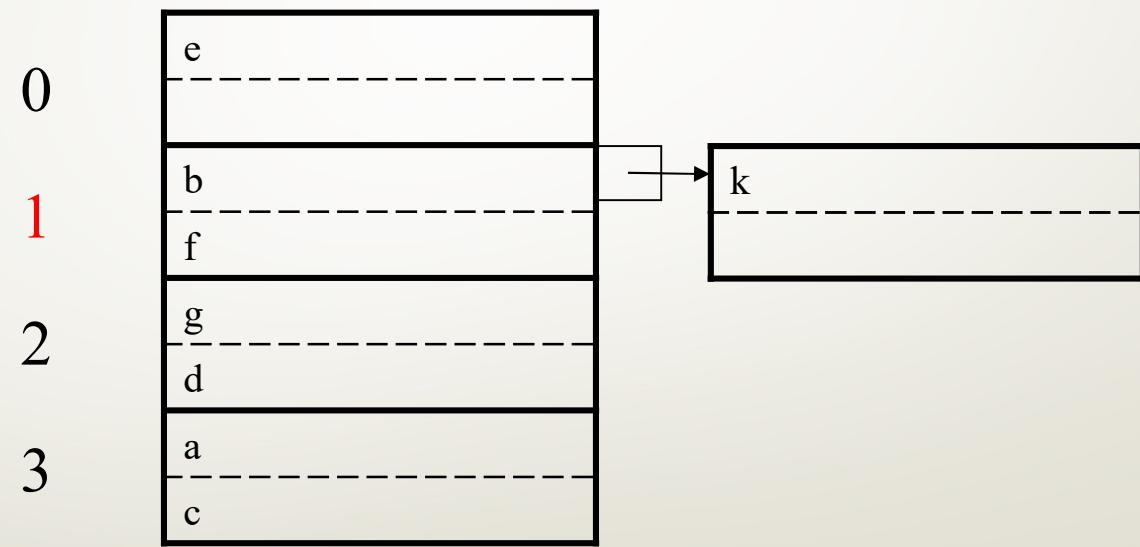
# Insertion in Hash Table

- Place in right bucket (block), if space
- E.g.  $h(d)=2$



# Insertion in Hash Table

- Create overflow block, if no space
- E.g.  $h(k)=1$



More over-flow  
blocks may be needed

# Hash Table Performance

- Fixed number of buckets
- Excellent, if no overflow blocks
- Degrades considerably when there are many overflow blocks.
  - Might need to go through a chain of overflow blocks

*Can improve this by allowing the number of buckets to grow*

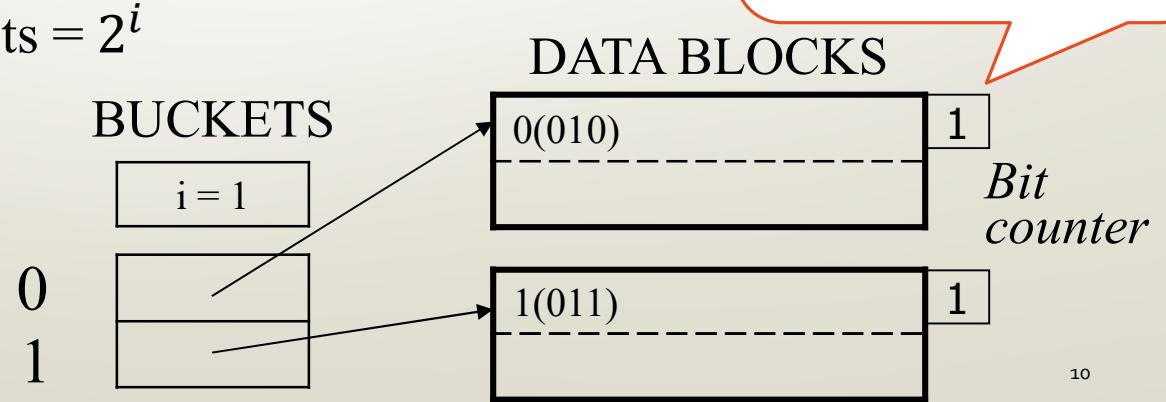
# Outline

- Hash Tables
  - ✓ Secondary storage HT
  - Extensible HT
  - Linear HT

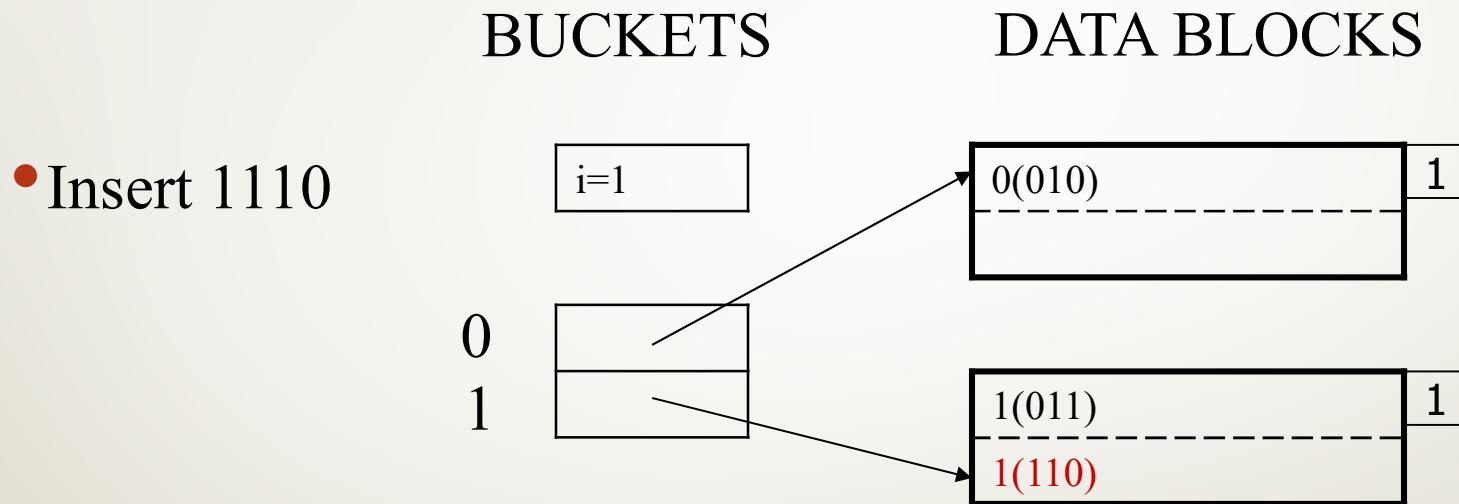
# Extensible Hash Table

- Array of pointers to blocks instead of array of blocks
- Size of array is allowed to grow. 2x size when it grows
- Don't need a block per bucket. Sparse buckets share a block
- Hash function returns k-bit integers (e.g., k=32)
  - Only use the first  $i \ll k$  bits to determine bucket
  - Number of buckets =  $2^i$

Bit counter on each block indicates how much bits are used for that block



# Insertion in Extensible Hash Table

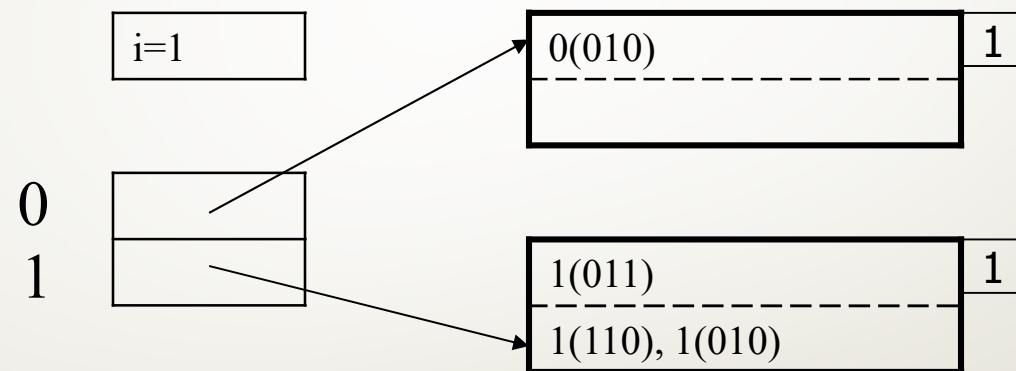


# Insertion in Extensible Hash Table

- Now insert 1010

BUCKETS

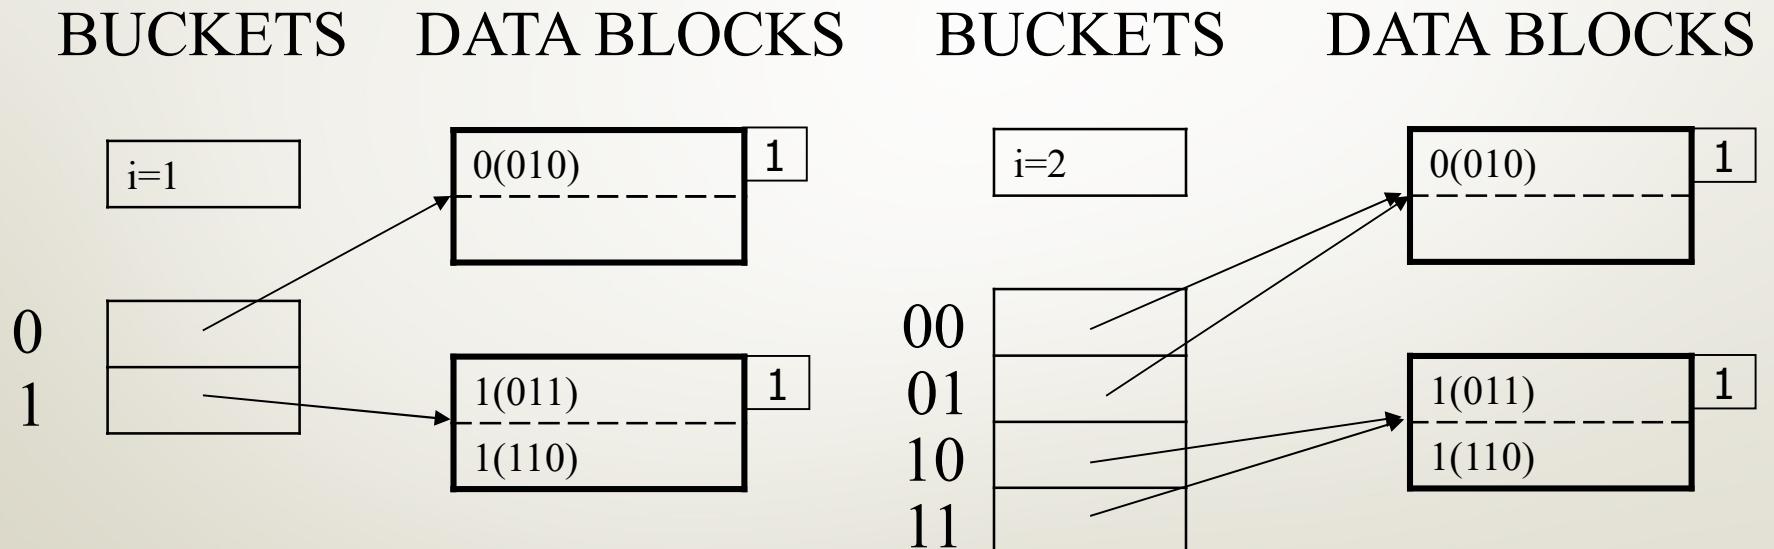
DATA BLOCKS



- Need to split block and extend bucket array
- $i$  becomes 2: done in two steps

# Insertion in Extensible Hash Table

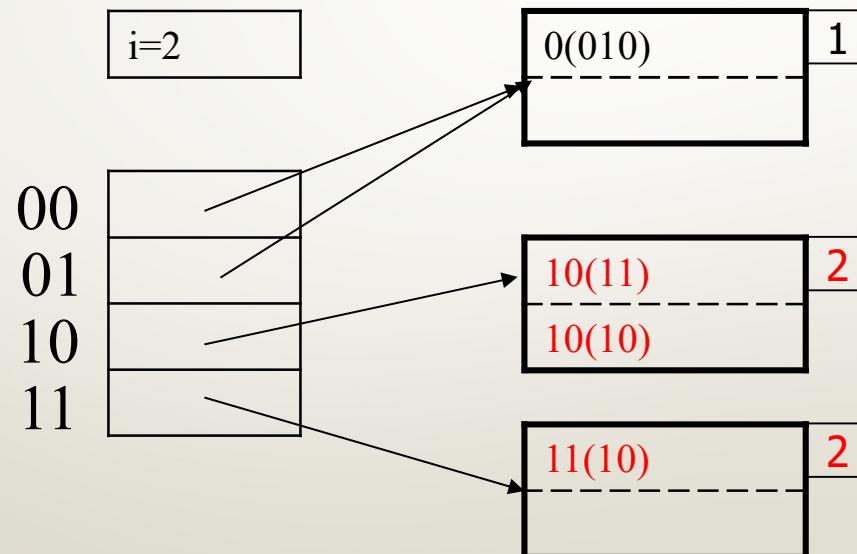
Step 1: Extend the buckets



# Insertion in Extensible Hash Table

Step 2: Now try to insert 1010

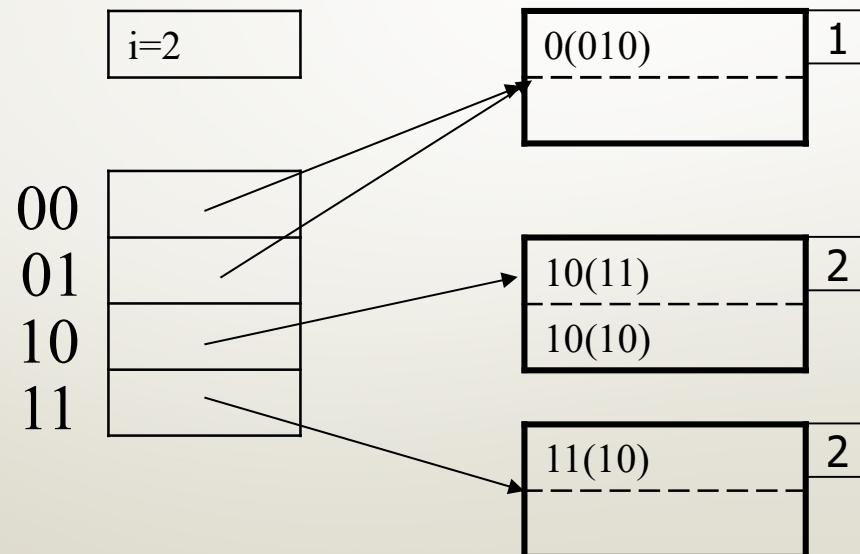
BUCKETS      DATA BLOCKS



# Insertion in Extensible Hash Table

- Now insert 0000: where would it go? Then 0101?
- Need to split block, but not bucket array

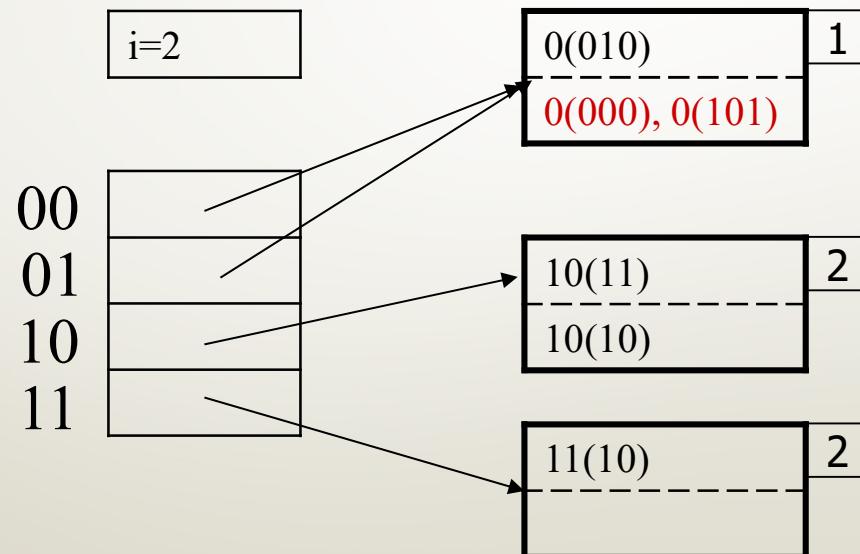
BUCKETS      DATA BLOCKS



# Insertion in Extensible Hash Table

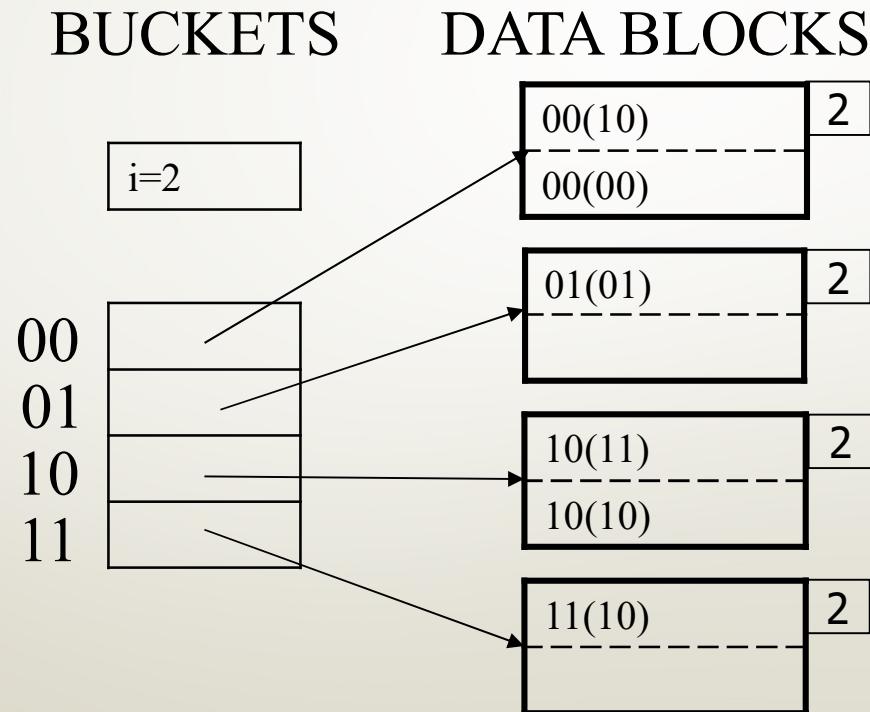
- Now insert 0000: where would it go? Then 0101?
- Need to split block, but not bucket array

BUCKETS      DATA BLOCKS



# Insertion in Extensible Hash Table

- Now insert 0000: where would it go? Then 0101?
- Need to split block, but not bucket array



# Performance: Extensible Hash Table

- No overflow blocks: access always one read for distinct keys
- BUT:
  - Extensions can be **costly and disruptive**
  - After an extension bucket table may no longer fit in memory
  - Imagine three records whose keys share the first 20 bits. These three records cannot be in same block (assume two records per block). But a block split would require setting  $i = 20$ , i.e., accommodating for  $2^{20} = 1 \text{ million buckets}$ , even though there may be only a few hundred records.

# Outline

- Hash Tables
  - ✓ Secondary storage HT
  - ✓ Extensible HT
- Linear HT

# Linear Hash Table

- Idea 1: add only one bucket at a time

Problem:  $n =$  no longer a power of 2

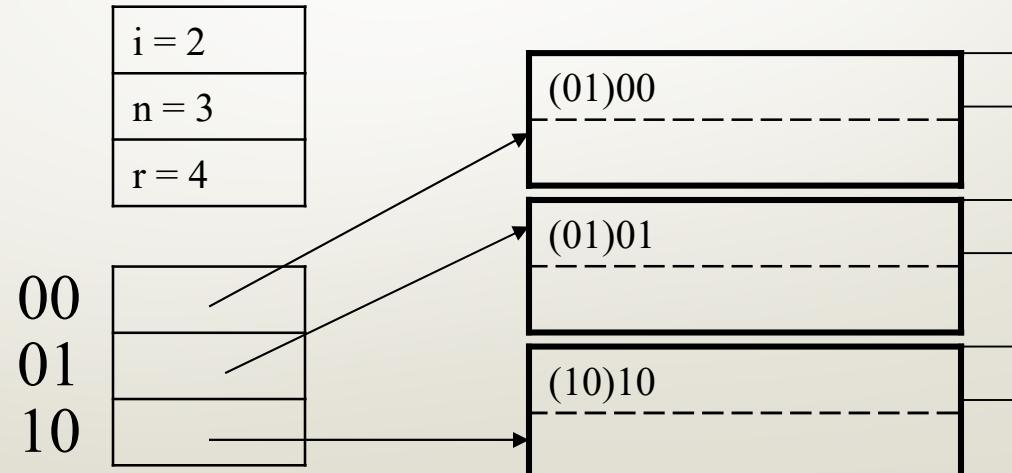
- Let  $i$  be # bits necessary to address  $n$  buckets.
  - $i = \text{ceil}(\log_2 n)$
- After computing  $h(k)$ , use *last*  $i$  bits:
  - If last  $i$  bits represent a number (say  $m$ )  $< n$ , store the key in bucket  $m$
  - If  $m \geq n$ , change msb from 1 to 0 (get a number  $< n$ )
- Idea 2: allow overflow blocks (not expensive to overflow)
- Convention: Read from the right (as opposed to the left)

# Linear Hash Table Example

- $N=3 \leq 2^2 = 4$ 
  - Therefore, only buckets until 10

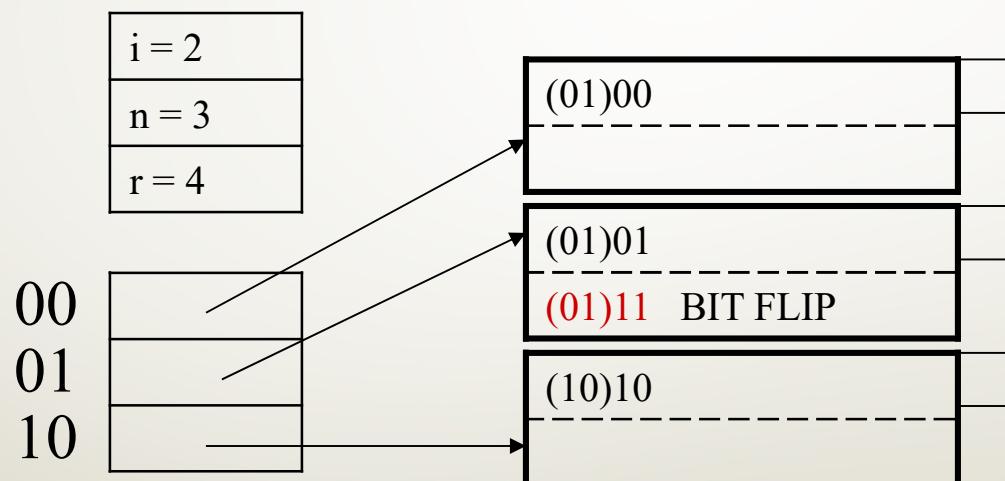
Try to insert 011

11 is flipped => 01



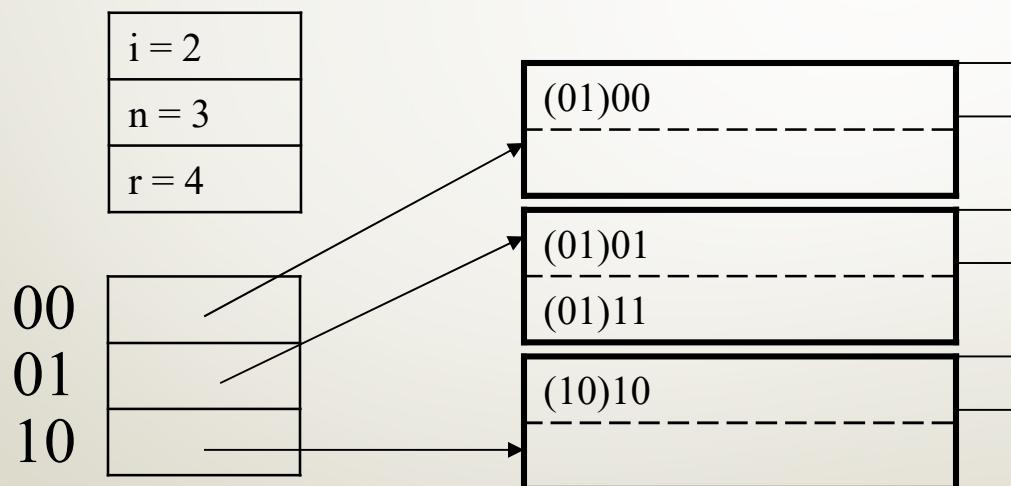
# Linear Hash Table Example

- After inserting 0111



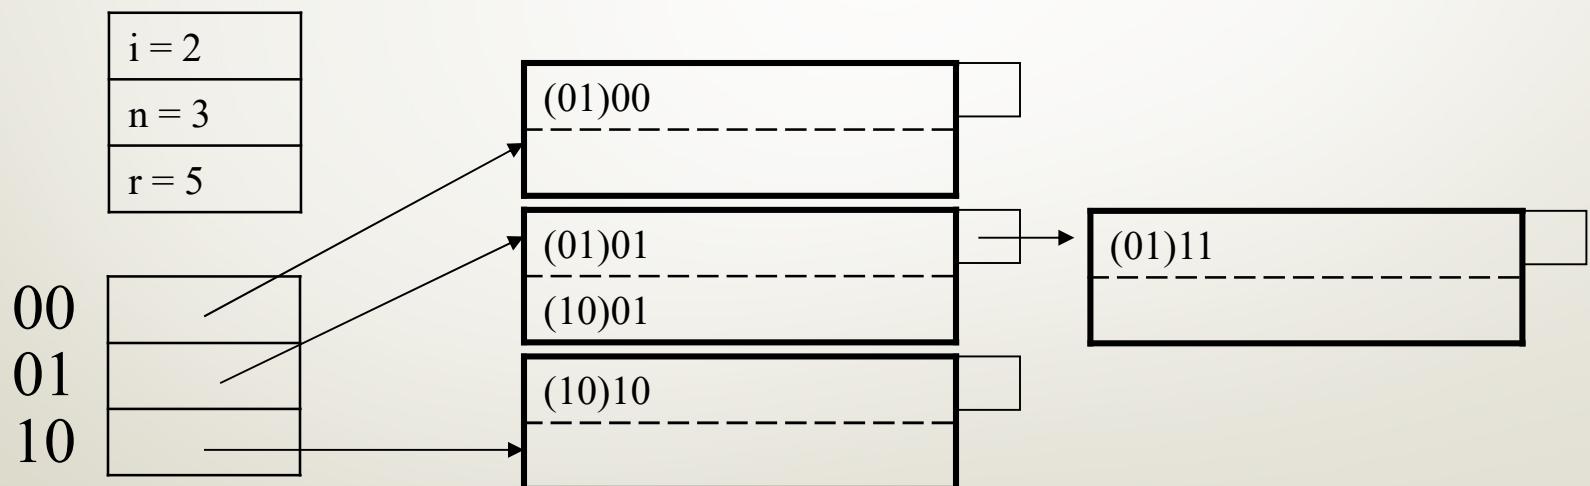
# Linear Hash Table Example

- Insert 1001:



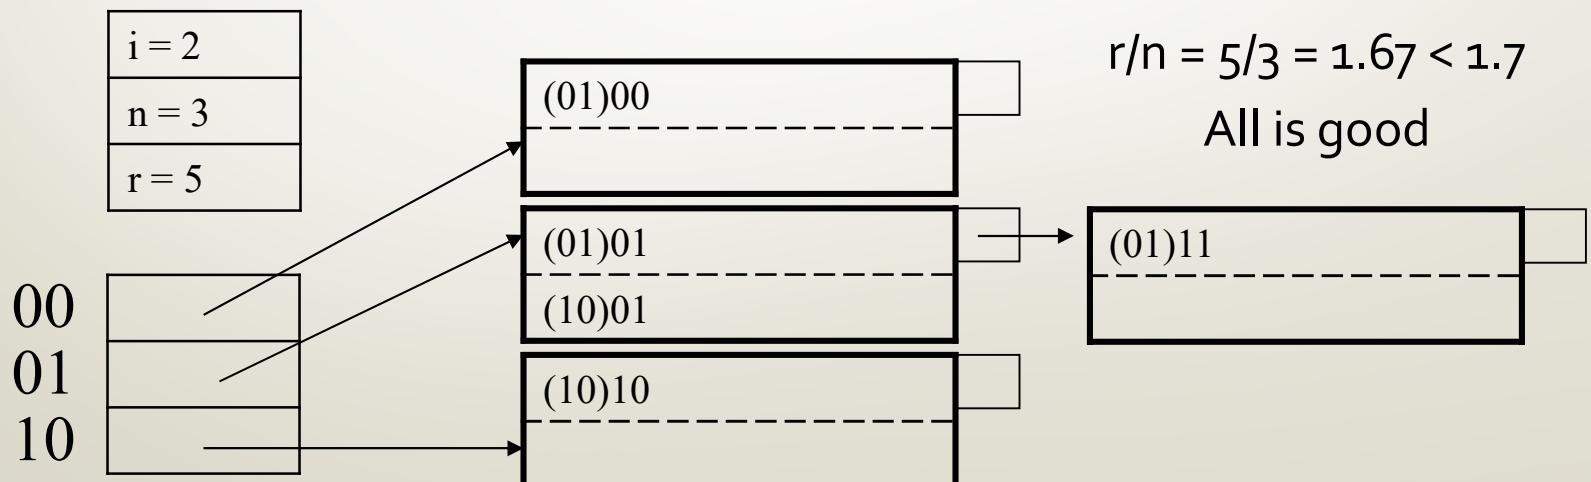
# Linear Hash Table Example

- Insert 1001: overflow blocks...



# Linear Hash Tables

- Extend  $n \rightarrow n+1$  when average number of records per bucket exceeds (say) 85% of total number of records per block
  - e.g.,  $r/n \leq 0.85 * 2 = 1.7$  (for block size = 2)
- Until then, use overflow blocks (cheaper than adding buckets)

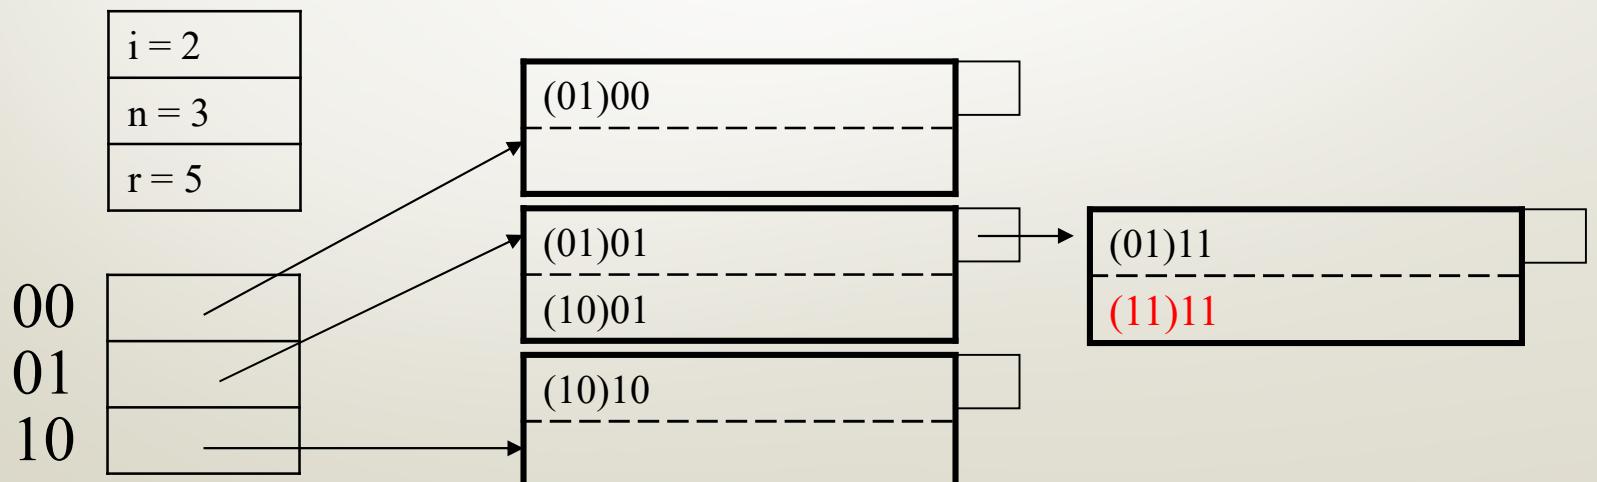


# Linear Hash Tables

- Try to insert 1111

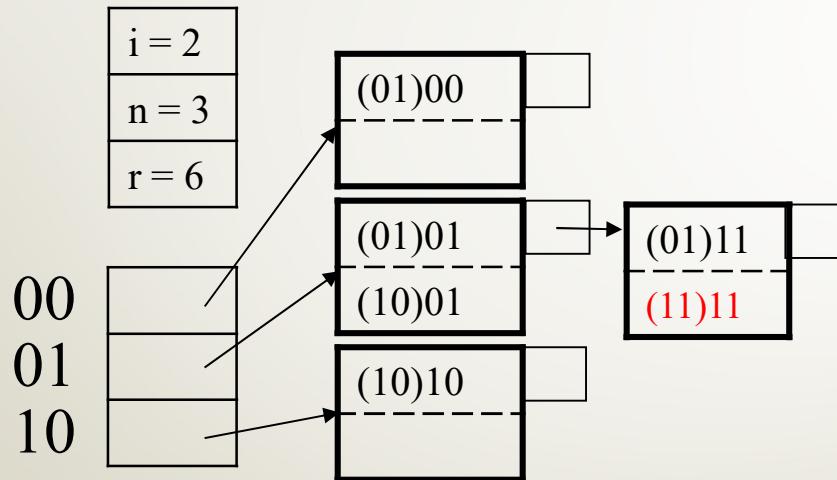
$$r/n = 6/3 = 2 > 1.7$$

→ Time to add a bucket

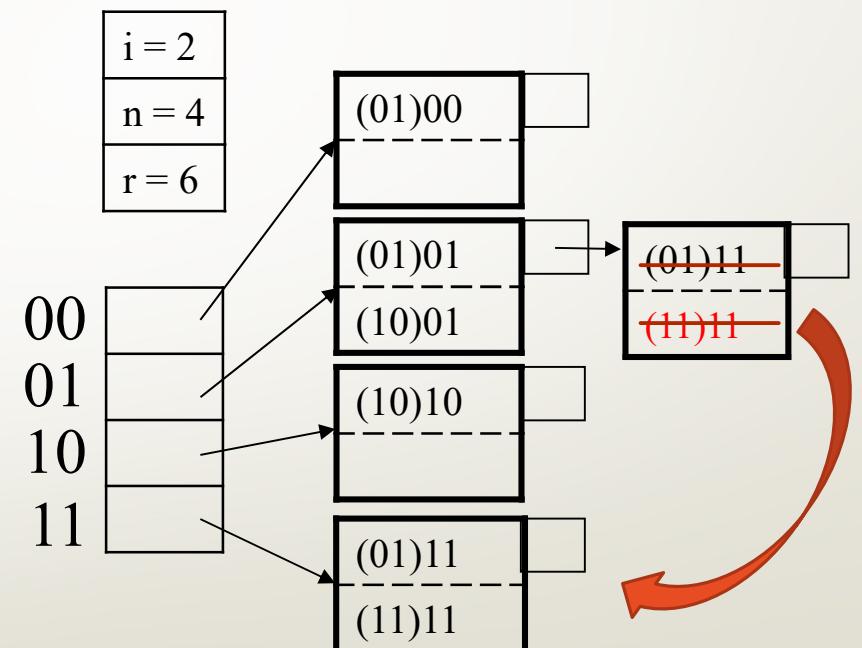


# Linear Hash Table Extension

- From  $n=3$  to  $n=4$



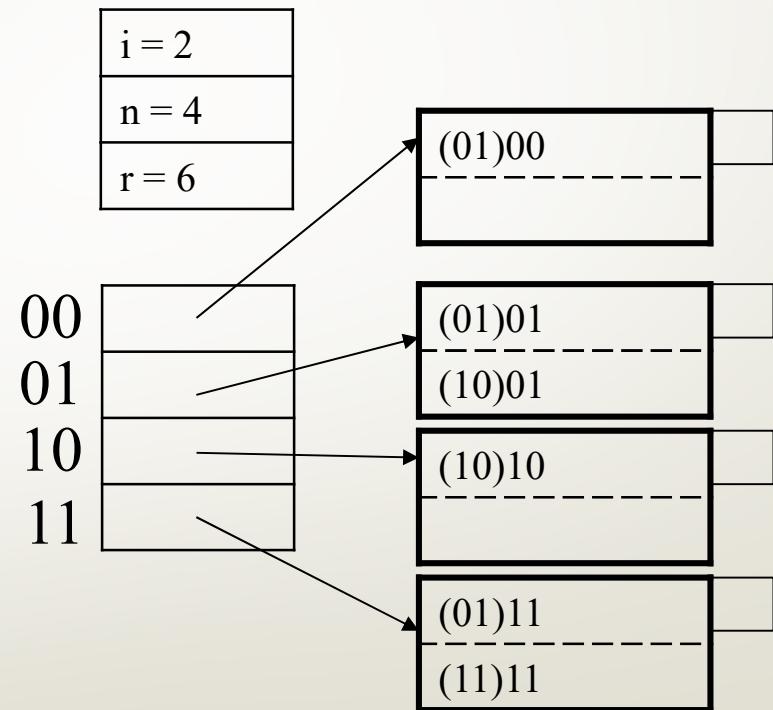
- Only need to touch one block (which one ?)



# Linear Hash Table Extension

- From n=3 to n=4 finished

$$r/n = 6/4 = 1.5 < 1.7 \quad \checkmark$$



# Indexing Summary

- B+ Trees (search, insertion, deletion)
  - Good for point and range queries
  - Log time lookup, insertion and deletion because of balanced tree
- Hash Tables (search, insertion)
  - Static hash tables: one I/O lookup, unless long chain of overflow
  - Extensible hash tables: one I/O lookup, extension can take long
  - Linear hash tables: ~ one I/O lookup, cheaper extension
- No panacea; dependent on data and use case