



# Transaction Management

**Ch 6.6.1 - 6.6.3  
and 18.1-18.4**

**Abdu Alawini**

University of Illinois at Urbana-Champaign

CS411: Database Systems

# Learning Objectives


After this lecture, you should be able:

- Describe transactions, ACID properties and issues related to concurrent execution of transactions.
- Describe how database systems implement isolation levels

# **Outline**

- Transactions and ACID properties
- Transactions in SQL
- Isolation levels
- Isolation Levels and Locking

# DBMS Architecture



# Motivating “Transactions”

- We've learned how to interact with DB using SQL.
- We assumed that:
  - each **operation** (e.g., UPDATE ... SET ... WHERE) is **executed one at a time**.
  - One operation executes, perhaps changes the DB state, then next operation executes. (**ISOLATION**)
  - each operation is executed in entirety or not executed at all. (**ATOMIC**)


# Motivating “Transactions”

- Complications arise if these assumptions are violated
  - multiple operations acting on the same table simultaneously?
  - system crash in the middle of an operation
    - e.g., half the tuples have been updated the others not

## Example 1: flight seat selection

- Program:
  - Check if seat is available : SELECT FROM WHERE (SFW)
  - Book seat (change availability to “occupied”): UPDATE SET WHERE (USW)

Two simultaneous runs



Two executions of the same “UPDATE ... SET ... WHERE” leads to seat being double-booked

## Example 1: lesson

- Group the SFW (which retrieved seat availability) and the USW (which reserved a seat) into one **TRANSACTION**.
- Transaction is a sequence of statements that are considered a “**unit of operation**” on a database
- Either user 1’s transaction executes first and then user 2’s transaction, or the other way, but not in parallel. ***Serializability of transactions.***

## Example 2: bank inter-account transfer

Problems can also occur if a crash occurs in the middle of executing a transaction:

Transfer

read(X.bal)  
read(Y.bal)  
 $\{X.bal = X.bal - \$100\}$   
write(X.bal)

CRASH 

$Y.bal = Y.bal + \$100$   
write(Y.bal)

Need to guarantee that the write to X does not persist (**ABORT**)

- Default assumption if a transaction doesn't commit

## Example 2: lesson

- Steps 1 and 2 must be done as one unit. *Atomically*.
- Either they both execute, or neither does.
- *Transactions must be atomic*.

# Transactions

**Definition:** **Transaction** is a sequence of read and write operations on the database with the property that either all actions complete or none completes.

a **transaction** may

- Succeed: the effects of all write operations persist (*COMMIT*)
- Or fail, no effects persist (*ABORT* or *ROLLBACK*)
- These guarantees are made despite concurrent activity in the system, and despite failures that may occur

# Transaction Manager

The part DBMS that ensures a transaction is executed as expected.

- Purpose 1: Ensure that transactions that execute in parallel don't *interfere* with each other.
- Purpose 2: Talks to "Log Manager", ensures that steps inside a transaction are being "logged".
- Purpose 3: Performs recovery after crashes, using logs.

We will not cover  
recovery in this class.

# ACID Properties

## Atomicity

- either all of the actions of a transaction are executed, or none are.

## Consistency

- each transaction executed in isolation keeps the database in a consistent state

## Isolation

- Transactions are isolated from the effects of other, concurrently executing, transactions.

## Durability

- updates stay in the DBMS!!!

# Outline

- ✓ Transactions and ACID properties
- Transactions in SQL
- Isolation levels
- Isolation Levels and Locking

# Transactions in SQL

- A transaction begins when any SQL statement that queries the db begins.
- To end a transaction, the user issues a **COMMIT** or **ROLLBACK** statement.

May be omitted if  
autocommit is off

**START TRANSACTION**

[SQL statements]

**COMMIT or ROLLBACK;**

# Read-Only Transactions

- When a transaction only reads information, we have more freedom to let the transaction execute concurrently with other transactions.
- We signal this to the system by stating:

```
SET TRANSACTION READ ONLY;  
SELECT * FROM Accounts  
WHERE account#= '1234' ;  
...
```

# Read-Write Transactions

- If we state “read-only”, then the transaction cannot perform any updates.

ILLEGAL!

```
SET TRANSACTION READ ONLY;  
UPDATE Accounts  
SET balance = balance - $100  
WHERE account#= '1234';...
```

- Instead, we must specify that the transaction may update (the default):

```
SET TRANSACTION READ WRITE;  
update Accounts  
set balance = balance - $100  
where account#= '1234';...
```

# Concurrent Execution Problems

- **Write-Read conflict (WR): dirty read, inconsistent read**

A transaction reads a value written by another transaction that has not yet committed

- **Read-Write conflict (RW): unrepeatable read**

A transaction reads the value of the same object twice.

Another transaction modifies that value in between the two reads

- **Write-Write conflict (WW): lost update**

Two transactions update the value of the same object. The second one to write the value overwrites the first change

# Example of Dirty Reads

sid	name	zip
123	Qin	14001
321	Kristin	14104

## Transaction 1

```
START TRANSACTION;
```

```
SELECT zip FROM student WHERE sid = 123;  
/* will read 14111*/
```

## Transaction 2

```
START TRANSACTION;
```

```
UPDATE student  
SET zip = 14111 WHERE sid = 123;
```

```
ROLLBACK  
/* zip code will revert to 14001*/
```

# Example of Non-repeatable Reads

sid	name	zip
123	Qin	14001
321	Kristin	14104

Transaction 1

```
START TRANSACTION;
```

```
SELECT zip FROM student WHERE sid = 123;  
/* will read 14001*/
```

Transaction 2

```
START TRANSACTION;
```

```
UPDATE student  
SET zip = 14111 WHERE sid = 123;  
COMMIT;
```

```
SELECT zip FROM student WHERE sid = 123;  
/* will read 14111*/
```

## Another concurrency problem: Phantom Reads

A **phantom read** occurs when, in the course of a transaction, new rows are added by another transaction to the records being read.\*

\* [https://en.wikipedia.org/wiki/Isolation\\_\(database\\_systems\)](https://en.wikipedia.org/wiki/Isolation_(database_systems))

# Example of Phantom Reads

sid	name	zip
123	Qin	14001
321	Kristin	14104

## Transaction 1

```
START TRANSACTION;  
  
SELECT COUNT(*) FROM students WHERE sid  
BETWEEN 100 AND 400; /* will return 2 */
```

## Transaction 2

```
START TRANSACTION;  
  
INSERT INTO students(sid,name,zip)  
VALUES ( 100, 'Abdu', 14444 );  
  
COMMIT;
```

```
SELECT COUNT(*) FROM students WHERE sid  
BETWEEN 100 AND 400; /* will return 3 */
```

# Outline

- ✓ Transactions and ACID properties
- ✓ Transactions in SQL
- Isolation levels
- Isolation Levels and Locking

# Isolation

- The problems we've seen are all related to *isolation*
- General rules of thumb w.r.t. isolation:
  - Fully serializable isolation is more expensive than “no isolation”
    - We can't do as many things concurrently (or we have to undo them frequently)
  - For performance, we generally want to specify the most relaxed *isolation level* that's acceptable
    - Note that we're “slightly” violating a correctness constraint to get performance!

# Specifying Acceptable Isolation Levels

- To signal to the system that a **dirty read** is acceptable,

```
SET TRANSACTION  
ISOLATION LEVEL READ UNCOMMITTED;
```

- In addition, there are

```
SET TRANSACTION  
ISOLATION LEVEL READ COMMITTED;  
SET TRANSACTION  
ISOLATION LEVEL REPEATABLE READ;
```

# Specifying Acceptable Isolation Levels

- **READ COMMITTED:**
  - Forbids the reading of dirty (uncommitted) data, but allows a transaction T to issue the same query several times and get different answers
  - No value written by T can be modified until T completes
- **REPEATABLE READ**
  - If a tuple is retrieved once it will be retrieved again if the query is repeated
  - But it is **NOT** a guarantee that the same query will get the same answer!

# Summary of Isolation Levels

Level	Dirty Read	Unrepeatable Read	Phantoms
READ UN-COMMITTED	Maybe	Maybe	Maybe
READ COMMITTED	No	Maybe	Maybe
REPEATABLE READ	No	No	Maybe
SERIALIZABLE	No	No	No

# Outline

- ✓ Transactions and ACID properties
- ✓ Transactions in SQL
- ✓ Isolation levels
- Isolation Levels and Locking

# Implementing Isolation Levels

- One approach – use locking at some level (tuple, page, table, etc.):
  - each data item is either locked (in some mode, e.g. **shared** or **exclusive**) or is available (**no lock**)
  - an action on a data item can be executed if the transaction holds an appropriate lock
  - consider *granularity* of locks – how big of an item to lock
    - Larger granularity = fewer locking operations but more contention!
- Appropriate locks:
  - Before a read, a **shared lock** must be acquired
  - Before a write, an **exclusive lock** must be acquired

# Lock Compatibility Matrix

Locks on a data item are granted based on a **lock compatibility matrix**:

		Mode of Data Item		
		None	Shared	Exclusive
Request mode {	Shared	Y	Y	N
	Exclusive	Y	N	N

When a transaction requests a lock, it must wait (**block**) until the lock is granted

# Locks Prevent “Bad” Execution

If the system used locking, the first “bad” execution could have been avoided: X = Seat 22A

User 1	User 2
<code>xlock(X)</code> <code>read(X.avail)</code> <code>IF (!X.avail)</code> <code>{X.avail := 1}</code> <code>write(X.avail)</code> <code>release(X)</code>	<code>xlock(X)</code> is not granted <code>read(X.avail)</code> <code>IF(!X.avail)</code> <code>{X.avail := 1}</code> <code>write(X.avail)</code> <code>release(X)</code>

Is this a “good”  
concurrent execution?

User 1                  User 2

slock(X)

read(X.avail)

slock(X)

read(X.avail)

release(X)

release(X)

IF (!X.avail)

{X.avail := 1}

IF (!X.avail)

{X.avail := 1}

xlock(X)

write(X.avail)

release(X)

xlock(X)

write(X.avail)

release(X)

No, this is a “bad” execution.  
Locks are not enough

User 1	User 2
slock(X)	
read(X.avail)	
	slock(X)
	read(X.avail)
release(X)	release(X)
	IF (!X.avail) {X.avail := 1}
IF (!X.avail) {X.avail := 1}	
xlock(X)	
write(X.avail)	
release(X)	xlock(X)
	write(X.avail)
	release(X)

# Isolation Levels and Locking

**READ UNCOMMITTED** allows queries in the transaction to **read** data *without acquiring any lock*

- Access mode READ ONLY, no updates are allowed

**READ COMMITTED** requires a read-lock to be obtained for all tuples touched by queries, but it **releases the locks immediately after the read**

- Exclusive locks must be obtained for updates and held to end of transaction

## Isolation levels and locking, cont.

**REPEATABLE READ** places shared locks on tuples retrieved by queries, holds them until the end of the transaction

- Exclusive locks must be obtained for updates and held to end of transaction

**SERIALIZABLE** places shared locks on tuples retrieved by queries *as well as the index*, holds them until the end of the transaction

- Exclusive locks must be obtained for updates and held to end of transaction

Holding locks to the end of a  
trxn is called “**strict**” locking

# A Notation for Transactions and Schedules

- We only care about reads and writes
  - Transaction is a sequence of  $R(A)$ 's and  $W(A)$ 's

T1:  $r_1(A); w_1(A); r_1(B); w_1(B)$

T2:  $r_2(A); w_2(A); r_2(B); w_2(B)$

- E.g. of a schedule:

$r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B)$

# Concurrent execution: Example 1

- Suppose we have two transactions:

T1: SET TRANSACTION READ WRITE

ISOLATION LEVEL READ COMMITTED;

SQL code that translates to: R1(A), R1(B) W1(B) W1(C)

T2: SET TRANSACTION READ WRITE

ISOLATION LEVEL READ COMMITTED;

SQL code that translates to: R2(C), R2(A) W2(A)

One possible interleaved execution of the transactions above:

R1(A) R2(C) R2(A) W2(A) R1(B) W1(B) W1(C)

S1(A) R1(A) REL1(A) S2(C) R2(C) REL2(C) S2(A) R2(A) X2(A) W2(A)  
REL2(A) S1(B) R1(B) X1(B) W1(B) X1(C) W1(C) REL1(B,C)

## Concurrent execution: Example 2

- Now suppose that T1 is REPEATABLE READ:

T1: SET TRANSACTION READ WRITE

ISOLATION LEVEL REPEATABLE READ;

SQL code that translates to: R1(A), R1(B) W1(B) W1(C)

T2: SET TRANSACTION READ WRITE

ISOLATION LEVEL READ COMMITTED;

SQL code that translates to: R2(C), R2(A) W2(A)

One possible interleaved execution of the transactions above:

R1(A) R2(C) R2(A) R1(B) W1(B) W1(C) W2(A)

S1(A) R1(A) S2(C) R2(C) REL2(C) S2(A) R2(A) Rel2(A) S1(B) R1(B) X1(B) W1(B) X1(C)  
W1(C) REL1(A, B, C) X2(A) W2(A) REL2(A)

# Outline

- ✓ Transactions and ACID properties
- ✓ Transactions in SQL
- ✓ Isolation levels
- ✓ Isolation Levels and Locking