



# Transaction Management: **Theory of Serializability**

**Abdu Alawini**

University of Illinois at Urbana-Champaign

CS411: Database Systems

# Learning Objectives

After this lecture, you should be able to:

- Describe the theory of serializability, including Serial, Serializable, Conflict-S Serializable, and View-S Serializable schedules.
- Use the Precedence Graph Algorithm to test for conflict-serializability.
- Describe how commercial database systems use two-phase locking to guarantee serializability of concurrent transactions.

# Schedules

- **Schedule**: A sequence of interleaved actions from a set of transactions, **where the actions of each individual transaction are in the original order**.
- Represents an actual sequence of database actions.
- In a *complete* schedule, each transaction ends in *commit* or *abort*.
- Initial State of DB + Schedule → Final State of DB

## Example

A and B are elements  
in the database  
t and s are variables  
in tx source code

T <sub>1</sub>	T <sub>2</sub>
READ(A, t)	READ(A, s)
t := t+100	s := s*2
WRITE(A, t)	WRITE(A,s)
READ(B, t)	READ(B,s)
t := t+100	s := s*2
WRITE(B,t)	WRITE(B,s)

## Serial Schedules

Run transactions one at a time, in a series. (Different orders might give different results.)

T <sub>1</sub>	T <sub>2</sub>
READ(A, t)	
t := t+100	
WRITE(A, t)	
READ(B, t)	
t := t+100	
WRITE(B,t)	
	READ(A, s)
	s := s*2
	WRITE(A,s)
	READ(B,s)
	s := s*2
	WRITE(B,s)

## Serializable schedules

A schedule is **serializable** if it is equivalent to a serial schedule.

Final state must be the same as the state produced by **one of the serial schedules**.

T <sub>1</sub>	T <sub>2</sub>
READ(A, t)	
$t := t + 100$	
WRITE(A, t)	
	READ(A, s)
	$s := s^* 2$
	WRITE(A, s)
READ(B, t)	
$t := t + 100$	
WRITE(B, t)	
	READ(B, s)
	$s := s^* 2$
	WRITE(B, s)

$T_1$	$T_2$
READ(A, t)	
$t := t + 100$	
WRITE(A, t)	
	READ(A, s)
	$s := s^* 2$
	WRITE(A, s)
	READ(B, s)
	$s := s^* 2$
	WRITE(B, s)
READ(B, t)	
$t := t + 100$	
WRITE(B, t)	

**Is this a Serializable schedule?**

## No, it's a **non-Serializable** schedule

T <sub>1</sub>	T <sub>2</sub>
READ(A, t)	
t := t+100	
WRITE(A, t)	
	READ(A, s)
	s := s*2
	WRITE(A,s)
	READ(B,s)
	s := s*2
	WRITE(B,s)
READ(B, t)	
t := t+100	
WRITE(B,t)	

# Outline

- ✓ Theory of Serializability
  - ✓ Serial and serializable schedules
  - Conflict-Serializable schedules
  - Two-Phased Locking Theorem
    - Two-Phased Locking (2PL)
    - Strict two-phased locking (S2PL)

# **Conflicts**

- Write-Read – WR
- Read-Write – RW
- Write-Write – WW

# Conflict-Serializability

**Definition.** A schedule is *conflict-serializable* if it can be transformed into a serial schedule by a series of swappings of adjacent non-conflicting actions

a *conflict-serializable* schedule is a *Serializable* schedule

- The converse is not true in general (see textbook page 893)

**Conflict-serializability** is a condition that the schedulers in commercial systems generally use when they need to guarantee serializability.

# Conflict Serializability

Swapping Conflicts:

- Two actions by same transaction  $T_i$ :

$R_i(X), W_i(Y)$

- Two writes by  $T_i, T_j$  to same element:

$W_i(X), W_j(X)$


- Read/write by  $T_i, T_j$  to same element:

$W_i(X), R_j(X)$

$R_i(X), W_j(X)$

# Conflict-Serializability

Example:



# Testing for Conflict-Serializability

Given a schedule  $S$ , we can construct a directed graph  $G=(V,E)$  called a **precedence graph**

- $V$  : all transactions in  $S$
- $E$  :  $T_i \rightarrow T_j$  whenever an action of  $T_i$  precedes and conflicts with an action of  $T_j$  in  $S$  (RW, WR, WW)


**Theorem:** A schedule  $S$  is **conflict serializable** iff its precedence graph contains no cycles

## Example 1

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$



## Example 1



This Schedule is **conflict-serializable**

# View Serializable!

Is this schedule  
conflict-serializable?

In both schedules, database is in  
the same state:  
A written by T2 and B written by T3

Both schedules are equivalent  
But not conflict-equivalent

A schedule is **view serializable** if it is  
**view-equivalent to a serial schedule**

T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
W(A)		
	W(A)	
		W(B)
	W(B)	
		W(B)
T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
W(A)		
	W(B)	
		W(A)
		W(B)
		W(B)

# How to find if two schedules are View Equivalent

Two schedules  $S_1$  and  $S_2$  are *view equivalent* if:

- If  $T_i$  reads an initial value of A in  $S_1$ , then  $T_i$  reads the initial value of A in  $S_2$
- If  $T_i$  reads a value of A written by  $T_j$  in  $S_1$ , then  $T_i$  reads a value of A written by  $T_j$  in  $S_2$
- If  $T_i$  writes the final value of A in  $S_1$ , then  $T_i$  writes the final value of A in  $S_2$

A schedule is **view serializable** if it is  
**view-equivalent** to a serial schedule

# Outline

- ✓ Theory of Serializability
  - ✓ Serial and serializable schedules
  - ✓ Conflict-Serializable schedules
- Two-Phased Locking Theorem
  - Two-Phased Locking (2PL)
  - Strict two-phased locking (S2PL)

## Well-Formed, Two-Phased Transactions

A transaction is **well-formed** if it acquires at least a shared lock on Q before reading Q or an exclusive lock on Q before writing Q and doesn't release the lock until the action is performed

A transaction is **two-phased** if it never acquires a lock after unlocking one

- ❑ There are two phases:
  - ❑ a *growing phase* in which the transaction acquires locks
  - ❑ a *shrinking phase* in which locks are released

## Two-Phased Locking Theorem


If all transactions are **well-formed** and **two-phase**,  
then any schedule in which conflicting locks are  
never granted ensures serializability

# Two Phase Locking Protocol (2PL)

2PL is a way of managing locks during a transaction T

- T gets (S and X) locks gradually, as needed
- T cannot request any additional locks once it releases any locks

# of locks held by a transaction T



## Can this schedule arise under 2PL?

T1:	R(A)	W(A)	R(B)	W(B)	
T2:		R(A)	W(A)	R(B)	W(B)

## Enforce 2PL

T1:	R(A),	R(C)	W(B)
T2:	R(B)	W(C)	R(B)
T3:	R(A)		R(B)

**S1(A); R1(A); S2(B) R2(B); S3(A); R3(A);  
X2(C); W2(C); REL2(C); S1(C); R1(C);  
R2(B); REL2(B); X1(B); W1(B); REL1(A,B,C);  
S3(B); R3(B); REL3(A,B);**

# A problem with 2PL

T <sub>1</sub> : R(A) W(A)	R(B) W(B)
T <sub>2</sub> :	R(A) W(A) R(B) W(B)

**Unrecoverable schedule!**

T <sub>1</sub>	T <sub>2</sub>
S(A)	
R(A)	
X(A)	
W(A)	
	S(A) <Denied, wait>
S(B)	
R(B)	
X(B)	
W(B)	
Rel(A)	
	S(A) <Granted>
	R(A)
	X(A)
	W(A)
Rel(B)	
	S(B)
	R(B)
	X(B)
	W(B)
ABORT	COMMIT


# Outline

- ✓ Theory of Serializability
  - ✓ Serial and serializable schedules
  - ✓ Conflict-Serializable schedules
- Two-Phased Locking Theorem
  - ✓ Two-Phased Locking (2PL)
  - Strict two-phased locking (S2PL)

## Strict Two Phase Locking Protocol (S2PL)

**Strict 2PL** is a way of managing locks during a transaction T

- T gets (S and X) locks gradually, as needed
- T holds all locks until end of transaction (commit/abort)



## Enforce S2PL

T1:	R(A),	R(C)	W(B)
T2:	R(B)	W(C)	R(B)
T3:	R(A)		R(B)

**S1(A); R1(A); S2(B) R2(B); S3(A); R3(A);  
X2(C); W2(C); S1(C) <T1 WAIT ON C>;  
R2(B); COMMIT REL2(B,C); S1(C); R1(C);  
X1(B); W1(B); COMMIT; REL1(A,B,C); S3(B);  
R3(B); COMMIT; REL3(A,B);**

# Unrecoverable schedule solved with S2PL

T <sub>1</sub> :	R(A)	W(A)		R(B)	W(B)
T <sub>2</sub> :		R(A)	W(A)	R(B)	W(B)

Recoverable schedule!

T <sub>1</sub>	T <sub>2</sub>
S(A)	
R(A)	
X(A)	
W(A)	
	S(A) <Denied, wait>
S(B)	
R(B)	
X(B)	
W(B)	
ABORT Rel(A, B)	
	S(A) <Granted>
	R(A)
	X(A)
	W(A)
	S(B)
	R(B)
	X(B)
	W(B)
	COMMIT, Rel(A, B)

# Strict 2PL guarantees serializability

- Can prove that a Strict 2PL schedule is equivalent to the serial schedule in which each transaction runs instantaneously at the time that it commits
- This is **huge**: A property of each transaction (2PL) implies a property of any set of transactions (serializability)
  - No need to check serializability of specific schedules
- Most DBMSs use 2PL to enforce serializability

# Summary

- Transactions are all-or-nothing units of work guaranteed despite concurrency or failures in the system.
- Theoretically, the “correct” execution of transactions is **serializable** (i.e. equivalent to some serial execution).
- Practically, this may adversely affect throughput  $\Rightarrow$  **isolation levels**.
- With isolation levels, users can specify the level of “incorrectness” they are willing to tolerate.