



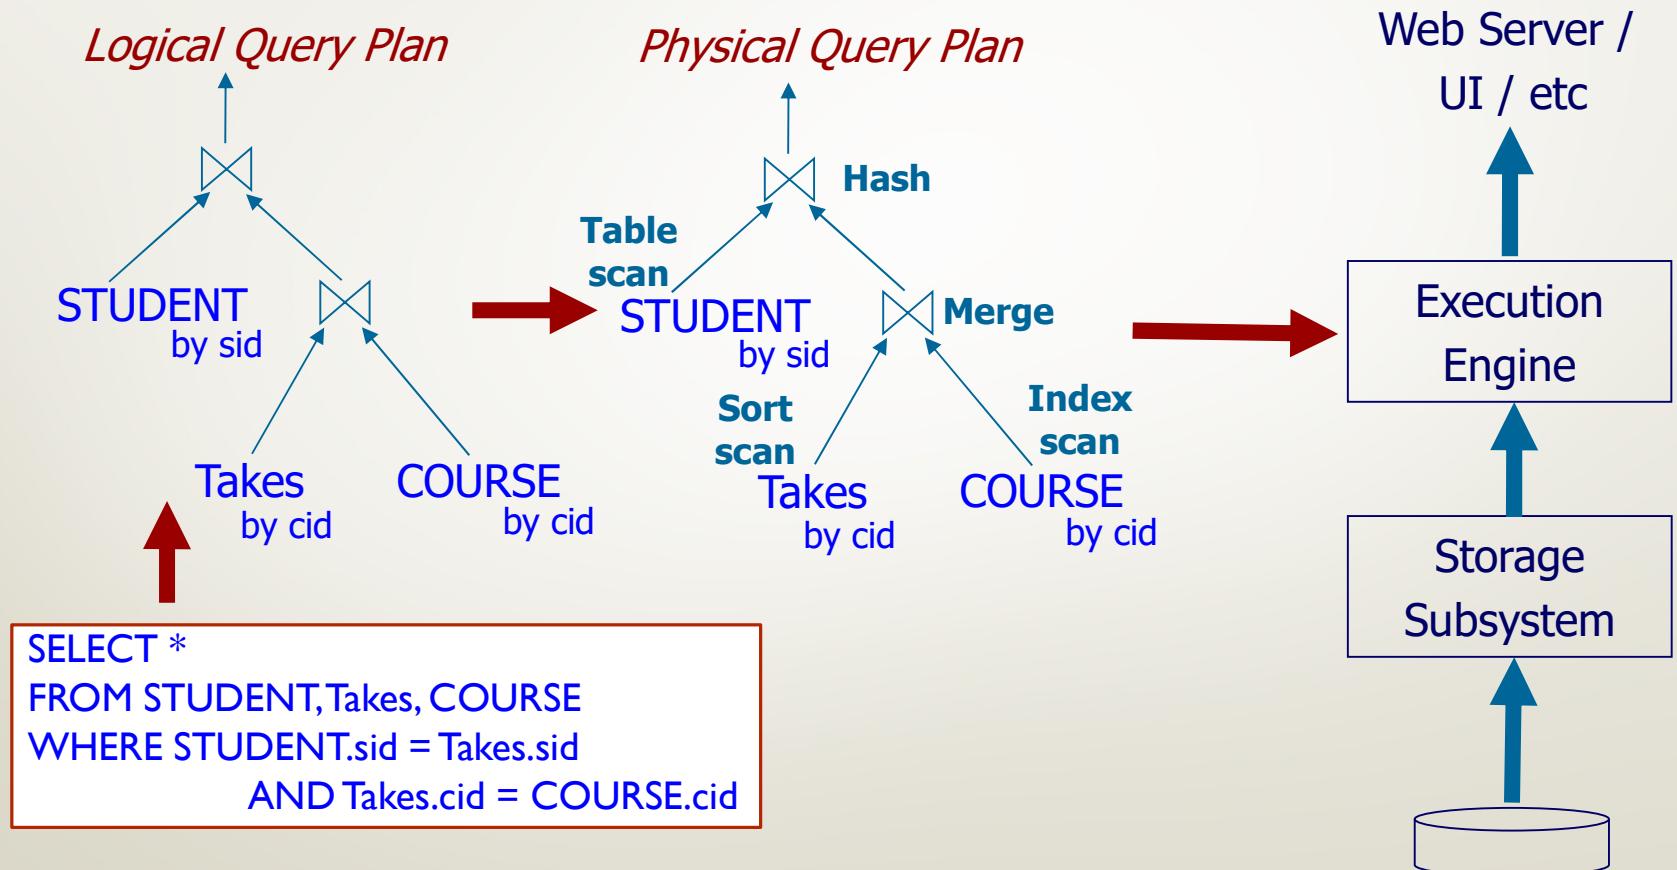
Query Processing: Physical Operators

Abdu Alawini

University of Illinois at Urbana-Champaign

CS411: Database Systems

QP: The Big Picture: SQL → Logical Query Plan Plan → Physical Query Plan



Logical v.s. Physical Operators

- Logical operators
 - what they do
 - e.g., union, selection, project, join, grouping
- Physical operators
 - how they do it
 - e.g., nested loop/sort-merge/hash/index join
 - In other words, physical operators are particular implementations of relational algebra operations
 - Physical operators also pertain to none RA operations, such as “scanning” a table.

Physical operators and costs

Before we describe implementations, we need to examine the *cost* involved

- Often, we have to make choices about which physical operators to use.
- For this, we need to estimate the “cost” of each physical operator.

Cost Parameters (or “statistics”)

Estimating the cost:

- Important in optimization (next topic: Qry Opt.)
 - Picking between plans
- Compute disk I/O cost only
 - Main memory operations are cheap
- We compute the **cost to *read* the arguments** of the operator
- We don’t compute the **cost to *write* the “final” result**
 - Same for all types of operators

Cost Parameters (or “statistics”)

- Cost parameters
 - M = number of blocks that fit in main memory
 - $B(R)$ = number of blocks needed to hold R (best case # blocks)
 - $T(R)$ = number of tuples in R (worst case # of blocks)
 - $V(R,a)$ = number of distinct values of the attribute a

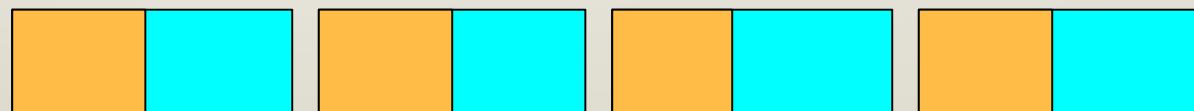
$T(R) / (\text{number of tuples that fits in a block}) = B(R)$

Here $T(R) = 4$; $B(R) = 2$

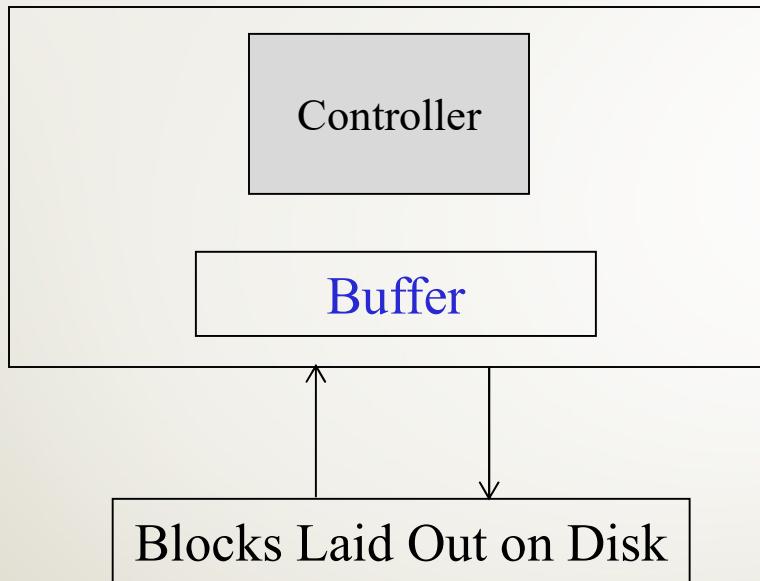
Clustered
File



Unclustered
File



A Mental Model for Costing Operator Implementations



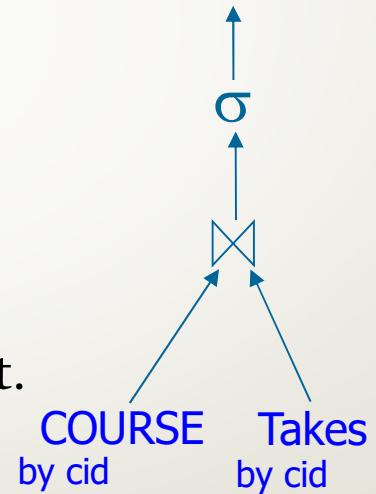
Temp computation
space in main memory

We compute the cost to *read* the arguments of the operator
We don't compute the cost to *write* the result

The iterator model for implementing operators

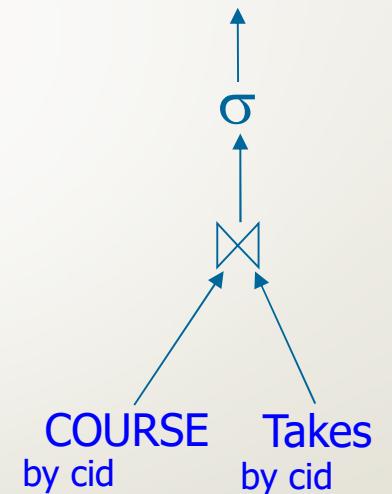
Each (physical) operation is implemented by 3 functions:

- **Open**: sets up the data structures and performs initializations
- **GetNext**: returns the the next tuple of the result.
- **Close**: ends the operations. Cleans up the data structures.



The iterator model for implementing operators

- Enables pipelining!
 - As opposed to: execute each operator in entirety, store its results on disk or in main memory
 - Many operators can be “active” simultaneously
- Pipelining: Not always possible (or meaningful):
 - E.g., “sort scan”.
 - “Blocking” operators

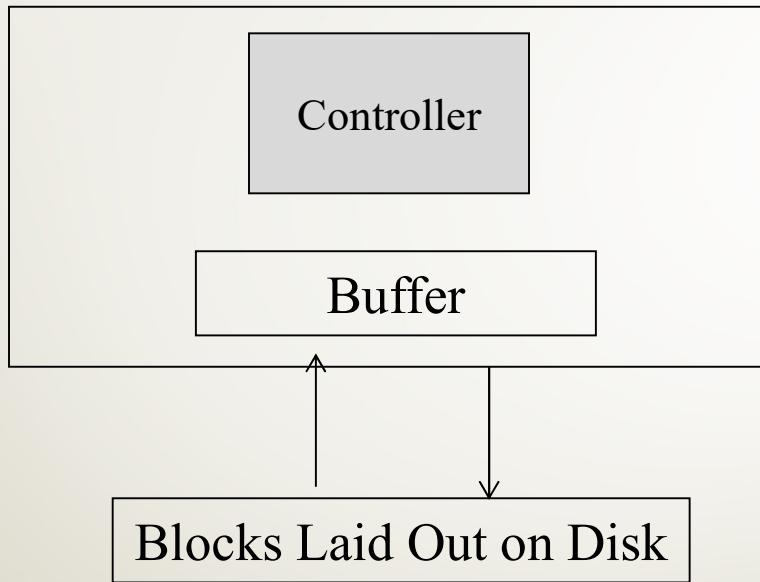


Overview of operator implementations: first classification

Operator algorithms mostly of one of these types:

- Sorting-based
- Hash-based
- Index-based

Another classification of operator implementations



One pass:

- reading the data only once from disk.
- Typically, at least one of the arguments must fit in memory.

Two pass:

- Data need not fit in memory but is not “too large”.

Multipass:

- No limit on data size.

How to think about operators

- Classification 1: Hash/Index/Sort
- Classification 2: One/Two/Many pass
- Metric 1: I/O Cost
- Metric 2: Buffer requirements (how large should M be)